

CMSC216: Introduction

Chris Kauffman

Last Updated:
Tue Jan 27 10:50:50 AM EST 2026

CMSC216 2xx: Logistics

Introductions

- ▶ Prof Kauffman: profk@umd.edu
- ▶ Office Hours Tue/Wed 4-5pm in IRB 2226
- ▶ Slides: Linked from Canvas “Course Schedule/Materials”
- ▶ Static link: <https://www.umd.edu/~profk/216/>

NOTE: Videos posted for Course Mechanics, Coding Environment Setup; will add Lectures and Lab01 at end of week

Reading

- ▶ Bryant/O'Hallaron: Ch 1
- ▶ C references: basic syntax, types, compilation

Goals

- ▶ Basic Model of Computation
- ▶ Begin discussion of C

Assignments

- Due Sun 31-Jan 11:59pm
- ▶ Lab01: Setup, submit to Gradescope
 - ▶ Set up your rig, attend discussion on Wed, and get coding!

Announcements

Not registered yet?

If you are in lecture but are not yet registered for the course, please consider registering for a different section with open seats: the longer you wait, the further you will be behind as the Kauffman 2xx sections have different assignments / topic order from the other sections.

Wednesday Discussion Sections will Meet

Staff will be on hand to assist students getting their environment configured and solving Lab01; please attend!

Zoom Etiquette

- ▶ Mute yourself unless called on to speak (then unmute)
- ▶ Video is optional at all times
- ▶ “Raise Hand” for questions / answers
- ▶ Chat will NOT be monitored for questions

"Von Kauffman" Model: CPU, Memory, Screen, Program

Most computers have 4 basic, physical components¹

1. CPU: can execute “instructions”
2. CONTROL: CPU knows WHICH instruction to execute
3. MEMORY: data is stored and can change
4. Input/Output devices like a SCREEN (optional)

CPU understands some **set of instructions**; a sequence of instructions is a **program** that changes MEMORY and SCREEN

Example of a Running Computer Program

CPU: at instruction 10:

```
> 10: set #1024 to 195
      11: set #1028 to 21
      12: sum #1024,#1028 into #1032
      13: print #1024, "plus", #1028
      14: print "is", #1032
```

MEMORY:

Addr	Value
#1032	-137
#1028	12
#1024	19

SCREEN:

¹Of course it's a *little* more complex than this but the adage, “All models are wrong but some are useful” applies here. This class is about asking “what is really happening?” and going deep down the resulting rabbit hole.

Sample Run Part 1

CPU: at instruction 10:

```
> 10: set #1024 to 195
  11: set #1028 to 21
  12: sum #1024,#1028 into #1032
  13: print #1024, "plus", #1028
  14: print "is", #1032
```

MEMORY:

Addr	Value
#1032	-137
#1028	12
#1024	19

SCREEN:

CPU: at instruction 11:

```
10: set #1024 to 195
> 11: set #1028 to 21
  12: sum #1024,#1028 into #1032
  13: print #1024, "plus", #1028
  14: print "is", #1032
```

MEMORY:

Addr	Value
#1032	-137
#1028	12
#1024	195

SCREEN:

CPU: at instruction 12:

```
10: set #1024 to 195
  11: set #1028 to 21
> 12: sum #1024,#1028 into #1032
  13: print #1024, "plus", #1028
  14: print "is", #1032
```

MEMORY:

Addr	Value
#1032	-137
#1028	21
#1024	195

SCREEN:

Sample Run Part 2

CPU: at instruction 13:

```
10: set #1024 to 195
11: set #1028 to 21
12: sum #1024,#1028 into #1032
> 13: print #1024, "plus", #1028
14: print "is", #1032
```

MEMORY:

Addr	Value
#1032	216
#1028	21
#1024	195

SCREEN:

CPU: at instruction 14:

```
10: set #1024 to 195
11: set #1028 to 21
12: sum #1024,#1028 into #1032
13: print #1024, "plus", #1028
> 14: print "is", #1032
```

MEMORY:

Addr	Value
#1032	216
#1028	21
#1024	195

SCREEN:

195 plus 21

CPU: at instruction 15:

```
10: set #1024 to 195
11: set #1028 to 21
12: sum #1024,#1028 into #1032
13: print #1024, "plus", #1028
14: print "is", #1032
> 15: ....
```

MEMORY:

Addr	Value
#1032	216
#1028	21
#1024	195

SCREEN:

195 plus 21
is 216

Observations: CPU and Program Instructions

- ▶ Program instructions are usually small, simple operations:
 - ▶ Put something in a specific memory cell using its **address**
 - ▶ Copy the contents of one cell to another
 - ▶ Do arithmetic (+, -, *, /) on cells or constants
 - ▶ Print stuff to the screen
- ▶ The CPU keeps track of which instruction to execute next
- ▶ After executing an instruction, CPU advances to the next instruction BUT **jumping** around to distant instructions is also possible: conditional and iterative execution
- ▶ Previous program is in **pseudocode** in which instructions can have any meaning understood by a human reader²
- ▶ Real machines require more precise instruction definitions as there are no smart humans to interpret them, only dumb physics to blindly move around electrons

²The pseudocode shown resembles a low-level **assembly language** rather than a high level language like C or Java

Observations: Memory Cells and the Screen

Memory Cells

- ▶ Memory cells have Fixed **ADDRESS** Changeable **CONTENTS**
- ▶ Random Access Memory (RAM): the contents in any memory cell can be retrieved FAST using its address
- ▶ My laptop has 16GB of memory = 4,294,967,296 (4 billion) integer boxes (!)
- ▶ Cell ADDRESS #'s never change: always cell #1024
- ▶ Cell CONTENTS / Values often change: set #1024 to 42

Screen versus Memory

- ▶ Nothing is on the screen until it is explicitly print-ed by the program
- ▶ Don't get to see memory while the program runs:
print stuff while debugging programs so you can see it
- ▶ Forming a mental model of what values are in memory and how they relate to one another is a valuable skill which we will practice, often by drawing memory contents explicitly

Variables are Named Memory Cells

- ▶ Dealing with raw memory addresses is tedious
- ▶ Any programming language worth its salt will have **variables**: symbolic names associated with memory cells
- ▶ **You pick variable names**; compiler/interpreter automatically translates to memory cell/address

PROGRAM ADDRESSES ONLY

CPU: at instruction 50:

```
> 50: copy #1024 to #1032
      51: copy #1028 to #1024
      52: copy #1032 to #1028
      53: print "first",#1024
      54: print "second",#1028
```

MEMORY:

Addr	Value
#1032	?
#1028	31
#1024	42

PROGRAM WITH NAMED CELLS

CPU: at instruction 51:

```
> 50: copy x to temp
      51: copy y to x
      52: copy temp to y
      53: print "first",x
      54: print "second",y
```

MEMORY:

Addr	Name	Value
#1032	temp	?
#1028	y	31
#1024	x	42

Correspondence of C Programs to Memory

- ▶ C programs require memory cell names to be declared with the **type of data** they will hold (*a novel idea when C was invented*).
- ▶ The equal sign (=) means "store the result on the right in the cell named on the left"
- ▶ Creating a cell and giving it a value can be combined

```
int x;           // need a cell named x, holds an integer
x = 42;          // put 42 in cell x
int y = 31;      // need a cell named y and put 31 in it
int tmp = x + y; // cell named tmp, fill with sum of x and y
```

Other Rules

- ▶ C/Java compilers read whole functions to figure out how many memory cells are needed based on declarations like int a; and int c=20;
- ▶ Lines that only declare a variable do nothing except indicate a cell is needed to the compiler
- ▶ When a function starts, enough memory for **all function variables** is made available immediately (including variables declared later in the function); *we'll find out how in this course*
- ▶ In C, uninitialized variables may have arbitrary crud in them making them dangerous to use: *we'll find out why in this course*

Exercise: First C Snippet

- ▶ Lines starting with // are comments, not executed
- ▶ `printf("x: %d y: %d\n",x,y);` shows variable values on the screen as %d is substituted for the value of additional parameters and shown as a *decimal integer*

CPU: at line 50

```
> 50: int x;  
51: x = 42;  
52: int y = 31;  
53: // swap x and y (?)  
54: x = y;  
55: y = x;  
56: printf("x: %d y: %d\n",x,y);
```

MEMORY:

Addr	Name	Value
#1032	y	?
#1028	x	?
#1024		

SCREEN:

With your nearby colleagues:

1. Show what memory / screen look like after running the program
2. **Correct** the program if needed: make swapping work

I will chat with a couple folks about their answers which will earn participation credit leading to **Bonus Engagement Points**.

Answer: First C Snippet

CPU: at line 54

```
50: int x;  
51: x = 42;  
52: int y = 31;  
53: // swap x and y (?)  
> 54: x = y;  
55: y = x;  
56: printf("x: %d y: %d\n",x,y);
```

MEMORY:

Addr	Name	Value
#1032	y	31
#1028	x	42
#1024		

SCREEN:

SCREEN:

CPU: at line 55

```
50: int x;  
51: x = 42;  
52: int y = 31;  
53: // swap x and y (?)  
54: x = y;  
> 55: y = x;  
56: printf("x: %d y: %d\n",x,y);
```

MEMORY:

Addr	Name	Value
#1032	y	31
#1028	x	31
#1024		

CPU: at line 57

```
50: int x;  
51: x = 42;  
52: int y = 31;  
53: // swap x and y (?)  
54: x = y;  
55: y = x;  
56: printf("x: %d y: %d\n",x,y);  
> 57: ...
```

MEMORY:

Addr	Name	Value
#1032	x	31
#1028	y	31
#1024		

SCREEN:

x: 31 y: 31

Clearly **incorrect**: how does one swap values properly? (fix `swap_main_bad.c`)

First Full C Program: swap_main.c

```
1 /* First C program showing a main() function. Demonstrates proper
2    swapping of two int variables declared in main() using a third
3    temporary variable. Uses printf() to print results to the screen
4    (standard out). Compile run with:
5
6    > gcc swap_main.c
7    > ./a.out
8 */
9
10 #include <stdio.h>           // headers declare existence of functions
11                                // printf in this case
12
13 int main(int argc, char *argv[]){ // ENTRY POINT: always start in main()
14     int x;                      // declare a variable to hold an integer
15     x = 42;                     // set its value to 42
16     int y = 31;                 // declare and set a variable
17     int tmp = x;                // declare and set to same value as x
18     x = y;                     // put y's value in x's cell
19     y = tmp;                   // put tmp's value in y's cell
20     printf("x: %d y: %d\n",x,y); // print the values of x and y
21     return 0;                  // return from main(): 0 indicates success
22 }
```

- ▶ Swaps variables using tmp space ([exotic alternatives exist](#))
- ▶ Executables always have a `main()` function: starting point
- ▶ Note inclusion of `stdio.h` **header** to declare `printf()` exists, allusions to C's (limited and clunky) library system

Exercise: Functions in C, swap_func.c

```
1 // C program which attempts to swap using a function.  
2 //  
3 // > gcc swap_func.c  
4 // > ./a.out  
5  
6 #include <stdio.h>           // declare existence printf()  
7 void swap(int a, int b);      // function exists, defined below main  
8  
9 int main(int argc, char *argv[]){ // ENTRY POINT: start executing in main()  
10    int x = 42;  
11    int y = 31;  
12    swap(x, y);              // invoke function to swap x/y (?)  
13    printf("x: %d y: %d\n",x,y); // print the values of x and y  
14    return 0;  
15 }  
16  
17 // Function to swap (?) contents of two memory cells  
18 void swap(int a, int b){      // arguments to swap  
19    int tmp = a;              // use a temporary to save a  
20    a = b;                   // a <- b  
21    b = tmp;                // b <- tmp=a  
22    return;  
23 }
```

Does swap() “work”? Discuss with neighbors and justify why the code works or why not

Answers: Swapping in a Function is Tricky

```
# Check if 42 swaps with 31
>> gcc swap_func.c                      # compile source
>> ./a.out                                # run executable
x: 42 y: 31                                # not swapped?
```

swap_func.c will not print swapped values

- ▶ If you thought the values would print swapped, you're about to learn something interesting
- ▶ If you were confident they would not print swapped but had difficulty articulating why, that's great: this class is here to give the vocab to do so
- ▶ If you knew the values wouldn't swap and also knew how to explain it well, tune in anyway as the subsequent explanation will introduce conventions used for the rest of the course

Why No Swap??

Necessitates introducing the **Function Call Stack** which is where functions store their local variables and parameters

Answers: The Function Call Stack and swap()

```
9: int main(...){  
10:    int x = 42;  
11:    int y = 31;  
+-<12:    swap(x, y);  
| 13:    printf("%d %d\n",x,y);  
| 14:    return 0;  
V 15: }  
  
|  
| 18: void swap(int a, int b){  
+->19:    int tmp = a;  
20:    a = b;  
21:    b = tmp;  
22:    return;  
23: }
```

STACK: Caller main(), prior to swap()
| FRAME | ADDR | SYM | VALUE |
|-----+-----+-----+-----|
| main() | #2048 | x | 42 | stack frame
| line:12 | #2044 | y | 31 | for main()
|-----+-----+-----+-----|

STACK: Callee swap() takes control
| FRAME | ADDR | SYM | VALUE |
|-----+-----+-----+-----|
| main() | #2048 | x | 42 | main() frame
| line:12 | #2044 | y | 31 | now inactive
|-----+-----+-----+-----|
| swap() | #2040 | a | 42 | new frame
| line:19 | #2036 | b | 31 | for swap()
| #2032 | tmp | ? | now active
|-----+-----+-----+-----|

- ▶ **Caller function** main() and **Callee function** swap()
- ▶ The call **pushes** a stack frame onto the **function call stack**
- ▶ Frame has space for All Callee parameters/locals vars
- ▶ Caller tracks where it left off to resume later
- ▶ Caller copies values to Callee frame for parameters
- ▶ Callee begins executing at its first instruction

Answers: Function Call Stack: Returning from swap()

```
9: int main(...){  
10:    int x = 42;  
11:    int y = 31;  
12:    swap(x, y);  
+-->13:   printf("%d %d\n",x,y);  
| 14:    return 0;  
| 15: }  
|  
^ 18: void swap(int a, int b){  
| 19:    int tmp = a;  
| 20:    a = b;  
| 21:    b = tmp;  
+-<22:   return;  
23: }
```

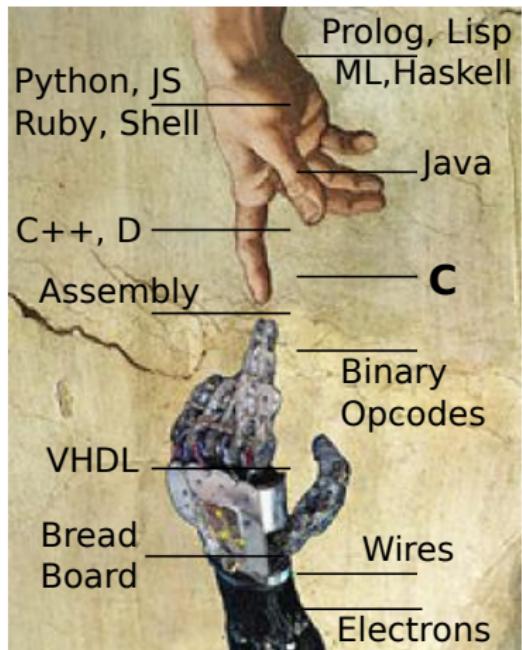
STACK: Callee swap() returning				
FRAME	ADDR	SYM	VALUE	
main()	#2048	x	42	inactive
line:12	#2044	y	31	
swap()	#2040	a	31	about to
line:22	#2036	b	42	return
	#2032	tmp	42	

STACK: Caller main() gets control back				
FRAME	ADDR	SYM	VALUE	
main()	#2048	x	42	now
line:13	#2044	y	31	active

- ▶ On finishing, Callee stack frame **pops** off, Control returns to Caller which resumes executing next instruction
- ▶ Callee may pass a return value to Caller but otherwise does not directly affect Caller stack frame on return
- ▶ `swap()` does NOT swap the variables `x,y` in `main()`, only its own local variables `a,b`

Motivation for C

Pure Abstraction



Bare Metal

[Source](#)

If this were Java, Python, many others, discussion would be over:

- ▶ Provide many safety and convenience features
- ▶ Insulate programmer from hardware for ease of use

C presents many CPU capabilities directly

- ▶ Very few safety features
- ▶ Little between programmer and hardware

You just have to know C. Why?

*Because for all practical purposes, every computer in the world you'll ever use is a **von Neumann machine**, and C is a lightweight, expressive syntax for the von Neumann machine's capabilities.*

—Steve Yegge, [Tour de Babel](#)

Von Neumann Machine Architecture ([Wikipedia](#))

Processing

- ▶ Wires/gates that accomplish fundamental ops
- ▶ +, -, *, AND, OR, move, copy, shift, etc.
- ▶ Ops act on contents of memory cells to change them

Control

- ▶ Memory address of next instruction to execute
- ▶ After executing, move ahead one unless instruction was to jump elsewhere

Memory

- ▶ Giant array of bits/bytes so **everything** is represented as 1's and 0's, including instructions
- ▶ Memory cell contents accessible by address number

Input/Output

- ▶ Allows humans to interpret what is happening
- ▶ Often special memory locations for screen and keyboard

Wait, these items seem kind of familiar...

Exercise: C allows direct use of memory cell addresses

Syntax	Meaning
<code>&x</code>	Address of: memory address of variable x
<code>int *a</code>	Pointer Variable: a stores a memory address
<code>*a</code>	Dereference: get/set the value pointed to by a

Where/how are these used in the code below?

```
1 // swap_pointer.c: swaps values using a function with pointer arguments.
2
3 #include <stdio.h>           // declare existence printf()
4 void swap_ptr(int *a, int *b); // function exists, defined below main
5
6 int main(int argc, char *argv[]){ // ENTRY POINT: start executing in main()
7     int x = 42;
8     int y = 31;
9     swap_ptr(&x, &y);          // call swap() with addresses of x/y
10    printf("x: %d y: %d\n",x,y); // print the values of x and y
11    return 0;
12 }
13
14 // Function to swap contents of two memory cells
15 void swap_ptr(int *a, int *b){ // a/b are addresses of memory cells
16     int tmp = *a;             // go to address a, copy value int tmp
17     *a = *b;                 // copy val at addr in b to addr in a
18     *b = tmp;                // copy temp into address in b
19     return;
20 }
```

Swapping with Pointers/Addresses: Call Stack

```
9: int main(...){          STACK: Caller main(), prior to swap()
10:   int x = 42;           | FRAME    | ADDR   | NAME  | VALUE  |
11:   int y = 31;           |-----+-----+-----+-----|
+-<12:   swap_ptr(&x, &y); | main()   | #2048  | x     | 42    |
| 13:   printf("%d %d\n",x,y); | line:12  | #2044  | y     | 31    |
| 14:   return 0;          |-----+-----+-----+-----|
V 15: }
|
| 18: void swap_ptr(int *a,int *b){ STACK: Callee swap() takes control
+->19:   int tmp = *a;   | FRAME    | ADDR   | NAME  | VALUE  |
20:   *a = *b;             |-----+-----+-----+-----|
21:   *b = tmp;            | main()   | #2048  | x     | 42    | <-+
22:   return;              | line:12  | #2044  | y     | 31    | <-|+
23: }
```

The code shows the state of the call stack before and after the swap operation. In the caller's frame (main()), the variable 'x' has address #2048 and value 42, while 'y' has address #2044 and value 31. In the callee's frame (swap()), the parameter 'a' has address #2048 and value 42, and 'b' has address #2044 and value 31. The local variable 'tmp' is also present at address #2024.

- ▶ Syntax `&x` reads “Address of cell associated with `x`” or just “Address of `x`”. Ampersand `&` is the address-of operator.
- ▶ Swap takes `int *a`: **pointer** to integer / memory address
- ▶ Values associated with `a/b` are the addresses of other cells

Swapping with Pointers/Addresses: Dereference/Use

```
9: int main(...){  
10:    int x = 42;  
11:    int y = 31;  
12:    swap_ptr(&x, &y);  
13:    printf("%d %d\n",x,y);  
14:    return 0;  
15: }  
  
18: void swap_ptr(int *a,int *b){  
19:    int tmp = *a; // copy val at #2048 to #2024  
>20:    *a = *b;  
21:    *b = tmp;  
22:    return;  
23: }
```

LINE	executed:	tmp	gets	42
	FRAME	ADDR	NAME	VALUE
19	main()	#2048	x	42 <-+
	line:12	#2044	y	31 <- +
	swap_ptr	#2036	a	#2048 --+
	line:20	#2028	b	#2044 ---+
		#2024	tmp	?->42

- ▶ Syntax `*a` reads “Dereference and operate on the cell pointed to by `a`” or just “Deref `a`”
- ▶ Line 19 dereferences via `*` operator:
 - ▶ Cell #2036 (`a`) contains address #2048,
 - ▶ Copy contents of #2048 (42) into #2024 (`tmp`)

Swapping with Pointers/Addresses: Dereference/Assign

```
9: int main(...){                                LINE 20 executed: alters x using a
10:    int x = 42;                            | FRAME   | ADDR   | NAME   | VALUE   |
11:    int y = 31;                            |-----+-----+-----+-----|
12:    swap_ptr(&x, &y);                     | main()  | #2048  | x      | 42->31 |<-+|
13:    printf("%d %d\n",x,y);                | line:12 | #2044  | y      | 31     |<-|+
14:    return 0;                           |-----+-----+-----+-----|   ||
15: }
16:
17: void swap_ptr(int *a,int *b){               |           | #2024  | tmp    | 42   |
18:    int tmp = *a;                         |           | #2036  | a      | #2048 |---+
19:    *a = *b;      // copy val at #2044 (31) to #2048 (was 42)
20:    *b = tmp;
21:    return;
22:
23: }
```

- ▶ Pointer Deref on Right Side **fetches** a value from a pointer location
- ▶ Pointer Deref on Left Side **stores** a value at a pointer location
- ▶ Line 20: Deref on both Left and right side of assignment
 - ▶ a and b contain pointers, not changed
 - ▶ x and y are pointed at, can change

Swapping with Pointers/Addresses: Deref 2

```
9: int main(...){                                LINE 21 executed: alters y using b
10:    int x = 42;                            | FRAME   | ADDR   | NAME   | VALUE   |
11:    int y = 31;                            |-----+-----+-----+-----+
12:    swap_ptr(&x, &y);                      | main()  | #2048  | x     | 31    |<-+
13:    printf("%d %d\n",x,y);                | line:12 | #2044  | y     | 31->42 |<-|+
14:    return 0;                           |-----+-----+-----+-----|  ||
15: }
16:                                     | swap_ptr| #2036  | a     | #2048  |--+
17:                                     | line:22 | #2028  | b     | #2044  |---+
18: void swap_ptr(int *a,int *b){           |         | #2024  | tmp   | 42   |
19:    int tmp = *a;
20:    *a = *b;
21:    *b = tmp;      // copy val at #2024 (42) to #2044 (was 31)
>22:    return;
23: }
```

- ▶ Line 21: dereference on left-hand side

*b = ...

stores new value at address #2044

- ▶ Use of variable **bare name** always retrieves value it that cell
 - ▶ tmp retrieves an int like 42
 - ▶ a retrieves a pointer like #2048

Swapping with Pointers/Addresses: Returning

```
9: int main(...){  
10:    int x = 42;  
11:    int y = 31;  
12:    swap_ptr(&x, &y);  
+-->13:   printf("%d %d\n",x,y);  
| 14:    return 0;  
| 15: }  
|  
| 18: void swap_ptr(int *a,int *b){  
| 19:    int tmp = *a;  
| 20:    *a = *b;  
| 21:    *b = tmp;  
+--<22:   return;  
23: }
```

LINE 22: prior to return

FRAME	ADDR	NAME	VALUE
main()	#2048	x	31
line:12	#2044	y	42
swap_ptr	#2036	a	#2048
line:22	#2028	b	#2044
	#2024	tmp	42

LINE 12 finished/return pops frame

FRAME	ADDR	NAME	VALUE
main()	#2048	x	31
line:13	#2044	y	42

- ▶ swap_ptr() finished so frame pops off
- ▶ Variables x,y in main() have changed due to use of references to them.

Aside: Star/Asterisk * has 3 uses in C

1. Multiply numbers as in

```
w = c*d;
```

2. **Declare** a pointer variable as in

```
int *x; // pointer to integer(s)
int b=4;
x = &b; // point x at b
int **r; // pointer to int pointer(s)
```

3. **Dereference** a pointer variable as in

```
int p = *x; // x must be an int pointer
             // retrieve contents at address
```

Three different context sensitive meanings for the same symbol makes * hard on humans to parse, a regrettable decision by Ritchie.

```
int z = *x * *y + *(p+2); // standard, 'unambiguous' C
The duck is ready to eat. // English is more ambiguous
```

Some Common Examples and Errors

- ▶ Learning syntax and semantics of pointers requires some practice, get started with below examples
- ▶ Won't go through these in much detail YET but over next couple weeks will discuss at length

```
// pointer_examples.c          // 4: proper pointer deref
// 1: proper pointer assignment
int a1 = 11;
int *p1 = &a1;    // cool
int b1 = 55;
p1 = &b1;        // cool

// 2: improper pointer assignment
int a2 = 13;
int *p2 = a2;    // ERROR

// 3: proper pointer copying
int a3 = 15;
int *p3 = &a3;
int *q3 = p3;    // cool

// 4: proper pointer deref
int a4 = 17;
int *p4 = &a4;
int b4 = *p4;    // cool

// 5: improper int assign (no deref)
int a5 = 19;
int *p5 = &a5;
int b5 = p5;    // ERROR
```

Important Principle: Non-local Changes

- ▶ Pointers allow functions to change variables associated with other running functions
- ▶ Common beginner example: `scanf()` family which is used to read values from terminal or files
- ▶ Snippet from `scanf_demo.c`

```
1 int main(...){  
2     int num = -1;  
3     scanf("%d", &num); // addr  
4     printf("%d\n", num); // val  
5 }
```

- ▶ See `scanf_error.c`: forgetting & yields great badness

`scanf()` called

FRAME	ADDR	NAME	VALUE
main():3	#2500	num	-1
scanf()	#2492	fmt	#400
	#2484	arg1	#2500

`scanf()` changes contents of #2500

FRAME	ADDR	NAME	VALUE
main():3	#2500	num	5
scanf()	#2492	fmt	#400
	#2484	arg1	#2500

`scanf()` returns

FRAME	ADDR	NAME	VALUE
main():4	#2500	num	5

Uncle Ben Said it Best...



- ▶ Pointers allow any line of C programs to modify any of its data
- ▶ A BLESSING: fine control of memory → efficiency, machine's true capability
- ▶ A CURSE: opens up many errors not possible in Java/Python which restrict use of memory

1972 - Dennis Ritchie invents a powerful gun that shoots both forward and backward simultaneously. Not satisfied with the number of deaths and permanent maimings from that invention he invents C and Unix.

- A Brief, Incomplete, and Mostly Wrong History of Programming Languages

All of these apply to our context..

Beneath the C

C is “high-level” as it abstracts away from a real machine. It must be translated to lower forms to be executed on a real machine.

Assembly Language

- ▶ Specific to each CPU architecture (Intel, etc)
- ▶ Still “human readable” but fairly directly translated to binary using Assemblers

INTEL x86-64 ASSEMBLY

```
cmpl    $1, %ecx
jle     .END
movl    $2, %esi
movl    %ecx,%eax
cqto
idivl   %esi
cmpl    $1,%edx
jne     .EVEN
```

Binary Opcodes

- ▶ 1's and 0's, represent the digital signal of the machine
- ▶ Codes corresponds to instructions directly understood by processor

HEXADECIMAL/BINARY OPCODES

1124:	83 f9 01
1127:	7e 1e = 0111 1110 0001 1110
1129:	be 02 00 00 00
112e:	89 c8
1130:	48 99
1132:	f7 fe
1134:	83 fa 01
1137:	75 07

Looks like **fun**, right? You bet it is! Assembly coding is 6 weeks away...

CMSC216: Course Goals

- ▶ Basic proficiency at C programming
- ▶ Knowledge of running programs in physical memory including the stack, heap, global, and text areas of memory
- ▶ Understanding of the essential elements of assembly languages
- ▶ Knowledge of the correspondence between high-level program constructs.
- ▶ Ability to use a symbolic debugger
- ▶ Basic understanding of how data is encoded in binary
- ▶ Understanding the process abstraction of running programs, ability to create and manipulate processes
- ▶ Basic understanding of execution threads, their relation to processes, the ability to create and manipulate threads