

# CMSC216: Assembly Basics and x86-64

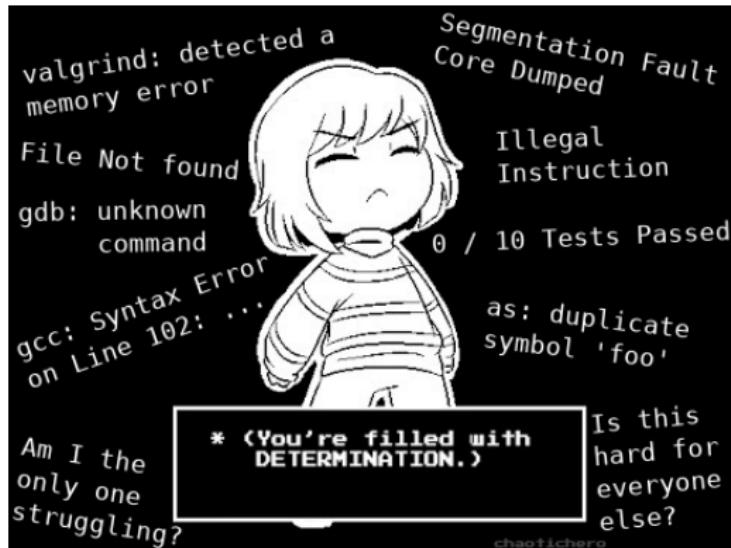
Chris Kauffman

*Last Updated:*  
*Sat Feb 28 12:06:26 PM EST 2026*

# Logistics

## Announcements

# Don't Give Up, Stay Determined!



- ▶ If Project 1 / Exam 1 went awesome, count yourself lucky
- ▶ If things did not go well, **Don't Give Up**
- ▶ Spend some time contemplating **why** things didn't go well, talk to course staff about it, learn from mistakes
- ▶ There is a LOT of semester left and time to recover from a bad start

## Wrapping Up Integer Affairs

Pre-exam one slides contain info on byte-ordering (little / big endian) and bit-wise operations to be discussed here prior to assembly.

# The Many Assembly Languages

- ▶ Most **microprocessors** are created to understand a **binary machine language**
- ▶ Machine Language provides means to manipulate internal memory, perform arithmetic, etc.
- ▶ The Machine Language of one processor is **not understood** by other processors

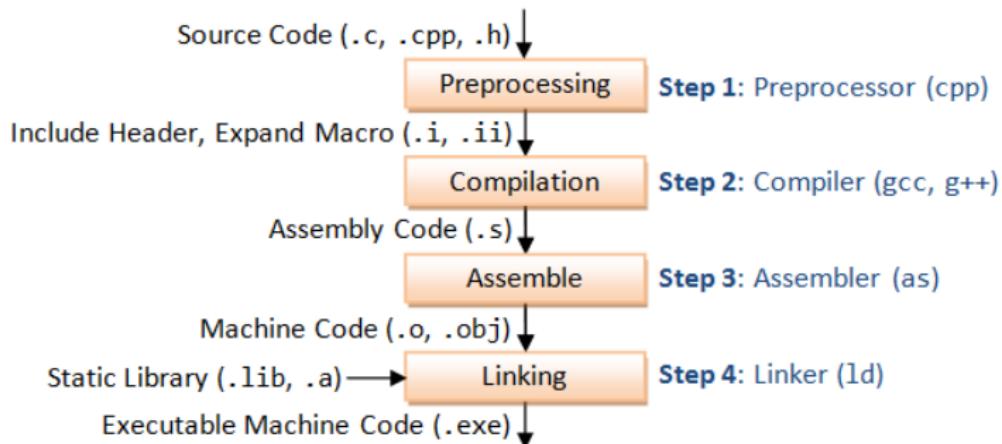
## MOS Technology 6502

- ▶ 8-bit operations, limited addressable memory, **1 general purpose register**, powered notable gaming systems in the 1980s
- ▶ Apple IIe, Atari 2600, Commodore
- ▶ Nintendo Entertainment System / Famicom

## IBM Cell Microprocessor

- ▶ Developed in early 2000s, 64-bit, many cores (execution elements), many registers (32 on the PPE), large addressable space, fast multimedia performance, is a **pain** to program
- ▶ Playstation 3 and Blue Gene Supercomputer

# Assemblers and Compilers



- ▶ **Compiler**: chain of tools that translate high level languages to lower ones, may perform optimizations
- ▶ **Assembler**: translates text description of the machine code to binary, formats for execution by processor, late compiler stage
- ▶ **Consequence**: The compiler can **generate assembly code**
- ▶ Generated assembly is a pain to read but is often quite fast
- ▶ **Consequence**: A compiler on an Intel chip can generate assembly code for a different processor, **cross compiling**

# Our focus: The x86-64 Assembly Language

- ▶ x86-64 Targets Intel/AMD chips with 64-bit word size  
*Reminder: 64-bit “word size” ≈ size of pointers/addresses*
- ▶ Lineage of x86 family
  - ▶ 1970s: 16-bit systems like Intel 8086
  - ▶ 1990s: IA32 (Intel 32-bit systems like 80386 and 80486)
  - ▶ 2000s: x86-64 (64-bit extension by AMD)
- ▶ x86-64 is backwards compatibility, consequently much cruft
  - ▶ Can run compiled code from the 70's / 80's on modern processors without much trouble BUT means 50-year-old instructions must be preserved
  - ▶ x86-64 is not the assembly language you would design from scratch today, it's the assembly you have to code against
  - ▶ RISC-V is a new assembly language that is “clean” as it has no history to support (and few CPUs run it)
- ▶ **Warning:** Lots of information available on the web for Intel assembly programming **BUT** some of it is dated, IA32 info which may not work on 64-bit systems

# x86-64 Assembly Language Syntax(es)

- ▶ Different assemblers understand different syntaxes for the same assembly language
- ▶ GCC use the GNU Assembler (GAS, command 'as file.s')
- ▶ GAS and Textbook favor AT&T syntax so **we will too**
- ▶ NASM assembler favors Intel, may see this online

## AT&T Syntax (Our Focus)

`multstore:`

```
pushq %rbx
movq %rdx, %rbx
call mult2@PLT
movq %rax, (%rbx)
popq %rbx
ret
```

- ▶ Use of % to indicate registers
- ▶ Use of q/l/w/b to indicate 64 / 32 / 16 / 8-bit operands

## Intel Syntax

`multstore:`

```
push rbx
mov rbx, rdx
call mult2@PLT
mov QWORD PTR [rbx], rax
pop rbx
ret
```

- ▶ Register names are bare
- ▶ Use of QWORD etc. to indicate operand size

## Generating Assembly from C Code

- ▶ `gcc -S file.c` will stop compilation at assembly generation
- ▶ Leaves assembly code in `file.s`
  - ▶ `file.s` and `file.S` conventionally assembly code though sometimes `file.asm` is used
- ▶ By default, compiler generates code that is often difficult for humans to interpret, may include re-arrangements, “conservative” compatibility assembly, etc. increasing size of assembly considerably
- ▶ `gcc -Og file.c`: optimize for debugging, generally makes it easier to read generated assembly, aligns somewhat more closely to C code

# Example of Generating Assembly from C

```
>> cat exchange.c                                # show C file to be translated
// exchange.c: sample C function
// to compile to assembly
long exchange(long *xp, long y){
    long x = *xp;
    *xp = y;
    return x;
}

>> gcc -Og -S exchange.c                         # Compile to show assembly
# -Og: debugging level optimization
# -S: only output assembly

>> cat exchange.s                               # show assembly output
      .file  "exchange.c"
      .text
      .globl exchange
      .type   exchange, @function
exchange:                                     # beginning of exchange function
.LFB0:
      .cfi_startproc
      movq    (%rdi), %rax
      movq    %rsi, (%rdi)
      ret
      .cfi_endproc
.LFE0:
      .size   exchange, .-exchange
      .ident  "GCC: (GNU) 11.1.0"
      .section .note.GNU-stack,"",@progbits
```

```
gcc -Og -S mstore.c
```

```
> cat mstore.c                                # show a C file
long mult2(long a, long b);
void multstore(long x, long y, long *dest){
    long t = mult2(x, y);
    *dest = t;
}

> gcc -Og -S mstore.c                         # Compile to show assembly
# -Og: debugging level optimization
# -S: only output assembly

# show assembly output

> cat mstore.s
      .file  "mstore.c"
      .text
      .globl  multstore
      .type   multstore, @function
multstore:                                # function symbol for linking
.LFB0:                                     # beginning of mulstore function
      .cfi_startproc
      pushq   %rbx
      .cfi_def_cfa_offset 16
      .cfi_offset 3, -16
      movq   %rdx, %rbx
      call   mult2@PLT
      movq   %rax, (%rbx)
      popq   %rbx
      .cfi_def_cfa_offset 8
      ret
      .cfi_endproc
```

# Every Programming Language

Look for the following as it should almost always be there

- Comments
- Statements/Expressions
- Variable Types
- Assignment
- Basic Input/Output
- Function Declarations
- Conditionals (if-else)
- Iteration (loops)
- Aggregate data (arrays, structs, objects, etc)
- Library System

## Exercise: Examine col\_simple\_asm.s

Take a simple sample problem to demonstrate assembly:

*Computes Collatz Sequence starting at n=10:*

*if n is ODD n=n\*3+1; else n=n/2.*

*Return the number of steps to converge to 1 as the **return code** from main()*

The following codes solve this problem

Code	Notes
col_simple_asm.s	Hand-coded assembly for obvious algorithm Straight-forward reading
col_unsigned.c	Unsigned C version Generated assembly is reasonably readable
col_signed.c	Signed C version Generated assembly is ... interesting

- ▶ Kauffman will Compile/Run code
- ▶ Students should **study the code and predict what lines do**
- ▶ Illustrate tricks associated with gdb and assembly

## Exercise: col\_simple\_asm.s

```
1 ##### Compute Collatz sequence starting at 10 in assembly.
2 .text
3 .global main
4 main:
5     movl    $0, %r8d      # int steps = 0;
6     movl    $10, %ecx     # int n = 10;
7 .LOOP:
8     cmpl    $1, %ecx     # while(n > 1){ // immediate must be first
9     jle     .END         #   n <= 1 exit loop
10    movl    $2, %esi      #   divisor in esi
11    movl    %ecx,%eax    #   prep for division: must use edx:eax
12    cltd
13    idivl   %esi        #   divide edx:eax by esi
14    #   eax has quotient, edx remainder
15    cmpl    $1,%edx      #   if(n % 2 == 1) {
16    jne     .EVEN        #       not equal, go to even case
17 .ODD:
18    imull   $3, %ecx     #       n = n * 3
19    incl    %ecx         #       n = n + 1 OR n++
20    jmp     .UPDATE
21 .EVEN:
22    sarl    $1,%ecx      #       else{
23 .UPDATE:
24    incl    %r8d         #       n = n / 2; via right shift
25    jmp     .LOOP        #       }
26 .END:
27    movl    %r8d, %eax    # r8d is steps, move to eax for return value
28    ret
29
```

## Answers: x86-64 Assembly Basics for AT&T Syntax

- ▶ *Comments* are one-liners starting with #
- ▶ *Statements*: each line does ONE thing, frequently text representation of an assembly instruction

```
    movq    %rdx, %rbx      # move rdx register to rbx
```

- ▶ Assembler directives and labels are also possible:

```
.global multstore          # notify linker of location multstore
multstore:                 # label beginning of multstore section
    blah blah blah         # instructions in this section
```

- ▶ *Variables*: mainly **registers**, also memory ref'd by registers maybe some named global locations
- ▶ *Assignment*: instructions like movX that put bits into registers and memory
- ▶ *Conditionals/Iteration*: assembly instructions that jump to code locations
- ▶ *Functions*: code locations that are **labeled** and global
- ▶ *Aggregate data*: none, use the stack/multiple registers
- ▶ *Library System*: link to other code

## So what are these Registers?

- ▶ Memory locations directly wired to the CPU
- ▶ Usually *very* fast to access, faster than **main memory**
- ▶ Most instructions involve registers, access or change reg val

### Example: Adding Together Integers

- ▶ Ensure registers have desired values in them
- ▶ Issue an addX instruction involving the two registers
- ▶ Result will be stored in a register

```
addl %eax, %ebx  
# add ints in eax and ebx, store result in ebx
```

```
addq %rcx, %rdx  
# add longs in rcx and rdx, store result in rdx
```

- ▶ Note instruction and register names indicate whether 32-bit int or 64-bit long are being added

# x86-64 “General Purpose” Registers

Many “general purpose” registers have special purposes and conventions associated such as

- ▶ Return Value:  
%rax / %eax / %ax
- ▶ Function Args 1 to 6:  
%rdi, %rsi, %rdx,  
%rcx, %r8, %r9
- ▶ Stack Pointer (top of stack): %rsp
- ▶ Old Code Base Pointer:  
**%rbp, historically start of current stack frame but is not used that way in modern codes**

Note: There are also Special Registers like %rip and %eflags which we will discuss later.

64-bit	32-bit	16-bit	8-bit	Notes
%rax	%eax	%ax	%al	Return Val
%rbx	%ebx	%bx	%bl	
%rcx	%ecx	%cx	%cl	Arg 4
%rdx	%edx	%dx	%dl	Arg 3
%rsi	%esi	%si	%sil	Arg 2
%rdi	%edi	%di	%dil	Arg 1
%rsp	%esp	%sp	%spl	Stack Ptr
%rbp	%ebp	%bp	%bpl	Base Ptr?
%r8	%r8d	%r8w	%r8b	Arg 5
%r9	%r9d	%r9w	%r9b	Arg 6
%r10	%r10d	%r10w	%r10b	
%r11	%r11d	%r11w	%r11b	
%r12	%r12d	%r12w	%r12b	
%r13	%r13d	%r13w	%r13b	
%r14	%r14d	%r14w	%r14b	
%r15	%r15d	%r15w	%r15b	
<b>Caller Save:</b>		Restore after calling func		
<b>Callee Save:</b>		Restore before returning		

# Register Naming Conventions

- ▶ AT&T syntax identifies registers with prefix %
- ▶ Naming convention is a historical artifact
- ▶ Originally 16-bit architectures in x86 had
  - ▶ General registers ax,bx,cx,dx,
  - ▶ Special Registers si,di,sp,bp
- ▶ Extended to 32-bit: eax,ebx,...,esi,edi,...
- ▶ Grew again to 64-bit: rax,rbx,...,rsi,rdi,...
- ▶ Added Eight 64-bit regs r8,r9,...,r14,r15 with 32-bit portion r8d,r9d,..., 16-bit r8w,r9w..., etc.
- ▶ Instructions must match registers sizes:

```
addw %ax, %bx #      word (16-bit)
addl %eax, %ebx # long word (32-bit)
addq %rax, %rbx # quad-word (64-bit)
```
- ▶ When hand-coding assembly, easy to mess this up, assembler will error out

## Hello World in x86-64 Assembly : Not that Easy

- ▶ Non-trivial in assembly because **output is involved**
  - ▶ Try writing helloworld.c without printf()
- ▶ Output is the business of the **operating system**, always a request to the almighty OS to put something somewhere
  - ▶ **Library call:** printf("hello"); mangles some bits but eventually results with a ...
  - ▶ **System call:** Unix system call directly implemented in the OS **kernel**, puts bytes into files / onto screen as in

```
write(1, buf, 5); // file 1 is screen output
```

This gives us several options for hello world in assembly:

1. `hello_printf64.s`: via calling `printf()` which means the C standard library must be (painfully) linked
2. `hello64.s` via direct system `write()` call which means no external libraries are needed: OS knows how to write to files/screen. Use the 64-bit Linux calling convention.
3. `hello32.s` via direct system call using the older 32 bit Linux calling convention which “traps” to the operating system.

## (Optional): The OS Privilege: System Calls

- ▶ Most interactions with the outside world happen via Operating System Calls (or just “system calls”)
- ▶ User programs indicate what service they want performed by the OS via making system calls
- ▶ System Calls differ for each language/OS combination
  - ▶ x86-64 Linux: set %rax to system call number, set other args in registers, issue `syscall`
  - ▶ IA32 Linux: set %eax to system call number, set other args in registers, issue an **interrupt**
  - ▶ C Code on Unix: make system calls via `write()`, `read()` and others (studied in CSCI 4061)
  - ▶ Tables of Linux System Call Numbers
    - ▶ [64-bit \(335 calls\)](#)
    - ▶ [32-bit \(190 calls\)](#)
  - ▶ Mac OS X: very similar to the above (it's a Unix)
  - ▶ Windows: use OS wrapper functions
- ▶ OS executes **privileged** code that can manipulate any part of memory, touch internal data structures corresponding to files, do other fun stuff discussed in OS courses

# Basic Instruction Classes

- ▶ **Remember:** Goal is to understand assembly as a *target* for higher languages, not become expert “assemblists”
- ▶ Means we won’t hit all 4,834 pages of the Intel x86-64 Manual
- ▶ Brown University’s [x64 Cheat Sheet](#) has a good overview
- ▶ [x86 Assembly Guide from Yale](#) is also good but is limited to 32-bit coverage

Kind	Assembly Instructions
<i>Fundamentals</i>	
- Memory Movement	mov
- Stack manipulation	push, pop
- Addressing modes	(%eax), 12(%eax,%ebx)...
<i>Arithmetic/Logic</i>	
- Arithmetic	add, sub, mul, div, lea
- Bitwise Logical	and, or, xor, not
- Bitwise Shifts	sal, sar, shr
<i>Control Flow</i>	
- Compare / Test	cmp, test
- Set on result	set
- Jumps (Un)Conditional	jmp, je, jne, jl, jg,...
- Conditional Movement	cmove, cmovg,...
<i>Procedure Calls</i>	
- Stack manipulation	push, pop
- Call/Return	call, ret
- System Calls	syscall
<i>Floating Point Ops</i>	
- FP Reg Movement	vmov
- Conversions	vcvts
- Arithmetic	vadd, vsub, vmul, vdiv
- Extras	vmins, vmaxs, sqrts

# Data Movement: movX instruction

```
movX SOURCE, DEST      # move/copy source value to dest
```

## Overview

- ▶ Moves data...
  - ▶ Reg to Reg
  - ▶ Mem to Reg
  - ▶ Reg to Mem
  - ▶ Imm to ...
- ▶ Reg: register
- ▶ Mem: main memory
- ▶ Imm: “immediate” value (constant) specified like
  - ▶ \$21 : decimal
  - ▶ \$0x2f9a : hexadecimal
  - ▶ **NOT** 1234 (mem adder)
- ▶ More info on operands next

## Examples

```
## 64-bit quadword moves
movq $4, %rbx      # rbx = 4;
movq %rbx,%rax    # rax = rbx;
movq $10, (%rcx)   # *rcx = 10;
```

```
## 32-bit longword moves
movl $4, %ebx      # ebx = 4;
movl %ebx,%eax    # eax = ebx;
movl $10, (%rcx)   # *rcx = 10;
```

## Note variations

- ▶ movq for 64-bit (8-byte)
- ▶ movl for 32-bit (4-byte)
- ▶ movw for 16-bit (2-byte)
- ▶ movb for 8-bit (1-byte)

# Operands and Addressing Modes

In many instructions like `movX`, operands can have a variety of forms called **addressing modes**, may include constants and memory addresses

Style	Address Mode	C-like	Notes
\$21	immediate	21	value of constant like 21
\$0xD2			or 0xD2 = 210
%rax	register	rax	to/from register contents
(%rax)	indirect	*rax	reg holds memory address, deref
8(%rax)	displaced	*(rax+2)	base plus constant offset, often
4(%rdx)		rdx->field	used for struct field derefs
(%rax,%rbx)	indexed	*(rax+rbx) char_arr[rbx]	base plus offset in given reg actual value of rbx is used, NOT multiplied by sizeof()
(%rax,%rbx,4)	scaled index	rax[rbx]	like array access with sizeof(..)=4
(%rax,%rbx,8)		rax[rbx]	"" with sizeof(..)=8
1024	absolute	...	Absolute address #1024 <b>Rarely used</b>

## Exercise: Show movX Instruction Execution

### Code movX\_exercise.s

```
movl $16, %eax
movl $20, %ebx
movq $24, %rbx
## POS A

movl %eax,%ebx
movq %rcx,%rax
## POS B

movq $45,(%rdx)
movl $55,16(%rdx)
## POS C

movq $65,(%rcx,%rbx)
movq $3,%rbx
movq $75,(%rcx,%rbx,8)
## POS D
```

### Registers/Memory

#### INITIAL

REG	%rax	0
	%rbx	0
	%rcx	#1024
	%rdx	#1032
MEM	#1024	35
	#1032	25
	#1040	15
	#1048	5

#### Lookup...

May need to look up addressing conventions for things like...

```
movX %y,%x    # reg y to reg x
movX $5,(%x)  # 5 to address in %x
```

## Answers Part 1/2: movX Instruction Execution

INITIAL		## POS A		## POS B	
REG	VALUE	REG	VALUE	REG	VALUE
%rax	0	%rax	16	%rax	#1024
%rbx	0	%rbx	24	%rbx	16
%rcx	#1024	%rcx	#1024	%rcx	#1024
%rdx	#1032	%rdx	#1032	%rdx	#1032
MEM	VALUE	MEM	VALUE	MEM	VALUE
#1024	35	#1024	35	#1024	35
#1032	25	#1032	25	#1032	25
#1040	15	#1040	15	#1040	15
#1048	5	#1048	5	#1048	5

#WARNING!: On 64-bit systems, ALWAYS use a 64-bit reg name like %rdx and movq to copy memory addresses; using smaller name like %edx will miss half the memory addressing leading to major memory problems

## Answers Part 2/2: movX Instruction Execution

			movq \$65,(%rcx,%rbx)	
			#1024+16 = #1040	
movl %eax,%ebx	movq \$45,(%rdx)	#1032	movq \$3,%rbx	
movq %rcx,%rax #!	movq \$55,16(%rdx)	16+#1032=#1048	movq \$75,(%rcx,%rbx,8)	#1024 + 3*8 = #1048
## POS B	## POS C		## POS D	
-----+-----	-----+-----		-----+-----	
REG   VALUE	REG   VALUE		REG   VALUE	
%rax   #1024	%rax   #1024		%rax   #1024	
%rbx   16	%rbx   16		%rbx   3	
%rcx   #1024	%rcx   #1024		%rcx   #1024	
%rdx   #1032	%rdx   #1032		%rdx   #1032	
-----+-----	-----+-----		-----+-----	
MEM   VALUE	MEM   VALUE		MEM   VALUE	
#1024   35	#1024   35		#1024   35	
#1032   25	#1032   45		#1032   45	
#1040   15	#1040   15		#1040   65	
#1048   5	#1048   55		#1048   75	
-----+-----	-----+-----		-----+-----	

# gdb Assembly: Examining Memory

gdb commands print and x allow one to print/examine memory memory of interest. Try on movX\_exercises.s

```
(gdb) tui enable          # TUI mode
(gdb) layout asm          # assembly mode
(gdb) layout reg          # show registers
(gdb) stepi               # step forward by single Instruction
(gdb) print $rax          # print register rax
(gdb) print *($rdx)        # print memory pointed to by rdx
(gdb) print (char *) $rdx  # print as a string (null terminated)
(gdb) x $r8                # examine memory at address in r8
(gdb) x/3d $r8             # same but print as 3 4-byte decimals
(gdb) x/6g $r8             # same but print as 6 8-byte decimals
(gdb) x/s $r8              # print as a string (null terminated)
(gdb) print *((int*) $rsp)  # print top int on stack (4 bytes)
(gdb) x/4d $rsp            # print top 4 stack vars as ints
(gdb) x/4x $rsp            # print top 4 stack vars as ints in hex
```

Many of these tricks are needed to debug assembly.

# Register Size and Data Movement

- ▶ %rax is 64-bit register, %eax is its lower 32 bits
- ▶ Data movement involving small registers **may NOT overwrite** higher bits in extended register
- ▶ Moving data to low 32-bit regs automatically zeros high 32-bits

```
movabsq $0x1122334455667788, %rax    # 8 bytes to %rax
movl $0xAABBCCDD, %eax                 # 4 bytes to %eax
## %rax is now 0x00000000AABBCCDD
```

- ▶ Moving data to other small regs DOES NOT ALTER high bits
- ▶ Gives rise to two other families of movement instructions for moving little registers (X) to big (Y) registers, see `movz_examples.s`

```
## movzXY move zero extend, movsXY move sign extend
movabsq $0x112233445566AABB,%rdx
movzqw %dx,%rax          # %rax is 0x000000000000AABB
movswq %dx,%rax          # %rax is 0xFFFFFFFFFFFFAABB
```

## Exercise: movX differences in Main Memory

Instr	# bytes
movb	1 byte
movw	2 bytes
movl	4 bytes
movq	8 bytes

Show the result of each of the following copies to main memory in sequence.

```
movl    %eax,    (%rsi) #1
movq    %rax,    (%rsi) #2
movb    %cl,     (%rsi) #3
movw    %cx,     2(%rsi) #4
movl    %ecx,    4(%rsi) #5
movw    4(%rsi), %ax  #6
```

INITIAL		
REG		
rax	0x00000000DDCCBBAA	
rcx	0x00000000000000FFEE	
rsi		#1024
MEM		
#1024		0x00
#1025		0x11
#1026		0x22
#1027		0x33
#1028		0x44
#1029		0x55
#1030		0x66
#1031		0x77
#1032		0x88
#1033		0x99

## Answers: movX to Main Memory 1/2

		movl %eax, (%rsi) #1	4 bytes	rax -> #1024
REG		movq %rax, (%rsi) #2	8 bytes	rax -> #1024
rax	0x00000000DDCCBAA	movb %cl, (%rsi) #3	1 byte	rcx -> #1024
rcx	0x000000000000FFEE	movw %cx, 2(%rsi) #4	2 bytes	rcx -> #1026
rsi	#1024	movl %ecx, 4(%rsi) #5	4 bytes	rcx -> #1028
		movw 4(%rsi), %ax #6	2 bytes	#1024 -> rx

INITIAL	#1 movl %eax, (%rsi)		#2 movq %rax, (%rsi)		#3 movb %cl, (%rsi)	
MEM	MEM	MEM	MEM	MEM	MEM	MEM
#1024	0x00	#1024	0xAA	#1024	0xAA	#1024
#1025	0x11	#1025	0xBB	#1025	0xBB	#1025
#1026	0x22	#1026	0xCC	#1026	0xCC	#1026
#1027	0x33	#1027	0xDD	#1027	0xDD	#1027
#1028	0x44	#1028	0x44	#1028	0x00	#1028
#1029	0x55	#1029	0x55	#1029	0x00	#1029
#1030	0x66	#1030	0x66	#1030	0x00	#1030
#1031	0x77	#1031	0x77	#1031	0x00	#1031
#1032	0x88	#1032	0x88	#1032	0x88	#1032
#1033	0x99	#1033	0x99	#1033	0x99	#1033

Answers: movX to Main Memory 2/2

		movl	%eax, (%rsi)	#1	4 bytes	rax -> #1024
REG		movq	%rax, (%rsi)	#2	8 bytes	rax -> #1024
rax	0x000000000DDCCBBAA	movb	%cl, (%rsi)	#3	1 byte	rcx -> #1024
rcx	0x0000000000000FFEE	movw	%cx, 2(%rsi)	#4	2 bytes	rcx -> #1026
rsi	#1024	movl	%ecx, 4(%rsi)	#5	4 bytes	rcx -> #1028
		movw	4(%rsi), %ax	#6	2 bytes	#1024 -> rxax
#3		#4		#5		#6
movb %cl, (%rsi)		movw %cx, 2(%rsi)		movl %ecx, 4(%rsi)		movw 4(%rsx), %ax
MEM		MEM		MEM		MEM
#1024	0xEE	#1024	0xEE	#1024	0xEE	#1024
#1025	0xBB	#1025	0xBB	#1025	0xBB	#1025
#1026	0xCC	#1026	0xEE	#1026	0xEE	#1026
#1027	0xDD	#1027	0xFF	#1027	0xFF	#1027
#1028	0x00	#1028	0x00	#1028	0xEE	#1028
#1029	0x00	#1029	0x00	#1029	0xFF	#1029
#1030	0x00	#1030	0x00	#1030	0x00	#1030
#1031	0x00	#1031	0x00	#1031	0x00	#1031
#1032	0x88	#1032	0x88	#1032	0x88	#1032
#1033	0x99	#1033	0x99	#1033	0x99	#1033

# addX : A Quintessential ALU Instruction

- addX B, A # A = A+B ► Addition represents most 2-operand ALU instructions well
- OPERANDS:**
- addX %reg, %reg
  - addX (%mem), %reg
  - addX %reg, (%mem)
  - addX \$con, %reg
  - addX \$con, (%mem)
- # No mem+mem or con+con ► Second operand A is modified by first operand B, No change to B
- Variety of register, memory, constant combinations honored
- addX has variants for each register size: addq, addl, addw, addb

## EXAMPLES:

addq %rdx, %rcx	# rcx = rcx + rdx
addl %eax, %ebx	# ebx = ebx + eax
addq \$42, %rdx	# rdx = rdx + 42
addl (%rsi), %edi	# edi = edi + *rsi
addw %ax, (%rbx)	# *rbx = *rbx + ax
addq \$55, (%rbx)	# *rbx = *rbx + 55

addl (%rsi,%rax,4),%edi # edi = edi+rsi[rax] (int)

## Optional Exercise: Addition

Show the results of the following addX/movX ops at each of the specified positions

```
addq $1,%rcx      # con + reg
addq %rbx,%rax    # reg + reg
## POS A

addq (%rdx),%rcx  # mem + reg
addq %rbx,(%rdx)  # reg + mem
addq $3,(%rdx)    # con + mem
## POS B

addl $1,(%r8,%r9,4) # con + mem
addl $1,%r9d        # con + reg
addl %eax,(%r8,%r9,4) # reg + mem
addl $1,%r9d        # con + reg
addl (%r8,%r9,4),%eax # mem + reg
## POS C
```

### INITIAL

REGS	
%rax	15
%rbx	20
%rcx	25
%rdx	#1024
%r8	#2048
%r9	0

MEM	
#1024	100
...	...
#2048	200
#2052	300
#2056	400

## Answers: Addition

INITIAL	POS A		POS B		POS C	
REG	REG		REG		REG	
%rax	15	%rax	35	%rax	35	%rax 435
%rbx	20	%rbx	20	%rbx	20	%rbx 20
%rcx	25	%rcx	26	%rcx	126	%rcx 126
%rdx	#1024	%rdx	#1024	%rdx	#1024	%rdx #1024
%r8	#2048	%r8	#2048	%r8	#2048	%r8 #2048
%r9	0	%r9	0	%r9	0	%r9 2
MEM	MEM		MEM		MEM	
#1024	100	#1024	100	#1024	123	#1024 123
...	...	...	...	...	...	...
#2048	200	#2048	200	#2048	200	#2048 201
#2052	300	#2052	300	#2052	300	#2052 335
#2056	400	#2056	400	#2056	400	#2056 400

```
addq $1,%rcx      addq (%rdx),%rcx    addl $1,(%r8,%r9,4)
addq %rbx,%rax    addq %rbx,(%rdx)    addl $1,%r9d
                    addq $3,(%rdx)    addl %eax,(%r8,%r9,4)
                                         addl $1,%r9d
                                         addl (%r8,%r9,4),%eax
```

# The Other ALU Instructions

- ▶ Most ALU instructions follow the same pattern as addX: two operands, second gets changed.
- ▶ Some one operand instructions as well.

Instruction	Name	Effect	Notes
addX B, A	Add	$A = A + B$	Two Operand Instructions
subX B, A	Subtract	$A = A - B$	
imulX B, A	Multiply	$A = A * B$	<i>Has a limited 3-arg variant</i>
andX B, A	And	$A = A \& B$	
orX B, A	Or	$A = A   B$	
xorX B, A	Xor	$A = A ^ B$	
salX B, A	Shift Left	$A = A << B$	B is constant or %cl reg
shlX B, A		$A = A << B$	
sarX B, A	Shift Right	$A = A >> B$	Arithmetic: Sign carry
shrX B, A		$A = A >> B$	Logical: Zero carry
incX A	Increment	$A = A + 1$	One Operand Instructions
decX A	Decrement	$A = A - 1$	
negX A	Negate	$A = -A$	
notX A	Complement	$A = \sim A$	

# leaX: Load Effective Address

- ▶ Memory addresses must often be loaded into registers
- ▶ Often done with a leaX, usually leaq in 64-bit platforms
- ▶ Sort of like “address-of” op & in C but a bit more general

INITIAL

REG	VAL
rax	0
rcx	2
rdx	#1024
rsi	#2048

```
## leaX_examples.s:  
movq 8(%rdx),%rax      # rax = *(rdx+1)      = 25  
leaq 8(%rdx),%rax      # rax = rdx+1          = #1032  
movl (%rsi,%rcx,4),%eax # rax = rsi[rcx]       = 400  
leaq (%rsi,%rcx,4),%rax # rax = &(rsi[rcx])     = #2056
```

Compiler sometimes uses leaX for multiplication as it is usually faster than imulX but less readable.

```
# Odd Collatz update n = 3*n+1  
#READABLE with imulX    #OPTIMIZED with leaX:  
imul $3,%eax  
addl $1,%eax  
# eax = eax*3 + 1        # eax = eax + 2*eax + 1,  
# 3-4 cycles              # 1 cycle
```



## Division: It's a Pain (1/2)

- ▶ idivX operation has some special rules
- ▶ Dividend must be in the `rax` / `eax` / `ax` register
- ▶ Sign extend to `rdx` / `edx` / `dx` register with “preparation” instructions like `cltd`
- ▶ `idivX` takes one **register** argument which is the divisor
- ▶ At completion
  - ▶ `rax` / `eax` / `ax` holds quotient (integer part)
  - ▶ `rdx` / `edx` / `dx` holds the remainder (leftover)

```
### division_example.s: start with uninitialized registers
## rax: 0x???????? = ? rdx: 0x???????? = ?
movl $5, %ecx          # will divide by 5
movl $19, %eax         # divide `int a=19;` by 5
cltd                  # prepare for 32-bit division
## combined 64-bit register %edx:%eax is
## rax: 0x0000000F = 15  rdx: 0x00000000 = 0
idivl %ecx            # 32-bit signed division
## 19 div 5 = 3 rem 4
## %eax == 3, quotient  %edx == 4, remainder
```

Compiler avoids division whenever possible: compile `col_unsigned.c` and `col_signed.c` to see some tricks involving shifts.

## Division: It's a Pain (2/2)

Division / preparation is irregular. See `division_examples.s` for more details which are summarized below.

```
movq $23, %rax          # divide `long a=23;` by 5
cqto                   # prepare for 64-bit division
idivq %rcx             # rax = 4, rdx = 3

movl $19, %eax          # divide `int a=19;` by 5
cltd                   # prepare for 32-bit division
idivl %ecx              # eax = 3, edx = 4

movw $37, %ax            # divide `short a=37;` by 5
cwtd                   # prepare for 16-bit division
idivw %cx                # ax = 7, dx = 2

## 8-Bit is a SPECIAL CASE
movb $42, %al            # divide `char a=42;` by 5
cbtw                   # prepare for 8-bit division
idivb %cl                # al = 8, ah = 2
## NOTE: quotient AND remainder in ax-family reg
```

Preparation instructions like `cltd` copy the sign bit from the AX-family register elsewhere: 0 for positive numbers, 1 for negative numbers, crucial for hardware division algorithm.