

CMSC216: x86-64 Assembly Extras and Wrap

Chris Kauffman

Last Updated:
Wed Oct 29 09:10:32 PM EDT 2025

Logistics

Reading Bryant/O'Hallaron

Read in Full

- ▶ Ch 3.7 Procedure Calls
- ▶ Ch 3.8-3.9: Arrays, Structs

Skim the following

- ▶ Ch 3.10: Pointers/Security
- ▶ Ch 3.11: Floating Point

Next

- ▶ Ch 8: Processes in Unix

Goals

- ▶ ~~Finish Function Calls / Stack~~
- ▶ ~~Data/Structs/Arrays in ASM~~
- ▶ Security Concerns
- ▶ Begin OS / Processes

Schedule

Date	Events
Mon 27-Oct	Dis: Lab08 1 / 2
Tue 28-Oct	Lec: Assembly Wrap
Wed 29-Oct	Dis: Lab08 2 / 2
Thu 30-Oct	Lec: Unix Processes
Mon 03-Nov	Dis: Bonus Review 2A P3 Due
Tue 04-Nov	Lec: Practice Exam 2B
Wed 05-Nov	Dis: Bonus Review 2B P3 Late Deadline
Thu 06-Nov	Exam 2

Announcements

Shell Slayer Tonight

<https://piazza.com/class/meuaquvb4xy46i/post/442>

- ▶ ACM Shell Slayer: A Shell Tutorial with a Daemonic Twist
- ▶ Tue 28-Oct 6-7pm
- ▶ Kim Engineering Building Room 1110

Costumes, Candy, Pizza, Shell Programming... Need I say more?

Reminders of Techniques for Puzzlebin

GDB Tricks from Quick Guide to GDB

Command	Effect
break *0x1248f2	Break at specific instruction address
break *func+24	Break at instruction with decimal offset from a label
break *func+0x18	Break at instruction with hex offset from a label
p \$rax	print: Print value in register rax
x \$rax	examine: Print memory pointed at by register rax
x /gx \$rax	Print as "giant" 64-bit numbers in hexadecimal format
x /5gd \$rax	Print 5 64-bit numbers starting where rax points in decimal format

Disassembling Binaries: objdump -d prog > code.txt

```
>> objdump -d a.out          # DISASSEMBLE BINARY
00000000000000001119 <main>:
    1119:   48 c7 c0 00 00 00 00    mov    $0x0,%rax
    1120:   48 c7 c1 01 00 00 00    mov    $0x1,%rcx
    1127:   48 c7 c2 64 00 00 00    mov    $0x64,%rdx
0000000000000000112e <LOOP>:
    112e:   48 39 d1              cmp    %rdx,%rcx
    1131:   7f 08                jg     113b <END>
    1133:   48 01 c8              add    %rcx,%rax
...
>> objdump -d a.out > code.txt # STORE RESULTS IN FILE
```

Accessing Global Variables in Assembly

Global data can be set up in assembly in .data sections with labels and assembler directives like .int and .short

```
.data
an_int:           # single int
    .int 17
some_shorts:      # array of shorts
    .short 10   # some_shorts[0]
    .short 12   # some_shorts[1]
    .short 14   # some_shorts[2]
```

Modern Access to Globals

```
movl an_int(%rip), %eax
leaq some_shorts(%rip), %rdi
```

- ▶ Uses %rip relative addressing
- ▶ Default in gcc as it plays nice with OS security features like *Address Space Randomization*
- ▶ May discuss again later during Linking/ELF coverage
(Bryan/O'Hallaron Chapter 7)

Traditional Access to Globals

```
movl an_int, %eax          # ERROR
leaq (some_shorts), %rdi  # ERROR
```

- ▶ Not accepted by gcc by default
- ▶ Yields compile/link errors

```
/usr/bin/ld: /tmp/ccocSiw5.o:
relocation R_X86_64_32S against `'.data'
can not be used when making a PIE object;
recompile with -fPIE
```

Aggregate Data In Assembly (Arrays + Structs)

Arrays

Usually: base + index × size

```
arr[i] = 12;  
movl $12,(%rdi,%rsi,4)
```

```
int x = arr[j];  
movl (%rdi,%rcx,4),%r8d
```

- ▶ Array starting address often held in a register
- ▶ Index often in a register
- ▶ Compiler inserts appropriate size (1,2,4,8)

Structs

Usually base+offset

```
typedef struct {  
    int i; short s;  
    char c[2];  
} foo_t;  
foo_t *f = ...;
```

```
short sh = f->s;  
movw 4(%rdi),%si  
  
f->c[i] = 'X';  
movb $88, 6(%rdi,%rax)
```

Packed Structures as Procedure Arguments

- ▶ Passing pointers to structs is 'normal': registers contain addresses to main memory
- ▶ Passing actual structs may result in *packed structs* where several fields are in a single register
- ▶ Assembly must *unpack* these through **shifts and masking**

```
1 // packed_struct_main.c
2 typedef struct {
3     short first;
4     short second;
5 } twoshort_t;
6
7 short sub_struct(twoshort_t ti);
8
9 int main(){
10    twoshort_t ts = {.first=10,
11                      .second=-2};
12    int sum = sub_struct(ts);
13    printf("%d - %d = %d\n",
14           ts.first, ts.second, sum);
15    return 0;
16 }
```

```
1 ### packed_struct.s
2 .text
3 .globl sub_struct
4 sub_struct:
5 ## first arg is twoshort_t ts
6 ## %rdi has 2 packed shorts in it
7 ## bits 0-15 are ts.first
8 ## bits 16-31 are ts.second
9 ## upper bits could be anything
10
11 movl %edi,%eax          # eax = ts;
12 andl $0xFFFF,%eax       # eax = ts.first;
13 sarl $16,%edi           # edi = edi >> 16;
14 andl $0xFFFF,%edi       # edi = ts.second;
15 subw %di,%ax             # ax = ax - di
16 ret                      # answer in ax
```

Example: coins_t in Lab06

```
// Type for collections of coins
typedef struct { // coint_t has the following memory layout
    char quarters; // 
    char dimes; // | Pointer | Packed | Packed |
    char nickels; // | Memory | Struct | Struct |
    char pennies; // | Field | Offset | Arg# | Bits |
} coins_t;
// |-----+-----+-----+
// | quarters | +0 | #1 | 0-7 |
// | dimes | +1 | #1 | 8-15 |
// | nickels | +2 | #1 | 16-23 |
// | pennies | +3 | #1 | 24-31 |

# EXAMPLE: coins_t *c = <#2048>;
# | #2048 | c->quarters | 2 |
# | #2049 | c->dimes | 1 |
# | #2050 | c->nickels | - |
# | #2051 | c->pennies | - |
# set_coins(c);

set_coins:
### int set_coins(int cents, coins_t *coins)
### %edi = int cents
### %rsi = coints_t *coins
...
# rsi: #2048
# al: 0 %dl: 3
movb    %al,2(%rsi)      # coins->nickels = al;
movb    %dl,3(%rsi)      # coins->pennies = dl;

## | #2048 | c->quarters | 2 |
## | #2049 | c->dimes | 1 |
## | #2050 | c->nickels | 0 |
## | #2051 | c->pennies | 3 |

# EXAMPLE:
# coins_t c = { .quarters=2, dimes=1,
#                 .nickels=0, .pennies=3 };
# int tot = total_coins(c);

total_coins:
### args are
### %rdi packed coin_t struct with struct fields
### { 0-7: quarters, 8-15: dimes,
###   16-23: nickels, 24-31: pennies}

### rdi: 0x00 00 00 00 03 00 01 02
###           p n d q
    movq    %rdi,%rdx      # extract dimes
### rdx: 0x00 00 00 00 03 00 01 02
###           p n d q
    sarq    $8,%rdx      # shift dimes to low bits
### rdx: 0x00 00 00 00 00 03 00 01
###           p n d
    andq    $0xFF,%rdx      # rdx = dimes
### rdx: 0x00 00 00 00 00 00 00 01
###           p n d
```

Large Packed Structs

Structs that don't fit into single registers may be packed across several argument registers

```
typedef struct{
    int    day_secs;   // 4
    short time_secs; // 2
    short time_mins; // 2
    short time_hours; // 2
    char   ampm;      // 1+1 pad
} tod_t;                      // 12 bytes
```

```
int set_display_from_tod(tod_t tod, ...)
//                                              ^^^ Large packed struct
```

C Field Access	Register	Bits in reg	Shift Required	Size
tod.day_secs	%rdi	0-31	None	4 bytes
tod.time_secs	%rdi	32-47	Right by 32	2 bytes
tod.time_mins	%rdi	48-63	Right by 48	2 bytes
tod.time_hours	%rsi	0-15	None	2 bytes
tod_ampm	%rsi	16-23	Right by 16	1 bytes

At a certain size, compiler stores Very Large packed structs in the stack and passes pointers to it to functions

General Cautions on Structs

Struct Layout by Compilers

- ▶ Compiler honors order of source code fields in struct
- ▶ BUT compiler may add padding between/after fields for alignment
- ▶ Compiler determines total struct size

Struct Layout Algorithms

- ▶ Baked into compiler
- ▶ **May change from compiler to compiler**
- ▶ May change through history of compiler

Structs in Mem/Regs

- ▶ Local var structs spread across several registers
- ▶ Don't need a struct on the stack at all in some cases (just like don't need local variables on stack)
- ▶ Struct arguments packed into 1+ registers

Stay Insulated

- ▶ Programming in C insulates you from all of this
- ▶ Feel the **warmth** of gcc's abstraction blanket

Security Risks in C

Buffer Overflow Attacks

- ▶ No default bounds checking in C:
Performance favored over safety
- ▶ Allows classic security flaws:

```
char buf[1024];
printf("Enter your name:");
fscanf(file,"%s",buf); // BAD
// or
gets(buf);           // BAD
// my name is 1500 chars
// long, what happens?
```

- ▶ For data larger than buf, begin overwriting other parts of the stack
 - ▶ Clobber return addresses
 - ▶ Insert executable code and run it

Counter-measures

- ▶ **Stack protection** is default in gcc in the modern era
- ▶ Inserts “canary” values on the stack near return address
- ▶ Prior to function return, checks that canaries are unchanged
- ▶ **Stack / Text Section Start randomized** by kernel, return address and function addresses difficult to predict ahead of time
- ▶ Kernel may also vary virtual memory address as well
- ▶ Disabling protections is risky

Stack Smashing

- ▶ Explored in a recent homework
- ▶ See stack_smash.c for a similar example
- ▶ Demonstrates detection of changes to stack that could be harmful / security threat

```
// stack_smash.c
void demo(){
    int arr[4];    // fill array off the end
    for(int i=0; i<8; i++){
        arr[i] = (i+1)*2;
    }

    for(int i=0; i<8; i++){
        printf("[%d]: %d\n", i, arr[i]);
    }
}

int main(){
    printf("About to do the demo\n");
    demo();
    printf("Demo Complete\n");
    return 0;
}
```

> cd 08-assembly-extras-code/
> gcc stack_smash1.c
> ./a.out

About to do the demo

[0]: 2
[1]: 4
[2]: 6
...
[7]: 16

*** stack smashing detected ***:
terminated
Aborted (core dumped)

Demonstration of Buffer Overflow Attack

- ▶ See the code `buffer_overflow.c`
- ▶ Presents an easier case to demo stack manipulations
- ▶ Prints addresses of functions `main()` and `never()`
- ▶ Reads long values which are 64-bits, easier to line up data in stack than with strings; still overflowing the buffer by reading too much data as in:

```
void always(){  
    long buf[1] = {0xABCD};           // room for 1  
    ...  
    printf("Enter 4 hex values: ");  
    fscanf(stdin,"%lx %lx %lx %lx", // reads 4  
           &buf[0], &buf[1], &buf[2], &buf[3]);
```

- ▶ When compiled via

```
>> gcc -fno-stack-protector buffer_overflow.c
```

can get `never()` to run by entering its address as input which will overwrite the return address

Sample Buffer Overflow Code

```
#include <stdio.h>
void print_all_passwords(){
    ...
}

int main(){
    printf("file to open: ");
    char buf[128];
    fscanf(stdin,"%s",buf);
    printf("You entered: %s\n",buf);

    ...
    return 0;
// By entering the correct length of string followed by the ASCII
// representation of the address of print_all_passwords(), one might
// be able to get that function when "return" is reached if there
// are no stack protection mechanisms at work ...
// (which was the case in 1999 on Windows :--)
}
```

Details of GCC / Linux Stack Security

- ▶ Programs compiled with GCC + Glibc on Linux for x86-64 will default to having stack protection
- ▶ This can be seen in compiled code as short blocks near the beginning and end of functions which
 1. At the beginning of the function uses an instruction like `movq %fs:40, %rax` and places a value in the stack beneath the return address
 2. At the end of the function again accesses `%fs:40` and the value earlier placed in the stack.
- ▶ The `%fs` register is a special [segment register](#) originally introduced in the 16-bit era to surmount memory addressing limitations; now used only for limited purposes
- ▶ The complete details are beyond the scope of our course BUT
- ▶ **A somewhat detailed explanation has been added to 08-assembly-extras-code/stack_protect.org**

Floating Point Operations

- ▶ Original Intel 8086 Processor **didn't do floating point ops**
- ▶ Had to buy a co-processor (Intel 8087) to enable FP ops
- ▶ Most modern CPUs support FP ops but they feel separate from the integer ops: FPU versus ALU

x86-64 “Media” Registers

512	256	128-bits	Use
%zmm0	%ymm0	%xmm0	FP Arg1/Ret
%zmm1	%ymm1	%xmm1	FP Arg2
...
%zmm7	%ymm7	%xmm7	FP Arg 8
%zmm8	%ymm8	%xmm8	Caller Save
...
%zmm15	%ymm15	%xmm15	Caller Save

Instructions

```
vaddss %xmm2,%xmm4,%xmm0  
# xmm0[0] = xmm2[0] + xmm4[0]  
# Add Scalar Single-Precision
```

```
vaddps %xmm2,%xmm4,%xmm0  
# xmm0[:] = xmm2[:] + xmm4[:]  
# Add Packed Single-Precision  
# "Vector" Instruction
```

- ▶ Can be used as “scalars” - single values but...
- ▶ `xmmI` is 128 bits big holding
 - ▶ 2 x 64-bit double's OR
 - ▶ 4 x 32-bit float's
- ▶ `ymmI` / `zmmI` extend further

- ▶ Operates on single values or “vectors” of packed values
- ▶ 3-operands common in more “modern” assembly languages

Example: float_ops.c to Assembly

```
// float_ops.c: original C Code
void array_add(float *arr1, float *arr2, int len){
    for(int i=0; i<len; i++){
        arr1[i] += arr2[i];
    }
}

# >> gcc -S -Og float_ops.c
# Minimal optimizations
array_add:    ## 16 lines asm
.LFB0:
    .cfi_startproc
    movl $0, %eax
    jmp .L2
.L3:
    movslq %eax, %r8
    leaq (%rdi,%r8,4), %rcx
    movss (%rsi,%r8,4), %xmm0
    addss (%rcx), %xmm0    ## add single
    movss %xmm0, (%rcx)    ## single prec
    addl $1, %eax
.L2:
    cmpl %edx, %eax
    jl .L3
    ret

# >> gcc -S -O3 -maxavx float_ops.c
# Max optimizations, Use AVX hardware
array_add:    ## 100 lines asm
...
.L5:          ## vector move/adds
    vmovups (%rcx,%rdx), %ymm1
    vaddps  (%rsi,%rdx), %ymm1, %ymm0
    vmovups %ymm0, (%rcx,%rdx)
    addq $32, %rdx
    cmpq %rdi, %rdx
    jne .L5
...
.L9:          ## single move/adds
    vmovss  (%rcx,%rax), %xmm0
    vaddss  (%rsi,%rax), %xmm0, %xmm0
    vmovss %xmm0, (%rcx,%rax)
    addq $4, %rax
    cmpq %rax, %rdx
    jne .L9
    ret
```

Floating Point and ALU Conversions

- ▶ Recall that bit layout of Integers and Floating Point numbers are quite different (**how?**)
- ▶ Leads to a series of assembly instructions to interconvert between types

```
# file:float_convert.c

# int eax = ...;
# double xmm0 = (double) eax;
cvtsi2sdl    %eax, %xmm0

# double xmm1 = ...
# long rcx = (long) xmm1;
cvttsd2siq   %xmm0, %rax
```

- ▶ These are non-trivial conversions: 5-cycle latency (delay) before completion, can have a performance impact on code which does conversions

Optional Exercise: All Models are Wrong...

- ▶ Rule #1: The Doctor Lies
- ▶ Below is our original model for memory layout of C programs
- ▶ Describe what is **incorrect** based on x86-64 assembly
- ▶ What is **actually** in the stack? How are registers likely used?

```
9: int main(...){          STACK: Caller main(), prior to swap()
10:   int x = 19;           | FRAME    | ADDR   | NAME   | VALUE  |
11:   int y = 31;           |-----+-----+-----+-----|
+-<12:   swap(&x, &y);    | main()   | #2048  | x      |     19  |
| 13:   printf("%d %d\n",x,y); | line:12 | #2044  | y      |     31  |
| 14:   return 0;          |-----+-----+-----+-----|
V 15: }
```

|

```
| 18: void swap(int *a,int *b){  STACK: Callee swap() takes control
|     int tmp = *a;             | FRAME    | ADDR   | NAME   | VALUE  |
+->19:     *a = *b;           |-----+-----+-----+-----|
|     *b = tmp;               | main()   | #2048  | x      |     19 |<-+
|     return;                 | line:12 | #2044  | y      |     31 |<-|+
| 23: }
```

```
|-----+-----+-----+-----|   ||
```

```
|     swap()  | #2036  | a      | #2048 |---+|
|     line:19 | #2028  | b      | #2044 |---+|
|             | #2024  | tmp    |       ?  |
```

Answers: All Models are Wrong, Some are Useful

```
9: int main(...){  
10:    int x = 19;  
11:    int y = 31;  
+-<12:    swap(&x, &y);  
| 13:    printf("%d %d\n",x,y);  
| 14:    return 0;  
V 15: }  
|  
| 18: void swap(int *a,int *b){  
+->19:    int tmp = *a;  
20:    *a = *b;  
21:    *b = tmp;  
22:    return;  
23: }
```

STACK: Callee swap() takes control			
FRAME	ADDR	NAME	VALUE
main()	#2048	x	19
	#2044	y	31
swap()	#2036	rip	Line 13

REGS as swap() starts

REG	VALUE	NOTE
rdi	#2048	for *a
rsi	#2044	for *b
rax	?	for tmp
rip	L19	line in swap

- ▶ main() must have stack space for locals passed by address
- ▶ swap() needs no stack space for arguments: in registers
- ▶ Return address is next value of rip register in main()
- ▶ Mostly don't need to think at this level of detail but **can be useful in some situations**