

CMSC 216: UNIX File Input/Output

Chris Kauffman

Last Updated:
Mon Nov 10 01:11:30 PM EST 2025

Logistics

Announcements

Exercise: C Standard I/O Functions

Recall basic I/O functions from the C Standard Library header
`stdio.h`

1. Printing things to the screen?
2. Opening a file?
3. Closing a file?
4. Printing to a file?
5. Scanning from terminal or file?
6. Get whole lines of text?
7. Names for standard input, output, error

Give samples of function calls

Answers: C Standard I/O Functions

Recall basic I/O functions from the C Standard Library header stdio.h

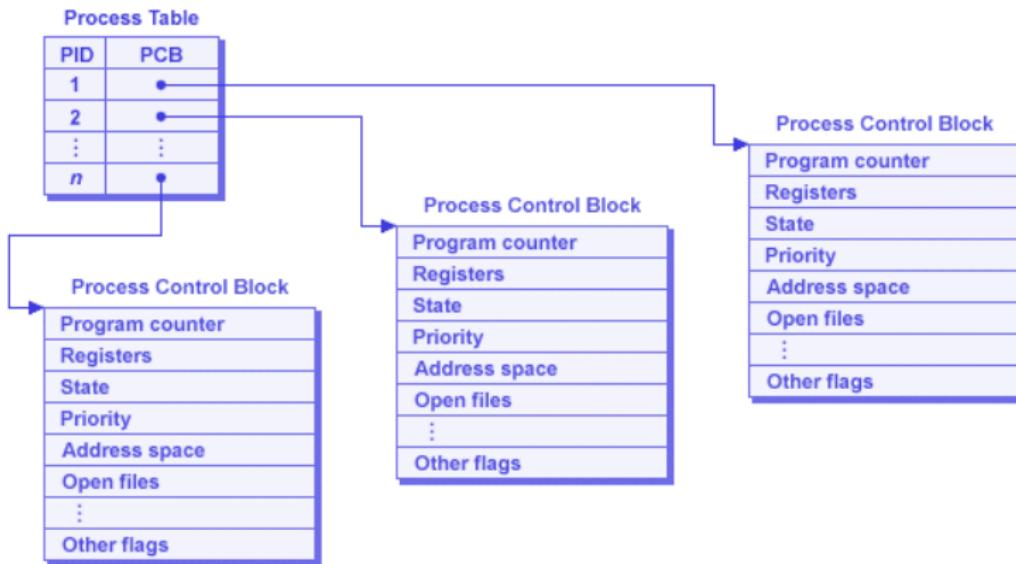
1	printf("%d is a number",5);	Printing things to the screen?
2	FILE *file = fopen("myfile.txt","r");	Opening a file?
3	fclose(file);	Close a file?
4	fprintf(file,"%d is a number",5);	Printing to a file?
5	scanf("%d %f",&myint,&mydouble); fscanf(file2,"%d %f",&myint,&mydouble);	Scanning from terminal or file?
6	result = fgets(charbuf, 1024, file);	Get whole lines of text?
7	FILE *stdin, *stdout, *stderr;	Names for standard input, etc

The standard I/O library was written by Dennis Ritchie around 1975.

–Stevens and Rago, Advanced Programming for the Unix Environment

- ▶ Assuming you are familiar with these and could look up others like fgetc() (single char) and fread() (read binary)
- ▶ Library Functions: available with any compliant C compiler
- ▶ On Unix systems, fscanf(), FILE*, and the like are backed by lower level System Calls and Kernel Data Structures

The Process Table




Source:

SO What is the Linux Process Table?

- ▶ OS maintains data on all processes in a Process Table
- ▶ Process Table Entry ≈ Process Control Block
- ▶ Contains info like PID, instruction that process is executing*, Virtual Memory Address Space and **Files in Use**

File Descriptors



- ▶ Each Process Table entry contains a table of open files
- ▶ A user program refers to these via **File Descriptors**
- ▶ File Descriptor is an integer index into Kernel's table
<MINDED>
- ▶ FD Table entry refers to other Kernel/OS data structures

File Descriptors (FDs) are Multi-Purpose

- ▶ Unix tries to provide most things via files/file descriptor
- ▶ Many Unix system actions are handled via `read()`-from or `write()`-to file descriptors
- ▶ FDs allow interaction with “normal” files like `myfile.txt` or `commando.c` to read/change them
- ▶ FDs also allow interaction with many other things
 - ▶ Pipes for interprocess communication
 - ▶ Sockets for network communication
 - ▶ Special files to manipulate terminal, audio, graphics, etc.
 - ▶ Raw blocks of memory for Shared Memory communication
 - ▶ Even processes themselves have special files in the file system:
`ProcFS` in `/proc/PID#`, provide info on running process
- ▶ We will focus on standard File I/O using FDs now and touch on some broader uses Later
- ▶ Also must discuss FD interactions with previous System Calls:
What happens with `open()` files when calling `fork()`?

Open and Close: File Descriptors for Files

<MINDED>

`open()` / `close()` show common features of many system calls

- ▶ Returns -1 on errors
- ▶ Show errors using the `perror()` function
- ▶ Use of vertical pipe (|) to bitwise-OR several options

read() from File Descriptors

<MINTED>

- ▶ Read up to SIZE from an open file descriptor
- ▶ Bytes stored in buffer, overwrite it
- ▶ Return value is number of bytes read, -1 for error
- ▶ SIZE commonly defined but can be variable, constant, etc
- ▶ **Examine `read_some.c`:** explain what's happening

Caution:

- ▶ Bad things happen if buffer is actually smaller than SIZE
- ▶ `read()` does NOT null terminate, add \0 manually if needed

Exercise: Behavior of read() in count_bytes.c

Run count_bytes.c on
file data.txt

<MINTED>

1. Explain control flow <MINTED>
within program
2. Predict output of
program

Answers: Behavior of read() in count_bytes.c

<MINTED>

<MINTED>

<MINTED>

Answers: Behavior of `read()` in `count_bytes.c`

Take-Aways from `count_bytes.c` include

- ▶ OS maintains **file positions** for each open File Descriptor
- ▶ I/O functions like `read()` use/change position **in a file**
- ▶ `read()`'ing into program arrays overwrites data there
- ▶ OS **does not** update positions in user arrays: programmer must do this in their program logic
- ▶ `read()` returns # of bytes read, may be less than requested
- ▶ `read()` returns 0 when at end of a file

Exercise: write() to File Descriptors

<MINDED>

- ▶ Write up to SIZE bytes to open file descriptor
- ▶ Bytes taken from buffer, leave it intact
- ▶ Return value is number of bytes written, -1 for error

Questions on write_then_read.c

- ▶ Compile and Run
- ▶ Explain Output, differences between write() / printf()

Answers: write() to File Descriptors

<MINTED>

read()/write() work with bytes

- ▶ In C, general correspondence between byte and the char type
- ▶ Not so for other types: int is often 4 bytes
- ▶ Requires care with non-char types
- ▶ All calls read/write actual bytes

<MINDED>

Questions

- ▶ Examine `write_read_ints.c`, compile/run
- ▶ Examine contents of `integers.dat`
- ▶ Explain what you see

Standard File Descriptors


- ▶ When a process is born, comes with 3 open file descriptors
- ▶ Related to FILE* streams in Standard C I/O library
- ▶ Traditionally have FD values given but use the Symbolic name to be safe

Symbol	#	FILE*	FD for...
STDIN_FILENO	0	stdin	standard input (keyboard)
STDOUT_FILENO	1	stdout	standard output (screen)
STDERR_FILENO	2	stderr	standard error (screen)


<MINTED>

See `low_level_interactions.c` to gain an appreciation for what `printf()` and its kin can do for you.


File Descriptors refer to Kernel Structures



```
my_fd1 = open("file.txt", O_RDONLY);
```



Processes Inherit Open FDs: Diagram



Typical sequence:

- ▶ Parent creates an `output_fd` and/or `input_fd`
- ▶ Call `fork()`
- ▶ Child changes standard output to `output_fd` and/or `input_fd`
- ▶ Changing means calls to `dup2()`

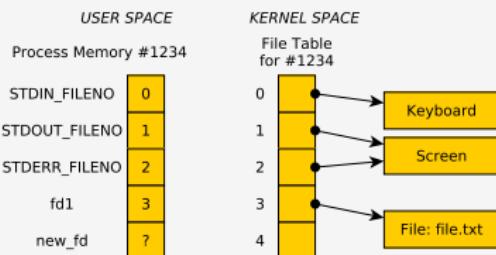
Shell I/O Redirection

- ▶ Shells can direct input / output for programs using < and >
- ▶ Most common conventions are as follows
<MINTED>
- ▶ Long output can be saved easily
- ▶ Can save typing input over and over
- ▶ Even more fun when you incorporate [Pipes to make Pipelines](#)
- ▶ **Goal:** Demonstrate systems calls to facilitate redirection


Manipulating the File Descriptor Table

- ▶ System calls `dup()` and `dup2()` manipulate the FD table
- ▶ `int backup_fd = dup(fd);` : copy a file descriptor
- ▶ `dup2(src_fd, dest_fd);` : `src_fd` copied to `dest_fd`


Effect of `dup()`: copy a file descriptor




`new_fd = dup(fd1);`



Effect of `dup2()`: change entry in FD table



`dup2(fd1, STDOUT_FILENO);
source destination`



Exercise: Redirecting Output with dup() / dup2()

- ▶ `dup()`, `dup2()`, and `fork()` can be combined in interesting ways
- ▶ **Diagram fork-dup.pdf** shows how to redirect standard out to a file like a shell does in: `ls -l > output.txt`

Write a program which

1. Prints PID to screen
2. Opens a file named `write.txt`
3. Forks a Child process
4. Child: **redirect standard output**
into `write.txt`
Parent: does no redirection <MINTED>
5. Both: `printf()` their PID
6. Child: **restore standard output to screen**
Parent: makes no changes
7. Both: `printf()` "All done"

Answers: Redirecting Output with dup() / dup2()

<MINTED>

C FILE Structs Use File Descriptors in UNIX

Typical Unix implementation of standard I/O library FILE is

- ▶ A file descriptor
- ▶ Some buffers with positions
- ▶ Some options controlling buffering

From /usr/include/bits/types/struct_FILE.h
<MINDED>

Exercise: Subtleties of Mixing Standard / Low-Level I/O

Sample compile/run:

<MINTED>

- ▶ Explain output of program given input file
- ▶ Use knowledge that **buffering** occurs internally for standard I/O library

<MINTED>

<MINTED>

Answers: Subtleties of Mixing Standard / Low-Level I/O

- ▶ C standard I/O calls like `printf` / `fprintf()` and `scanf`() / `fscanf()` use internal buffering
- ▶ A call to `fscanf(file, "%d", &x)` will read a large chunk from a file but only process part of it
- ▶ From OS perspective, associated file descriptor has advanced forwards / read a bunch
- ▶ The data is in a hidden “buffer” associated with a `FILE *file`, used by `fscanf()`

Output Also buffered, Always `fclose()`

- ▶ Output is also buffered: `output_buffering.c`
- ▶ Output may be lost if `FILE*` are not `fclose()`'d: closing will flush remaining output into a file
- ▶ See `fail_to_write.c`
- ▶ File descriptors always get flushed out by OS when a program ends BUT `FILE*` requires user action
- ▶ To force output, use `fflush(some_file);`

Controlling FILE Buffering

<MINTED>

Above functions change buffering behavior of standard C I/O
Examples:

<MINTED>

- ▶ When testing lab/project code, buffering is disabled as it makes it easier to understand some bugs

Basic File Statistics via stat

Command	C function	Effect
stat file	int ret = stat(file,&statbuf); int ret = lstat(file,&statbuf);	Get statistics on file Same, don't follow symlinks
	int fd = open(file,...); int ret = fstat(fd,&statbuf);	Same as above but with an open file descriptor

Shell command stat provides basic file info such as shown below

<MINDED>

See `stat_demo.c` for info on C calls to obtain this info

Attributes of Files from stat()

stat_demo.c shows some attributes that may be obtained about a file after a call to `stat(filename, &statbuf)` which fills in the `statbuff` struct. Attributes include:

Attribute	Notes
Size	In bytes via <code>st_size</code> field
File Type	Via <code>st_mode</code> field and macros like <code>S_ISREG(mode)</code> Limited number of fundamental types: regular, directory, socket, etc.
Permissions	Read/Write/Execute for Owner/Group/Others via <code>st_mode</code> field
Ownership	Via <code>st_uid</code> (user) and <code>st_gid</code> (group), numeric IDs for both
Time Data	Access / Change / Modification times via <code>st_atime</code> , <code>st_ctime</code> , ...

Permissions / Modes

- ▶ Unix enforces file security via *modes*: permissions as to who can read / write / execute each file
- ▶ See permissions/modes with `ls -l`
- ▶ Look for series of 9 permissions

<MINDED>

- ▶ Every file has permissions set from somewhere on creation

Changing Permissions

Owner of file (and sometimes group member) can change permissions via `chmod`

<MINDED>

- ▶ `chmod` also works via octal bits (suggest against this unless you want to impress folks at parties)
- ▶ Programs specify file permissions via system calls
- ▶ Curtailed by **Process User Mask** which indicates permissions that are disallowed by the process
 - ▶ `umask` shell function/setting: `$> umask 007`
 - ▶ `umask()` system call: `umask(S_IWGRP | S_IWOTH);`
- ▶ Common program strategy: create files with very liberal read/write/execute permissions, `umask` of user will limit this

Permissions / Modes in System Calls

open() can take 2 or 3 arguments

<MINTED>

Symbol	Entity	Sets
S_IRUSR	User	Read
S_IWUSR	User	Write
S_IXUSR	User	Execute
S_IRGRP	Group	Read
S_IWGRP	Group	Write
S_IXGRP	Group	Execute
S_IROTH	Others	Read
S_IWOTH	Others	Write
S_IXOTH	Others	Execute

Compare: write_readable.c VERSUS write_unreadable.c

<MINTED>

Movement within Files, Changing Sizes

- ▶ Can move OS internal position in a file around with `lseek()`
- ▶ Note that size is arbitrary: can seek to any positive position
- ▶ File automatically expands if position is larger than current size - fills holes with 0s (null chars)
- ▶ Can manually set size of a file with `ftruncate(fd, size)`
- ▶ Examine `file_hole1.c` and `file_hole2.c`

C function	Effect
<code>int res = lseek(fd, offset, option);</code>	Move position in file
<code>lseek(fd, 20, SEEK_CUR);</code>	Move 20 bytes forward
<code>lseek(fd, 50, SEEK_SET);</code>	Move to position 50
<code>lseek(fd, -10, SEEK_END);</code>	Move 10 bytes from end
<code>lseek(fd, +15, SEEK_END);</code>	Move 15 bytes beyond end
<code>ftruncate(fd, 64);</code>	Set file to be 64 bytes big If file grows, new space is zero-filled

Note: C standard I/O functions `fseek(FILE*)` and `rewind(FILE*)` mirror functionality of `lseek()`

Directory Access

- ▶ Directories are fundamental to Unix (and most file systems)
- ▶ Unix file system rooted at / (root directory)
- ▶ Subdirectories like bin, ~/home, and /home/kauffman
- ▶ Useful shell commands and C function calls pertaining to directories are as follows

Shell Command	C function	Effect
mkdir name	int ret = mkdir(path,perms);	Create a directory
rmdir name	int ret = rmdir(path);	Remove empty directory
cd path	int ret = chdir(path);	Change working directory
pwd	char *path = getcwd(buf,SIZE);	Current directory
ls	DIR *dir = opendir(path); struct dirent *file = readdir(dir); int ret = closedir(dir);	List directory contents Start reading filenames from dir Call in a loop, NULL when done After readdir() returns NULL

See `dir_demo.c` for demonstrations

Optional Exercise: Code for Total Size of Regular Files

- ▶ Code which will scan all files in a directory
- ▶ Will get file statistics on each file
- ▶ Skips directories, symlinks, etc. <MINTED>
- ▶ Totals bytes of all Regular files in current directory

Use techniques demoed in
`dir_demo.c` and `stat_demo.c`
from codepack

Answers: Sketch Code for Total Size of Regular Files

- ▶ Scans only current directory
- ▶ **Recursive scanning** is trickier and involves... recursion
- ▶ OR the very useful `nftw()` library function (read about this on your own if curious about systems programming)

<MINTED>

Extras: Processes Inherit Open FDs

- ▶ Child processes share all open file descriptors with parents
- ▶ By default, Child prints to screen / reads from keyboard input
- ▶ Redirection requires manipulation prior to `fork()`
- ▶ See: `open_fork.c`
- ▶ Experiment with order
 1. `open()` then `fork()`
 2. `fork()` then `open()`

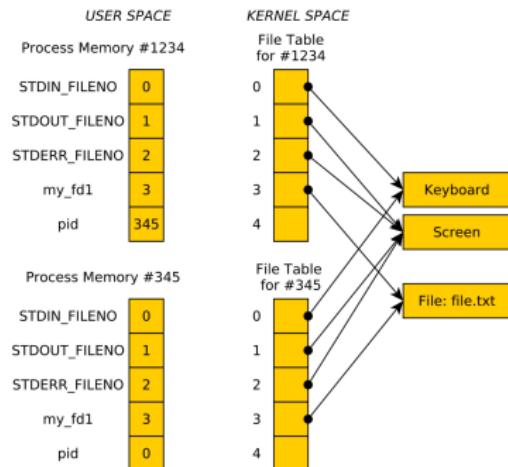
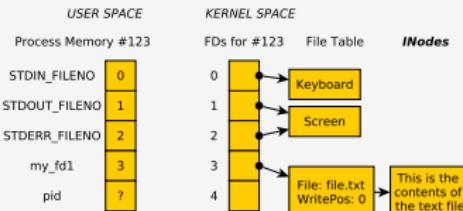


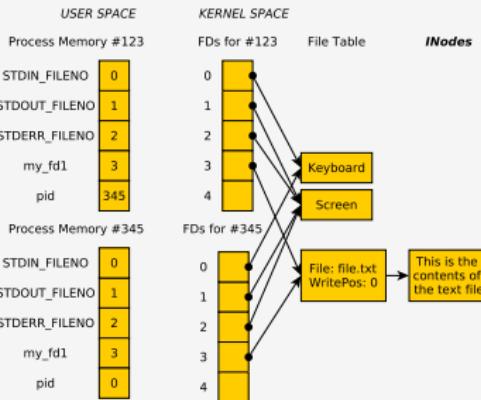
Diagram on next slide shows variations of open-then-fork vs fork-then-open from `open_fork.c`

open() normal file then call fork()

`my_fd = open("file.txt"); // called by parent`




`pid = fork();`




fork() then call open() normal file

`pid = fork();`



`my_fd = open("file.txt"); // called by parent and child`



(Review) Exercise: Regular File Creation Basics

C Standard I/O

- ▶ Write/Read data?
- ▶ Open a file, create it if needed?
- ▶ Result of opening a file?
- ▶ Close a file?
- ▶ Set permissions on file creation?

Unix System Calls

- ▶ Write/Read data?
- ▶ Open a file, create it if needed?
- ▶ Result of opening a file?
- ▶ Close a file?
- ▶ Set permissions on file creation?

Answers: Regular File Creation Basics

C Standard I/O

- ▶ Write/Read data?

```
fscanf(), fprintf()  
fread(), fwrite()
```

- ▶ Open a file, create it if needed?
- ▶ Result of opening a file?

```
FILE *out =  
    fopen("myfile.txt", "w");
```

- ▶ Close a file?

```
fclose(out);
```

- ▶ Set permissions on file creation?
Not possible... dictated by umask

Unix System Calls

- ▶ Write/Read data?

```
write(), read()
```

- ▶ Open a file, create it if needed?
- ▶ Result of opening a file?

```
int fd =  
    open("myfile.txt",  
          O_WRONLY | O_CREAT,  
          permissions);
```

- ▶ Close a file?

```
close(fd);
```

- ▶ Set permissions on file creation?

- ▶ Additional options to
open(), which brings us
to...