

MPI Basics

Chris Kauffman

Last Updated:
Thu Feb 12 01:46:03 PM EST 2026

Logistics

Reading: Grama Ch 6 + 4

- ▶ Ch 6: MPI basics
- ▶ Ch 4: Communication patterns
- ▶ Zaratan Setup Link on Schedule (Week 1)

Assignments

- ▶ A1 Due Soon
 - ▶ On-time by Thu 12-Feb
 - ▶ Late through Sat 14-Feb
- ▶ **Questions?**
- ▶ A2 up next week:
MPI Programming

Today

- ▶ Primitives for Distributed Memory Computing
- ▶ MPI Programming

Next Week

- ▶ Communication Patterns
- ▶ Thu 19-Feb: Mini-Exam 1

Announcements

Assignment Clarifications

Several clarifications regarding Assignment 1 discussed on Piazza posts, some of which are added to the assignment.

- ▶ Network embedding: Mesh into Ring example from last week
- ▶ Speedup definition
- ▶ Assumptions about perfect trees in networks
- ▶ Meaning of times

Coding Environment Setup

Zaratan Setup added to schedule in Week 1

<https://www.cs.umd.edu/~profk/416/setup-guide.html>

- ▶ Good idea to code your K-Means solution there to get acquainted as later codes MUST run there
- ▶ Provided test cases / testy run on Zaratan and most Linuxes, likely not on other platforms

Generic Send and Receive

- ▶ Distributed memory machines require explicit sharing of data
- ▶ Minimum required functionality is:

```
send(void *sendbuf, int nelems, int dst_proc);  
receive(void *recvbuf, int nelems, int src_proc);
```

- ▶ Referred to as a “point-to-point” communication
- ▶ **Sample Use**

```
1 // P0 runs           // P1 runs  
2 a = 100;            receive(&a, 1, 0)  
3 send(&a, 1, 1);     printf("%d\n", a);  
4 a=50;
```

- ▶ Proc 0 sends a single integer to Proc 1
- ▶ Proc 0 then changes that integer
- ▶ Proc 1 receives and prints the integer

Programmatic Send/Receive

- ▶ Typically write this as a single program which every processor runs: Single Program, Multiple Data (SPMD)
- ▶ Assume availability of a function giving logical Proc Number
- ▶ Branching on proc number to take different actions

```
1 void exchange(){
2     int a = -1;
3     int my_proc = get_processor_number();
4     if(my_proc == 0){
5         a = 100;
6         send(&a, 1, 1);      // send data 100 to Proc 1
7         a=50;
8     }
9     else if(my_proc == 1){
10        receive(&a, 1, 0);  // receive data from Proc 0
11        printf("%d\n", a);
12    }
13    printf("my_proc: %d  a: %d\n",my_proc,a);
14 }
```

Flavors of Send/Receive

- ▶ Hardware+OS may support copying message into a “buffer” space which allows sending program to proceed faster: copy to buffer, OS/Hardware handles the rest
- ▶ Functions usually available to do both Blocking send() and Nonblocking “immediate” isend() BUT without OS/Hardware support they are the same

| | Blocking Operations | Non-Blocking Operations |
|--------------|------------------------------------------------------------------------------|----------------------------------------------------------------------------------------------------------------|
| Buffered | Sending process returns after data has been copied into communication buffer | Sending process returns after initiating DMA transfer to buffer. This operation may not be completed on return |
| Non-Buffered | Sending process blocks until matching receive operation has been encountered | Send and Receive semantics assured by corresponding operation |

Figure 6.3 Space of possible protocols for send and receive operations.

Blocking + Unbuffered Send/Receive




Figure 6.1 Handshake for a blocking non-buffered send/receive operation. It is easy to see that in cases where sender and receiver do not reach communication point at similar times, there can be considerable idling overheads.


Blocking/Unbuffered: no extra buffer available to hold pending sends/receives so must wait until message is sent to proceed
Blocked processors are idle, do no work, which cuts into speedup

Improper Send/Receive Order May Deadlock

```
1 // P0                                // P1
2 send(&a, 1, 1);          send(&a, 1, 0); // both blocked
3 receive(&b, 1, 1);      receive(&b, 1, 0);
```

- ▶ Above ordering is a bug: both processors block until the other receives a message
- ▶ They may both wait forever, a form of Deadlock
- ▶ On some platforms for some message sizes, the above may NOT produce deadlock due to send() being **buffered** by OS/Hardware

Blocking + Buffered Send/Receive



Commonly finite-sized hardware communications buffers are available that offload data transfer work from the main processor; decouples `send()`/`receive()` operations on processors

Buffers are no Panacea: Deadlocks Still Possible

Receive always Blocks

```
1 // P0          // P1  
3 receive(&a, 1, 1);    receive(&a, 1, 0);  
3 send(&b, 1, 1);    send(&b, 1, 0);
```

- ▶ `receive()` always blocks until message is obtained
- ▶ Does the above code work even in the buffered setting?

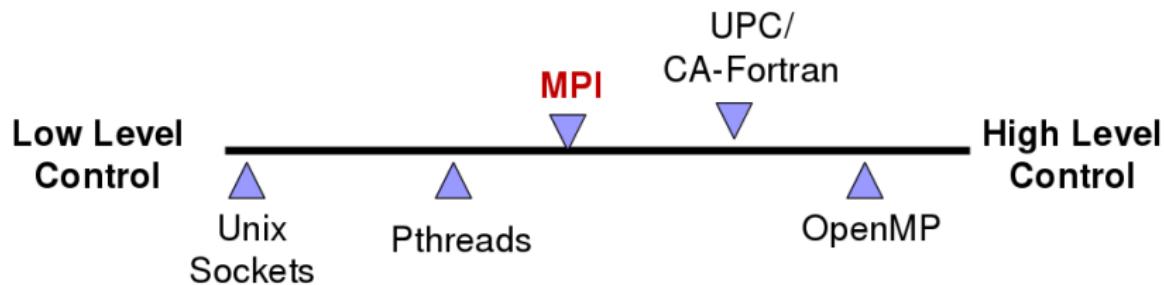
Finite Buffer Sizes

```
1 // P0          // P1  
2 send(&a, N, 1);    send(&a, N, 0);  
3 receive(&b, N, 1);    receive(&b, N, 0);
```

- ▶ For small N , does not deadlock due to buffers
- ▶ For large N , does deadlock due to finite buffer size

MPI: Message Passing Interface

- ▶ Standardized library of functions for C/C++/Fortran
- ▶ Communicate between processors in a distributed memory machine
- ▶ First appearing around 1992
- ▶ MPI Version 1.x universally deployed, Version 2.x less so
- ▶ Open Source Versions: MPICH, Open MPI
- ▶ Proprietary Versions: Intel, Platform, IBM, Platform, Cray
- ▶ Typically vendor configures MPI for particular architecture / network of a large-scale machine



MPI In a Nutshell: 6 Essential Functions

```
// Initialize and Terminate MPI
int MPI_Init(int *argc, char ***argv);
int MPI_Finalize();

// Get total number of processors
int MPI_Comm_size(MPI_Comm comm, int *size);

// Get logical proc number of calling process
int MPI_Comm_rank(MPI_Comm comm, int *rank);

// Send a message to dest processor
int MPI_Send(void *buf, int count, MPI_Datatype datatype,
            int dst_proc, int tag, MPI_Comm comm);

// Receive a message from source processor
int MPI_Recv(void *buf, int count, MPI_Datatype datatype,
             int src_proc, int tag, MPI_Comm comm,
             MPI_Status *status);
```

MPI Hello World

```
1 // mpi_hello.c: C Example of hello world with MPI. Compile and run as
2 // > mpicc -o mpi_hello mpi_hello.c
3 // > mpirun ./mpi_hello      # use number of processors equal to total machine
4 // > mpirun -np 2 mpi_hello  # use 2 processors
5 // > mpirun -np 8 mpi_hello  # use 8 processors
6
7 #include <stdio.h>
8 #include <mpi.h>
9
10 int main (int argc, char *argv[]){
11     int rank;                      // the id of this processor
12     int size;                     // the number of processors being used
13
14     MPI_Init (&argc, &argv);        // starts MPI
15     MPI_Comm_rank (MPI_COMM_WORLD, &rank); // get current process id
16     MPI_Comm_size (MPI_COMM_WORLD, &size); // get number of processes
17
18     // Say hello from this proc
19     printf( "Proc %d of %d says 'Hello world'\n", rank, size );
20
21     MPI_Finalize();
22
23 }
```

Compiling and Running

- ▶ Demo using openmpi implementation
- ▶ mpirun for interactive running
- ▶ mpirun -np 4 progr sets number of “processors” to 4

```
>> cd 04-mpi-code/
>> mpicc -o mpi_hello mpi_hello.c
>> ./mpi_hello
...
Proc 0 of 1 says 'Hello world'

>> mpirun -np 2 mpi_hello
...
Proc 0 of 2 says 'Hello world'
Proc 1 of 2 says 'Hello world'

>> mpirun mpi_hello
...
Proc 2 of 4 says 'Hello world'
Proc 0 of 4 says 'Hello world'
Proc 1 of 4 says 'Hello world'
Proc 3 of 4 says 'Hello world'
```

MPI Implementations and OpenMPI Warnings

- ▶ Several Implementations of MPI:
 - ▶ OpenMPI and MPICH are free, open source, widely available
 - ▶ HPC Vendors like IBM and Cray provide their own tailored MPI versions
- ▶ Recent Versions of OpenMPI can complain a LOT about various items missing; usually options / environment vars exist to suppress these warnings
 - ▶ Example: --mca btl_base_warn_component_unused 0 to warn about missing HPC network components during mpirun
 - ▶ Example: --mca opal_warn_on_missing_libcuda 0 if not intending to use GPU libraries
- ▶ Exact nature of warnings/errors varies a lot, look at messages which often dictate how to disable them
- ▶ Zaratan does not warn a lot as it seems configured well but course staff may provide specific instructions if warnings emerge

Warning Suppression in OpenMPI

```
>> mpicc mpi_hello_plus.c

# On the below machine, warnings abound
>> mpirun -np 2 a.out
-----
The library attempted to open the following supporting CUDA libraries,
but each of them failed. CUDA-aware support is disabled.
libcuda.so.1: cannot open shared object file: No such file or directory
libcuda.dylib: cannot open shared object file: No such file or directory
/usr/lib64/libcuda.so.1: cannot open shared object file: No such file or directory
/usr/lib64/libcuda.dylib: cannot open shared object file: No such file or directory
If you are not interested in CUDA-aware support, then run with
--mca opal_warn_on_missing_libcuda 0 to suppress this message. If you are interested
in CUDA-aware support, then try setting LD_LIBRARY_PATH to the location
of libcuda.so.1 to get passed this issue.
-----
P0000 [val]: Hello world from process    0 of 2
P0001 [val]: Hello world from process    1 of 2
[val:558294] 1 more process has sent help message help-mpi-common-cuda.txt / dlopen failed
[val:558294] Set MCA parameter "orte_base_help_aggregate" to 0 to see all help / error messages

# in BASH, sourcing provided script sets some environment variables to quiet things down
>> source mpiopts.sh

>> mpirun $MPIOPTS -np 2 a.out
P0001 [val]: Hello world from process    1 of 2
P0000 [val]: Hello world from process    0 of 2
```

MPI Oversubscribing

Default OpenMPI config uses all processors on a single machine,
fails for larger requests unless using option --oversubscribe

```
>> mpirun -np 2 a.out
...
P0001 [val]: Hello world from process    1 of 2
P0000 [val]: Hello world from process    0 of 2

>> mpirun -np 16 a.out
-----
There are not enough slots available in the system to satisfy the 16
slots that were requested by the application:

a.out

Either request fewer slots for your application, or make more slots
available for use.

...
Alternatively, you can use the --oversubscribe option to ignore the
number of available slots when deciding the number of processes to
launch.

-----
>> mpirun --oversubscribe -np 16 a.out
P0009 [val]: Hello world from process    9 of 16
...
P0014 [val]: Hello world from process    14 of 16
P0012 [val]: Hello world from process    12 of 16

# Newer MPI versions prefer `--map-by :OVERSUBSCRIBE` to `--oversubscribe`
```

Hostfiles

- ▶ For laptop/desktop MPI or on Zaratan, system config establishes the nodes with processors
- ▶ Some clusters instead use a **hostfile** in `mpirun` to indicate host names of other machines in cluster
- ▶ Simplest form of hostfile is a list of symbolic or IP addresses for machines to recruit for the run

```
myclust>> cat hostfile.txt
node01.mycluster.umd.edu
node02.mycluster.umd.edu
node03.mycluster.umd.edu
node04.mycluster.umd.edu
node05.mycluster.umd.edu
```

```
myclust>> mpirun -hostfile hostfile.txt -np 100 ./a.out
...
P0003 [node01]: Hello world from process      3 of 100
...
P0089 [node02]: Hello world from process      39 of 100
...
P0190 [node05]: Hello world from process      89 of 100
```

- ▶ Again, we won't need hostfiles as Zaratan is already configured for MPI

MPI On Zaratan

- ▶ UMD's HPC cluster is Zaratan, access via
>> ssh MYID@login.zaratan.umd.edu
- ▶ **Login Nodes** are login-1 login-2 login-3 and are meant for development, compilation, data manipulation, short ($\leq 1\text{min}$) tests
- ▶ Most nodes (including login) have 128 cores on them but are shared among all logged in useres
- ▶ Keep code runs on login short and sweet
- ▶ When experimenting with MPI, limit procs requested to **no more than 16**

```
profk@login-1 [04-mpi-code]% mpicc -o mpi_hello_plus mpi_hello_plus.c
profk@login-1 [04-mpi-code]% mpirun -np 16 mpi_hello_plus
P0011 [login-1.zaratan.umd.edu]: Hello world from process    11 of 16
P0013 [login-1.zaratan.umd.edu]: Hello world from process    13 of 16
P0014 [login-1.zaratan.umd.edu]: Hello world from process    14 of 16
...
```

- ▶ To request longer jobs and more processors, use the **SLURM Scheduler** discussed in more detail later

MPI Hello On Zaratan

```
# assumes correct environment configuration on Zaratan,
# that ~/.bashrc has lines below for mpicc / mpirun
#   source ~profk/bin/cmsc216-env
#   source ~profk/bin/cmsc416-env

>> ssh profk@login.zaratan.umd.edu
...
Last login: Thu Jan 29 17:18:51 2026
profk@login-1 [04-mpi-code]%
profk@login-1 [04-mpi-code]%
profk@login-1 [04-mpi-code]%
profk@login-1 [04-mpi-code]%
```

```
profk@login-1 [04-mpi-code]%
profk@login-1 [04-mpi-code]% mpirun -np 16 mpi_hello_plus
P0011 [login-1.zaratan.umd.edu]: Hello world from process    11 of 16
P0013 [login-1.zaratan.umd.edu]: Hello world from process    13 of 16
P0014 [login-1.zaratan.umd.edu]: Hello world from process    14 of 16
P0010 [login-1.zaratan.umd.edu]: Hello world from process    10 of 16
P0000 [login-1.zaratan.umd.edu]: Hello world from process     0 of 16
P0002 [login-1.zaratan.umd.edu]: Hello world from process     2 of 16
P0004 [login-1.zaratan.umd.edu]: Hello world from process     4 of 16
P0005 [login-1.zaratan.umd.edu]: Hello world from process     5 of 16
P0006 [login-1.zaratan.umd.edu]: Hello world from process     6 of 16
P0007 [login-1.zaratan.umd.edu]: Hello world from process     7 of 16
P0009 [login-1.zaratan.umd.edu]: Hello world from process     9 of 16
P0012 [login-1.zaratan.umd.edu]: Hello world from process    12 of 16
P0015 [login-1.zaratan.umd.edu]: Hello world from process    15 of 16
P0001 [login-1.zaratan.umd.edu]: Hello world from process     1 of 16
P0003 [login-1.zaratan.umd.edu]: Hello world from process     3 of 16
P0008 [login-1.zaratan.umd.edu]: Hello world from process     8 of 16
```

MPI Send and Recieve

MPI's most basic functionality is point-to-point message transfer via MPI_Send() / MPI_Recv()

```
1 {
2     int count = 5;
3     int a[count]={10,20,30,40,50};
4     int b[count];
5     int partner = 1;
6     int tag = 1;
7
8     // Send contents of a to partner proc with tag=1
9     MPI_Send(a, count, MPI_INT, partner, tag, MPI_COMM_WORLD);
10
11    // Receive message into b from partner proc
12    MPI_Recv(b, count, MPI_INT, partner, tag, MPI_COMM_WORLD,
13              MPI_STATUS_IGNORE); // ignore status of receipt
14 }
```

- ▶ Analyze the codepack program send_recv_demo.c
 - ▶ Study code then try several runs
 - ▶ Pick a number of procs that will cause it problems
- ▶ Compare with send_recv_bugs.c
 - ▶ Vary the ss and rr param to demo stalls
 - ▶ Vary the size of buffers sent to check implementation
 - ▶ Note Laptop vs Zaratan buffer sizes differ

Tags Make Messages Unique

```
int a[10], b[10], myrank;
MPI_Status status;

...
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
if (myrank == 0) {
    ////////////////////////////////////////////////// V : TAG to send
    MPI_Send(a, 10, MPI_INT, 1, 123, MPI_COMM_WORLD);
    MPI_Send(b, 10, MPI_INT, 1, 456, MPI_COMM_WORLD);
}
else if (myrank == 1) {
    ////////////////////////////////////////////////// V : TAG to receive
    MPI_Recv(b, 10, MPI_INT, 0, 456, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    MPI_Recv(a, 10, MPI_INT, 0, 123, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
}
```

- ▶ Send/Recv have an integer Tag parameter which must match between sender and receiver
- ▶ Above code may deadlock if not buffered due to the misordering of Tags
- ▶ Mostly we will use tag=1 for simplicity
- ▶ Alternatively MPI_ANY_TAG works for MPI_Recv()
- ▶ Receiver can capture the received status (last param) and query what tag was received (we won't bother with this)

MPI Data Types Supported

```
// Sends a message.  
int MPI_Send(void *buf, int count, MPI_Datatype datatype,  
             int dst_proc, int tag, MPI_Comm comm);  
  
// Receives a message.  
int MPI_Recv(void *buf, int count, MPI_Datatype datatype,  
             int src_proc, int tag, MPI_Comm comm,  
             MPI_Status *status);
```

- ▶ Buffer is always untyped (`void* buf`)
- ▶ To strive for slightly better safety, MPI has standard datatypes

| | |
|-------------------------|-----------------------------------------------|
| <code>MPI_CHAR</code> | <code>signed char</code> |
| <code>MPI_INT</code> | <code>signed int</code> |
| <code>MPI_LONG</code> | <code>signed long int</code> |
| <code>MPI_FLOAT</code> | <code>float</code> |
| <code>MPI_DOUBLE</code> | <code>double</code> |
| <code>MPI_BYTE</code> | Last two used for sending structure arrays |
| <code>MPI_PACKED</code> | |

Unsigned types also available

Sending Structs

Sending structs can be done via the MPI_BYTE type

```
{                                     // from send_structs.c
    typedef struct {
        double x;
        int a, b;
    } dint_t;
    dint_t mine[10] = { {.x=1.23, .a=5, .b=7}, {.x=...}, ... }
    ...
    // calculate data sizes "manually" just as is done in a malloc()
    MPI_Send(mine, 10*sizeof(dint_t), MPI_BYTE,
             partner, 1, MPI_COMM_WORLD);
}
```

- ▶ Simple and effective if all compute nodes **use the same binary layout** (the typical case)
- ▶ MPI also provides a (complex) method for situations where struct layout differs between nodes
- ▶ Must Dictate # of struct fields, types, and ordering into a MPI_Datatype and use MPI_Type_create_struct()
- ▶ Likely hurts performance if struct layout differs so will not discuss in detail

Exercise: Heat Transfer in MPI

- ▶ Discuss conversion of the following A1 code to an MPI version
- ▶ How is data in $H[] []$ divided up?
- ▶ Is communication required?
- ▶ How would one arrange `MPI_Send()` / `MPI_Recv()` calls?
- ▶ How much data needs to be transferred and between who?
- ▶ When the computation is finished, how can all data be printed / displayed / saved?

```
// Simulate the temperature changes for internal cells
for(t=0; t<max_time-1; t++){
    for(p=1; p<width-1; p++){
        double left_diff = H[t][p] - H[t][p-1];
        double right_diff = H[t][p] - H[t][p+1];
        double delta = -k*(left_diff + right_diff);
        H[t+1][p] = H[t][p] + delta;
    }
}
```


Some Patterns that occur in the problem

- ▶ Pair exchange of items: made easier with `MPI_Sendrecv()`
- ▶ Collecting final output for display: `MPI_Gather()`
 - ▶ Previewed here
 - ▶ Discussed in following lectures

Exchange: Sendrecv() for exchanging data between pairs

```
{  
    double send[10], recv[10]; int partner;  
    if(procid % 2 == 1){ // odd procs send left, receive left  
        partner = procid-1;  
        MPI_Send(send, 10, MPI_DOUBLE, partner, 1, MPI_COMM_WORLD);  
        MPI_Recv(recv, 10, MPI_DOUBLE, partner, 1, MPI_COMM_WORLD,  
                 MPI_STATUS_IGNORE);  
    }  
    else{ // even procs receive right, send right  
        partner = procid+1;  
        MPI_Recv(recv, 10, MPI_DOUBLE, partner, 1, MPI_COMM_WORLD,  
                 MPI_STATUS_IGNORE);  
        MPI_Send(send, 10, MPI_DOUBLE, partner, 1, MPI_COMM_WORLD);  
    }  
}  
{ // Sendrecv simplifies this pattern  
    double send[10], recv[10]; int partner;  
    partner = (procid % 2 == 1) ? procid-1 : procid+1;  
    MPI_Sendrecv(send, 10, MPI_DOUBLE, partner, 1,  
                recv, 10, MPI_DOUBLE, partner, 1,  
                MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
}
```

Take Care: Pair exchange can hang



```
{  
    double send[10], recv[10]; int partner;  
    partner = (procid % 2 == 1) ? procid-1 : procid+1;  
    MPI_Sendrecv(send, 10, MPI_DOUBLE, partner, 1,  
                recv, 10, MPI_DOUBLE, partner, 1,  
                MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
}
```

- ▶ With 9 processors, logic is broken
- ▶ Proc 8 will wait to communicate with absent partner
- ▶ Program never terminates
- ▶ Special checks for `procid == (nprocs-1)` can surmount this
- ▶ Heat problem is interesting as it involves communicating with a Left AND Right partner, yet more complex and part of A2...

Getting Answers to Root Proc: Gather

Before Call Begins

| P0 | P1 | P2 | P3 |
|-------------------------------|---------------------|---------------------|---------------------|
| send_buf[] 10 20 | send_buf[] 30 40 | send_buf[] 50 60 | send_buf[] 70 80 |
| recv_buf[] ? ? ? ? ? ? ? ? | recv_buf[] NULL | recv_buf[] NULL | recv_buf[] NULL |

`MPI_Gather(send_buf,2,MPI_INT, recv_buf,2,MPI_INT, 0,MPI_COMM_WORLD);`

| P0 | P1 | P2 | P3 |
|---------------------------------------|---------------------|---------------------|---------------------|
| send_buf[] 10 20 | send_buf[] 30 40 | send_buf[] 50 60 | send_buf[] 70 80 |
| recv_buf[] 10 20 30 40 50 60 70 80 | recv_buf[] NULL | recv_buf[] NULL | recv_buf[] NULL |

After Call Completes

- ▶ At the end of Heat, every processor has computed some columns
- ▶ One processor needs to gather all of the data for printing / saving, usually “root” processor with ID/rank 0
- ▶ Everyone calls `MPI_Gather()` to get data to root proc

MPI_Gather Sample

Use of Gather

```
// Preamble for any code
MPI_Comm comm = MPI_COMM_WORLD;
int sendarray[100];
int procid, total_procs, *rbuf;
...
// Only proc 0 needs space for to
// receive entirety of data
if(procid == 0) {
    rbuf = malloc(total_procs*100*
                  sizeof(int));
}
// Everyone calls gather
// proc 0 gets all data eventually
MPI_Gather(sendarray, 100, MPI_INT,
            rbuf, 100, MPI_INT,
            0, comm);
```

Equivalent Non-Gather Code

```
if(rank == 0){
    for(i=0; i<100; i++){
        rbuf[i] = sendarray[i];
    }
    for(i=1; i<total_procs; i++){
        int *rloc = &rbuf[i*100];
        MPI_Recv(rloc, 100,
                  MPI_INT, i,
                  tag, MPI_COMM_WORLD,
                  MPI_STATUS_IGNORE);
    }
}
else{
    MPI_Send(sendarray, 100,
              MPI_INT, 0,
              tag, MPI_COMM_WORLD);
}
```

Collective Communication Patterns Next

- ▶ `gather` is an example of a **Collective Communication Pattern**
- ▶ Will study more of these in subsequent lectures
- ▶ Using built-in collective comm. patterns simplifies programs
 - ▶ Without Collective Comm: hand-code proc 1 sends to proc 1/2 who receives, then to proc 1/4, then to...
 - ▶ With Collective Comm: `MPI_Reduce()`
- ▶ MPI implementation typically uses the most efficient underlying communications for a particular network

Non-Blocking Communication: “Immediate” Operations




Figure 6.4 Non-blocking non-buffered send and receive operations (a) in absence of communication hardware; (b) in presence of communication hardware.

- ▶ Non-Blocking calls return immediately
 - ▶ `isend(data,dest,request)`: send w/o waiting
 - ▶ `ireceive(data,dest,request)`: receive w/o waiting
- ▶ `wait(request)`: block until requested send/receive finishes, required to avoid race conditions
- ▶ Will revisit later if time permits:
More Control = More Potential Speedup + More Potential Bugs

Non-Blocking Send / Receive in MPI

| Effect | Blocking | Non-Blocking |
|---------|----------------------------|-------------------------------------------|
| Send | <code>MPI_Send(...)</code> | <code>MPI_Isend(..., &request)</code> |
| Receive | <code>MPI_Recv(...)</code> | <code>MPI_Irecv(..., &request)</code> |
| Sync | Automatic | <code>MPI_Wait(&request, ...)</code> |

- ▶ Non-blocking calls trigger send/receive to be initiated but do not block process(or) to completion
- ▶ `MPI_Request` struct tracks whether operation has completed
- ▶ Block via `MPI_wait()` until send/recv completes
- ▶ Prior to blocking, unsafe to alter/use data in buffers
- ▶ Can pair `MPI_ISend()` / `MPI_Recv()` and vice versa
- ▶ Allows for more overlap of computation and communication at the cost of more complexity

Faux Example of MPI_Isend() / MPI_IRecv()

```
1 int data_a[100] = {...};  
2 int data_b[100] = {...};  
3 int partner = ...;  
4 int tag = ...;  
5  
6 ...; // compute data_a[]  
7  
8 MPI_Request request;  
9 MPI_Isend(data_a, 100, MPI_INT, partner, tag, MPI_COMM_WORLD, &request);  
10  
11 ...; // unsafe to alter data_a[] so compute data_b[]  
12  
13 MPI_Wait(&request, MPI_STATUS_IGNORE); // block until data_a[] has been sent  
14 // now safe to alter data_a[]  
15 for(int i=0; i<100; i++){ // more computations on data_a[]  
16     data_a[i]++;  
17 }  
18 ...;  
19 MPI_Irecv(data_a, 100, MPI_INT, partner, tag, MPI_COMM_WORLD, &request);  
20  
21 ...; // unsafe to do anything with data_a[], compute on data_b[]  
22  
23 MPI_Wait(&request, MPI_STATUS_IGNORE); // block until data_a[] has been sent  
24 // now safe to alter data_a[]  
25 for(int i=0; i<100; i++){ // more computations on data_a[]  
26     data_a[i]++;  
27 }
```