

CSCI 2021: Binary, Integers, Arithmetic

Chris Kauffman

*Last Updated:
Mon Feb 8 09:45:09 AM CST 2021*

Logisitcs

Reading

- ▶ C references: whole language:
- ▶ Bryant/O'Hallaron Ch 2.1-3

Goals

- ▶ Binary Representations / Notation
- ▶ Integers in binary
- ▶ Arithmetic operations

Assignments

- ▶ P1 Ongoing
- ▶ Lab03 on File Input
- ▶ HW03 on Binary Ints

Date	Event
Mon 2/8	Binary Ints/Chars
Wed 2/10	Lec/Lab Review
	Project 1 Due
Fri 2/12	Exam 1
Mon 2/15	Bit-level Ops
Wed 2/17	Feedback Survey Due

NOTE: `font_load()` tutorial posted

Announcements: Course Feedback Survey on Canvas

- ▶ Will post a **Feedback Survey** tomorrow on Canvas, due a week later
- ▶ Anonymous, worth **1 Engagement Point** in place of Lab04
- ▶ Soliciting feedback on various aspects of the course, both numerically and free-form
- ▶ If you have had good/bad/ugly experiences so far, this will be your chance to weigh in

NOTE: An old version of the survey was accidentally left open on Canvas and has been deleted. If you took this version, you will still need to fill out the one released on Wed.

Unsigned Integers: Decimal and Binary

- ▶ Unsigned integers are always positive:
`unsigned int i = 12345;`
- ▶ To understand binary, recall how decimal numbers “work”

Decimal: Base 10 Example

Each digit adds on a power 10

$80,345 = 5 \times 10^0 +$	5 ones
$4 \times 10^1 +$	40 tens
$3 \times 10^2 +$	300 hundreds
$0 \times 10^3 +$	0 thousands
8×10^4	80,000 ...
$5 + 40 + 300 + 80,000$	

Binary: Base 2 Example

Each digit adds on a power 2

$11001_2 = 1 \times 2^0 +$	1 ones
$0 \times 2^1 +$	0 twos
$0 \times 2^2 +$	0 fours
$1 \times 2^3 +$	8 eights
$1 \times 2^4 +$	16 sixteens
$= 1 + 8 + 16 = 25$	

So, $11001_2 = 25_{10}$

Exercise: Convert Binary to Decimal

Base 2 Example:

$$\begin{aligned} 11001 &= 1 \times 2^0 + & 1 \\ &0 \times 2^1 + & 0 \\ &0 \times 2^2 + & 0 \\ &1 \times 2^3 + & 8 \\ &1 \times 2^4 + & 16 \\ &= 1 + 8 + 16 &= 25 \end{aligned}$$

So, $11001_2 = 25_{10}$

Try With a Pal

Convert the following two numbers from base 2 (binary) to base 10 (decimal)

► 111

► 11010

► 01100001

Answers: Convert Binary to Decimal

$$\begin{aligned}111_2 &= 1 \times 2^2 + 1 \times 2^1 + 1 \times 2^0 \\&= 1 \times 4 + 1 \times 2 + 1 \times 1 \\&= 7_{10}\end{aligned}$$

$$\begin{aligned}11010_2 &= 1 \times 2^4 + 1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 0 \times 2^0 \\&= 1 \times 16 + 1 \times 8 + 0 \times 4 + 1 \times 2 + 0 \times 1 \\&= 26_{10}\end{aligned}$$

$$\begin{aligned}01100001_2 &= 0 \times 2^7 + 1 \times 2^6 + 1 \times 2^5 + 0 \times 2^4 \\&\quad + 0 \times 2^3 + 0 \times 2^2 + 0 \times 2^1 + 1 \times 2^0 \\&= 0 \times 128 + 1 \times 64 + 1 \times 32 + 0 \times 16 \\&\quad + 0 \times 8 + 0 \times 4 + 0 \times 2 + 1 \times 1 \\&= 97_{10}\end{aligned}$$

Note: last example ignores leading 0's

The Other Direction: Base 10 to Base 2

Converting a number from base 10 to base 2 is easily done using repeated division by 2; keep track of **remainders**

Convert 124 to base 2:

$$124 \div 2 = 62 \qquad \text{rem } 0$$

$$62 \div 2 = 31 \qquad \text{rem } 0$$

$$31 \div 2 = 15 \qquad \text{rem } 1$$

$$15 \div 2 = 7 \qquad \text{rem } 1$$

$$7 \div 2 = 3 \qquad \text{rem } 1$$

$$3 \div 2 = 1 \qquad \text{rem } 1$$

$$1 \div 2 = 0 \qquad \text{rem } 1$$

- ▶ Last step got 0 so we're done.
- ▶ Binary digits are in **remainders in reverse**
- ▶ Answer: 1111100
- ▶ Check:

$$0 + 0 + 2^2 + 2^3 + 2^4 + 2^5 + 2^6 = 4 + 8 + 16 + 32 + 64 = 124$$

Decimal, Hexadecimal, Octal, Binary

- ▶ Numbers exist independent of any writing system
- ▶ Can write the same number in a variety of bases
- ▶ C provides syntax for most common bases used in computing

	Decimal	Binary	Hexadecimal	Octal
Base	10	2	16	8
Mathematical	125	1111101 ₂	7D ₁₆	175 ₈
C Prefix	None	0b...	0x..	0...
C Example	125	0b1111101	0x7D	0175

- ▶ **Hexadecimal** often used to express long-ish byte sequences
Larger than base 10 so for 10-15 uses letters A-F
- ▶ **Examine** `number_writing.c` and `table.c` for patterns
- ▶ **Expectation:** Gain familiarity with doing conversions between bases as it will be useful in practice

Hexadecimal: Base 16

- ▶ Hex: compact way to write bit sequences
- ▶ One byte is 8 bits
- ▶ Each hex character represents 4 bits
- ▶ **Each Byte is 2 hex digits**

Byte	Hex	Dec
0101 0111	$57 = 5 \cdot 16 + 7$	87
5 7		
0011 1100	$3C = 3 \cdot 16 + 12$	60
3 C=12		
1110 0010	$E2 = 14 \cdot 16 + 2$	226
E=14 2		

Hex to 4 bit equivalence

Dec	Bits	Hex
0	0000	0
1	0001	1
2	0010	2
3	0011	3
4	0100	4
5	0101	5
6	0110	6
7	0111	7
8	1000	8
9	1001	9
10	1010	A
11	1011	B
12	1100	C
13	1101	D
14	1110	E
15	1111	F

Exercise: Conversion Tricks for Hex and Octal

Examples shown in this week's HW, **What tricks are illustrated?**

Decimal	Byte = 8bits	Byte by 4	Hexadecimal
87	01010111	bin: 0101 0111 hex: 5 7	57 = $5 \cdot 16 + 7$ hex dec
60	00111100	bin: 0011 1100 hex: 3 C=12	3C = $3 \cdot 16 + 12$ hex dec
226	11100010	bin: 1110 0010 hex: E=14 2	E2 = $14 \cdot 16 + 2$ hex dec
Decimal	Byte = 8bits	Byte by 3	Octal
87	01010111	bin: 01 010 111 oct: 1 2 7	127 = $1 \cdot 8^2 + 2 \cdot 8 + 7$ oct dec
60	00111100	bin: 00 111 100 oct: 0 7 4	074 = $0 \cdot 8^2 + 7 \cdot 8 + 4$ oct dec
226	11100010	bin: 11 100 010 oct: 3 4 2	342 = $3 \cdot 8^2 + 4 \cdot 8 + 2$ oct dec

Answers: Conversion Tricks for Hex and Octal

- ▶ Converting between Binary and Hexadecimal is easiest when grouping bits by 4: each 4 bits corresponds to one hexadecimal digit

bin: 0101 0111	bin: 1110 0010
hex: 5 7	hex: E=14 2

- ▶ Converting between Binary and Octal is easiest when grouping bits by 3: each 3 bits corresponds to one octal digit

bin: 01 010 111	bin: 11 100 010
oct: 1 2 7	oct: 3 4 2

Unix Permissions with Octal

- ▶ Octal arises associated with **Unix file permissions**
- ▶ Every file has 3 permissions for 3 entities
- ▶ Permissions are true/false so a single bit will suffice
- ▶ Octal historically used for this
- ▶ `ls -l`: long list files, shows permissions
 - binary octal
 - 110110101 = 665
 - rw-rw-r-x somefile.txt
 - U G O
 - S R T
 - E O H
 - R U E
 - P R
- ▶ `chmod 665 somefile.txt`: change permissions of somefile.txt to those shown to the right
- ▶ `chmod 777 x.txt`: open to everyone
- ▶ Symbolic `chmod` invocations are often preferred

Readable `chmod` version:
`chmod u=rw,g=rw,o=rx somefile.txt`

Character Coding Conventions

- ▶ Would be hard for people to share words if they interpreted bits as letters differently
- ▶ **ASCII**: American Standard Code for Information Interchange
An old standard for bit/character correspondence
- ▶ 7 bits per character, includes upper, lower case, punctuation

Dec	Hex	Binary	Char	Dec	Hex	Binary	Char
65	41	01000001	A	78	4E	01001110	N
66	42	01000010	B	79	4F	01001111	O
67	43	01000011	C	80	50	01010000	P
68	44	01000100	D	81	51	01010001	Q
69	45	01000101	E	82	52	01010010	R
70	46	01000110	F	83	53	01010011	S
71	47	01000111	G	84	54	01010100	T
72	48	01001000	H	85	55	01010101	U
73	49	01001001	I	86	56	01010110	V
74	4A	01001010	J	87	57	01010111	W
75	4B	01001011	K	88	58	01011000	X
76	4C	01001100	L	89	59	01011001	Y
77	4D	01001101	M	90	5A	01011010	Z
91	5B	01011101	[97	61	01100001	a
92	5C	01011110	\	98	62	01100010	b

Unicode

- ▶ World: why can't I write
人
in my code/web address/email?

- ▶ America: ASCII has 128 chars.
Deal with it.

- ▶ World: Seriously?

- ▶ America: We invented
computers. 'Merica!

- ▶ World:



- ▶ America: ... Unicode?
- ▶ World: But my language takes
more bytes than American.
- ▶ America: Deal with it. 'Merica!

- ▶ ASCII Uses 7 bits per char,
limited to 128 characters
- ▶ UTF-8 uses **1-4 bytes per
character** to represent **many**
more characters
(1,112,064 *codepoints*)
- ▶ Uses 8th bit in a byte to
indicate extension to more than
a single byte
- ▶ Requires software to understand
coding convention allowing
broader language support
- ▶ ASCII is a proper subset of
UTF-8 making UTF-8
backwards compatible and
increasingly popular

Binary Integer Addition/Subtraction

Adding/subtracting in binary works the same as with decimal
EXCEPT that carries occur on values of 2 rather than 10

ADDITION #1

```
  1 11      <-carries
  0100 1010 = 74
+ 0101 1001 = 89
-----
  1010 0011 = 163
```

ADDITION #2

```
  1111   1 <-carries
  0110 1101 = 109
+ 0111 1001 = 121
-----
  1110 0110 = 230
```

SUBTRACTION #1

```
      ? <-carries
  0111 1001 = 121
- 0001 0011 =  19
-----
  VVVVVVVVVVVVVV
  VVVVVVVVVVVVVV
  VVVVVVVVVVVVVV
```

```
      x12 <-carries
  0111 0001 = 119
- 0001 0011 =  19
-----
  0110 0110 = 102
```

Two's Complement Integers: Representing Negative Values

- ▶ To represent negative integers, must choose a coding system
- ▶ **Two's complement** is the most common for this
- ▶ Alternatives exist
 - ▶ Signed magnitude: leading bit indicates pos (0) or neg (1)
 - ▶ One's complement: invert bits to go between positive negative
- ▶ Great advantage of two's complement: **signed and unsigned arithmetic are identical**
- ▶ Hardware folks only need to make one set of units for both unsigned and signed arithmetic

Summary of Two's Complement

Short explanation: most significant bit is associated with a negative power of two.

UNSIGNED BINARY

7654 3210 : position
ABCD EFGH : 8 bits
A: 0/1 * $+(2^7)$ *POS*
B: 0/1 * $+(2^6)$
C: 0/1 * $+(2^5)$
...
H: 0/1 * $+(2^0)$

UNSIGNED BINARY

7654 3210 : position
1000 0000 = +128
1000 0001 = +129
1000 0011 = +131
1111 1111 = +255
0000 0000 = 0
0000 0001 = +1
0000 0101 = +5
0111 1111 = +127

TWO's COMPLEMENT (signed)

7654 3210 : position
ABCD EFGH : 8-bits
A: 0/1 * $-(2^7)$ *NEG*
B: 0/1 * $+(2^6)$
C: 0/1 * $+(2^5)$
...
H: 0/1 * $+(2^0)$

TWO's COMPLEMENT (signed)

7654 3210 : position
1000 0000 = -128
1000 0001 = -127 = -128+1
1000 0011 = -125 = -128+1+2
1111 1111 = -1 = -128+1+2+4+...+64
0000 0000 = 0 [+127]
0000 0001 = +1
0000 0101 = +5
0111 1111 = +127

Two's Complement Notes

- ▶ Leading 1 indicates negative, 0 indicates positive
- ▶ All 0's = Zero
- ▶ Positive numbers are identical to unsigned

Conversion Trick

Positive -> Negative

- ▶ **Invert bits, Add 1**

Negative -> Positive

- ▶ **Invert bits, Add 1**

Same trick works both ways, implemented in hardware for the **unary minus** operator as in

`int y = -x;`

$\sim 1001\ 1000 = \text{negative, invert}$

0110 0111 = +103 inverted

+ 1

0110 1000 = +104 (original = -104)

$\sim 0110\ 1000$ pos to neg

1001 0111 inverted

+ 1

1001 1000 = -104

original bits

Add Pos/Neg should give 0

1 1111 <-carries

0110 1000 = +104

+ 1001 1000 = -104

x 0000 0000 = zero

Overflow

- ▶ Sums that exceed the representation of the bits associated with the integral type **overflow**
- ▶ Excess significant bits are **dropped**
- ▶ Addition can result in a sum smaller than the summands, even for two positive numbers (!?)
- ▶ Integer arithmetic in fixed bits is a mathematical **ring**

Examples of Overflow in 8 bits

ADDITION #3 OVERFLOW

```
1 1111 111 <-carries
  1111 1111 = 255
+ 0000 0001 =   1
-----
1 0000 0000 = 256
x drop 9th bit
-----
  0000 0000 = 0
```

ADDITION #4 OVERFLOW

```
1           1 <-carries
  1010 1001 = 169
+ 1100 0001 = 193
-----
1 0110 1010 = 362
x drop 9th bit
-----
  0110 1010 = 106
```

Underflow

- ▶ **Underflow** occurs in unsigned arithmetic when values go below 0 (no longer positive)
- ▶ Pretend that there is an extra significant bit to carry out subtraction
- ▶ Subtracting a positive integer from a positive integer may result in a **larger** positive integer (!!?)
- ▶ Integer arithmetic in fixed bits is a mathematical **ring**

Examples of 8-bit Underflow

SUBTRACTION #2 UNDERFLOW

?<-carries

0000 0000 = 0

- 0000 0001 = 1

VVVVVVVVVVVVVV

?<-carries

1 0000 0000 = 256 (pretend)

- 0000 0001 = 1

VVVVVVVVVVVVVV

x 2<-carries

0 1111 1110 = 256

- 0000 0001 = 1

1111 1111 = 255

Overflow and Underflow In C Programs

- ▶ See `over_under_flow.c` for demonstrations in a C program.
- ▶ **No runtime errors** for under/overflow
- ▶ Good for hashing and cryptography
- ▶ Bad for most other applications: system critical operations should use checks for over-/under-flow
- ▶ See textbook Arienne Rocket crash which was due to overflow of an integer converting from a floating point value
- ▶ At assembly level, there are condition codes indicating that overflow has occurred

Endianness: Byte ordering in Memory

- ▶ Single bytes like ASCII characters lay out sequentially in memory in increasing address
- ▶ Multi-byte entities like 4-byte ints require decisions on byte ordering
- ▶ We think of a 32-bit int like this

Binary: 0000 0000 0000 0000 0001 1000 1110 1001
 0 0 0 0 1 8 E 9
Hex : 000018E9
Decimal: 6377

- ▶ But need to assign memory addresses to each byte
 - ▶ Little Endian: least significant byte early
 - ▶ Big Endian: most significant byte early
- ▶ Example: Integer starts at address #1024

	Address			
LittleEnd:	#1027	#1026	#1025	#1024
Binary:	0000 0000	0000 0000	0001 1000	1110 1001
	0 0	0 0	1 8	E 9
BigEnd:	#1024	#1025	#1026	#1027
	Address			

Little Endian vs. Big Endian

- ▶ Most modern machines use **little endian** by default
- ▶ Processor may actually support big endian
- ▶ Both Big and Little Endian have engineering trade-offs
- ▶ At one time debated hotly among hardware folks: a la [Gulliver's Travels](#) conflicts
- ▶ Intel chips were little endian and “won” so set the basis for most modern use
- ▶ Big endian byte order shows up in **network programming**: sending bytes over the network is done in big endian ordering
- ▶ **Examine** `show_endianness.c` to see C code to print bytes in order
- ▶ Since most machines are little endian, will see bytes print in the reverse order usually think of them

Output of show_endianness.c

```
1 > cat show_endianness.c
2 // Show endiannes layout of a binary number in memory Most machines
3 // are little endian so bytes will print leas significant earlier.
4 #include <stdio.h>
5
6 int main(){
7     int bin = 0b0000000000000000000000001100011101001;    // 6377
8     //          |   |   |   |   |   |   |   |
9     //          0   0   0   0   1   8   e   9
10    printf("%d\n%x\n",bin,bin);                               // show decimal/hex of binary
11    unsigned char *ptr = (unsigned char *) &bin; // pointer to beginning of bin
12    for(int i=0; i<4; i++){                                   // print bytes of bin from low
13        printf("%x ", ptr[i]);                                // to high memory address
14    }
15    printf("\n");
16    return 0;
17 }
18 > gcc show_endianness.c
19
20 > ./a.out
21 6377
22 18e9
23 e9 18 0 0
```

Notice: num prints with value 18e9 but bytes appear in reverse order e9 18 when looking at memory

Integer Ops and Speed

- ▶ Along with Addition and Subtraction, **Multiplication and Division** can also be done in binary
- ▶ Algorithms are the same as base 10 but more painful to do by hand
- ▶ This pain is reflected in hardware speed of these operations
- ▶ The **Arithmetic and Logic Unit (ALU)** does integer ops in the machine
- ▶ A **clock** ticks in the machine at some rate like 3Ghz (3 billion times per second)

- ▶ Under ideal circumstances, typical ALU Op speeds are

Operation	Cycles
Addition	1
Logical	1
Shifts	1
Subtraction	1
Multiplication	3
Division	>30

- ▶ Due to disparity, it is worth knowing about **relation** between multiply/divide and **bitwise** operations
- ▶ Compiler often uses such tricks: shift rather than multiply/divide

Mangling bits puts hair on your chest

Below contrasts difference between logical and bitwise operations.

```
int x1 = 12 || 10; // truthy (Logical OR)
int xb = 12 | 10;  // 14      (Bitwise OR)
int y1 = 12 && 10; // truthy (Logical AND)
int yb = 12 & 10;  // 8       (Bitwise AND)
int zb = 12 ^ 10;  // 6       (Bitwise XOR)
int w1 = !12;      // falsey  (Logical NOT)
int wb = ~12;      // 3       (Bitwise NOT/INVERT)
```

- ▶ Bitwise ops evaluate on a per-bit level
- ▶ 32 bits for int, 4 bits shown

Bitwise OR	Bitwise AND	Bitwise XOR	Bitwise NOT
1100 = 12	1100 = 12	1100 = 12	
1010 = 10	& 1010 = 10	^ 1010 = 10	~ 1100 = 12
-----	-----	-----	-----
1110 = 14	1000 = 8	0110 = 6	0011 = 3

Bitwise Shifts

- ▶ **Shift** operations move bits within a field of bits
- ▶ Shift operations are
 - `x = y << k;` // left shift y by k bits, store in x
 - `x = y >> k;` // right shift y by k bits, store in x
- ▶ All integral types can use shifts: long, int, short, char
- ▶ **Not applicable** to pointers or floating point
- ▶ Examples in 8 bits

```
//          76543210
char x = 0b00010111; // 23
char y = x << 2;      // left shift by 2
// y = 0b01011100; // 92
// x = 0b00010111; // not changed
char z = x >> 3;      // right shift by 3
// z = 0b00000010; // 2
// x = 0b00010111; // not changed
char n = 0b10000000; // -128, signed
char s = n >> 4;      // right shift by 4
// s = 0b11111000; // -8, sign extension
// right shift >> is "arithmetic"
```

Shifty Arithmetic Tricks

- ▶ Shifts with add/subtract can be used instead of multiplication and division
- ▶ Turn on optimization: `gcc -O3 code.c`
- ▶ Compiler automatically does this if it thinks it will save cycles
- ▶ *Sometimes* programmers should do this but better to convince compiler to do it for you, **comment** if doing manually

Multiplication

```
//          76543210
char  x = 0b00001010;  // 10
char x2 = x << 1;       // 10*2
//  x2 = 0b00010100;   // 20
char x4 = x << 2;       // 10*4
//  x4 = 0b00101000;   // 40
char x7 = (x << 3)-x;   // 10*7
//  x7 = (x * 8)-x;    // 10*7
//  x7 = 0b01000110;   // 70
//          76543210
```

Division

```
//          76543210
char  y = 0b01101110;  // 110
char y2 = y >> 1;       // 110/2
//  y2 = 0b00110111;   // 55
char y4 = y >> 2;       // 110/4
//  y4 = 0b00011011;   // 27
char z = 0b10101100;   // -84
char z2 = z >> 2;       // -84/4
//  z2 = 0b11101011;   // -21
//  right shift sign extension
```

Exercise: Checking / Setting Bits

Use a combination of bit shift / bitwise logic operations to...

1. Check if bit `i` of `int x` is set (has value 1)
2. Clear bit `i` (set bit at index `i` to value 0)

Show C code for this

Answers: Checking / Setting Bits

1. Check if bit *i* of int *x* is set (has value 1)

```
int x = ...;
int mask = 1; // or 0b0001 or 0x01 ...
int shifted = mask << i; // shifted 0b00...010..00
if(x & shifted){          //          x & 0b10...010..01
    ...                  //          -----
}                        //          0b00...010..00
```

2. Clear bit *i* (set bit at index *i* to value 0)

```
int x = ...;
int mask = 1; // or 0b0001 or 0x01 ...
int shifted = mask << i; // shifted 0b00...010..00
int inverted = ~shifted; // inverted 0b11...101..11
x = x & inverted;        //          x & 0b10...010..01
...                     //          -----
                        //          0b10...000..01
```

Showing Bits

- ▶ `printf()` capabilities:
 - `%d` as Decimal
 - `%x` as Hexadecimal
 - `%o` as Octal
 - `%c` as Character
- ▶ No specifier for binary
- ▶ Can construct such with bitwise operations
- ▶ Code pack contains two codes to do this
 - ▶ `printbits.c`: single args printed as 32 bits
 - ▶ `showbits.c`: multiple args printed in binary, hex, decimal

- ▶ Showing bits usually involves shifting and bitwise AND &
- ▶ Example from `showbits.c`

```
#define INT_BITS 32

// print bits for x to screen
void showbits(int x){
    int mask = 0x1;
    for(int i=INT_BITS-1; i>=0; i--){
        int shifted_mask = mask << i;
        if(shifted_mask & x){
            printf("1");
        } else {
            printf("0");
        }
    }
}
```

Bit Masking

- ▶ Semi-common for functions to accept bit patterns which indicate true/false options
- ▶ Frequently makes use of bit **masks** which are constants associated with specific bits
- ▶ Example from earlier: Unix permissions might be...

```
#define S_IRUSR 0b100000000 // User   Read
#define S_IWUSR 0b010000000 // User   Write
#define S_IXUSR 0b001000000 // User   Execute
#define S_IRGRP 0b000100000 // Group  Read
...
#define S_IWOTH 0b000000010 // Others Write
#define S_IXOTH 0b000000001 // Others Execute
```

- ▶ Use them to create options to C functions like

```
int permissions = S_IRUSR|S_IWUSR|S_RGRP;
chmod("/home/kauffman/solution.zip",permissions);
```