

CSCI 2021: x86-64 Control Flow

Chris Kauffman

Last Updated:

Fri Mar 5 09:58:27 AM CST 2021

Logistics

Reading Bryant/O'Hallaron

- ▶ Ch 3.6: Control Flow
- ▶ Ch 3.7: Procedure calls

Goals

- ▶ Finish Assembly Basics
- ▶ Jumps and Control flow
- ▶ Comparison / Test Instructions
- ▶ Procedure calls

Lab07 / HW07

- ▶ Assembly Coding and debugging
- ▶ Chance to configure environment
- ▶ All techniques used in Project 3

Project 3: Post Today

- ▶ Problem 1: Battery Assembly Functions (50%)
- ▶ Problem 2: Binary Bomb via GDB (50%)

Control Flow in Assembly and the Instruction Pointer

- ▶ No high-level conditional or looping constructs in assembly
- ▶ Only `%rip`: Instruction Pointer or “Program Counter”: memory address of the next instruction to execute
- ▶ Don’t mess with `%rip` by hand: automatically increases as instructions execute so the next valid instruction is referenced
- ▶ Jump instructions modify `%rip` to go elsewhere
- ▶ Typically label assembly code with positions of instructions that will be the target of jumps
- ▶ **Unconditional Jump** Instructions always jump to a new location.
- ▶ **Comparison / Test** Instruction, sets EFLAGS bits indicating relation between registers/values
- ▶ **Conditional Jump** Instruction, jumps to a new location if certain bits of EFLAGS are set, ignored if bits not set

Examine: Loop Sum with Instruction Pointer (rip)

- ▶ Can see direct effects on rip in disassembled code
- ▶ rip increases corresponding to instruction length
- ▶ Jumps include address for next rip

```
// C Code equivalent
int sum=0, i=1, lim=100;
while(i<=lim){
    sum += i;
    i++;
}
return sum;
```

000000000000005fa <main>:

ADDR	HEX-OPCODES	ASSEMBLY	EFFECT ON RIP
5fa:	48 c7 c0 00 00 00 00	mov \$0x0,%rax	# rip = 5fa -> 601
601:	48 c7 c1 01 00 00 00	mov \$0x1,%rcx	# rip = 601 -> 608
608:	48 c7 c2 64 00 00 00	mov \$0x64,%rdx	# rip = 608 -> 60f
0000000000000060f <LOOP>:			
60f:	48 39 d1	cmp %rdx,%rcx	# rip = 60f -> 612
612:	7f 08	jg 61c <END>	# rip = 612 -> 614 OR 61c
614:	48 01 c8	add %rcx,%rax	# rip = 614 -> 617
617:	48 ff c1	inc %rcx	# rip = 617 -> 61a
61a:	eb f3	jmp 60f <LOOP>	# rip = 61a -> 60f
0000000000000061c <END>:			
61c:	c3	retq	# rip 61c -> return address

Disassembling Binaries

- ▶ Binaries hard to read on their own
- ▶ Many tools exist to work with them, notably `objdump` on Unix
- ▶ Can **disassemble** binary: show “readable” version of contents

```
> gcc -Og loop.s                                # COMPILE AND ASSEMBLE

> file a.out
a.out: ELF 64-bit LSB pie executable, x86-64, version 1 (SYSV),

> objdump -d a.out                             # DISASSEMBLE BINARY
a.out:      file format elf64-x86-64
...
Disassembly of section .text:
...
0000000000000119 <main>:
    1119:      48 c7 c0 00 00 00 00    mov     $0x0,%rax
    1120:      48 c7 c1 01 00 00 00    mov     $0x1,%rcx
    1127:      48 c7 c2 64 00 00 00    mov     $0x64,%rdx
000000000000012e <LOOP>:
    112e:      48 39 d1                  cmp     %rdx,%rcx
    1131:      7f 08                    jg      113b <END>
    1133:      48 01 c8                  add     %rcx,%rax
    1136:      48 ff c1                  inc     %rcx
    1139:      eb f3                    jmp     112e <LOOP>
000000000000013b <END>:
    113b:      c3                      retq
```

FLAGS: Condition Codes Register

- ▶ Most CPUs have a special register with “flags” for various conditions
- ▶ In x86-64 this register goes by the following names

Name	Width	Notes
FLAGS	16-bit	Most important bits in first 16
EFLAGS	32-bit	Name shown in gdb
RFLAGS	64-bit	Not used normally

- ▶ Bits in FLAGS register are **automatically** set based on results of other operations
- ▶ Pertinent examples with conditional execution

Bit	Abbrev	Name	Description
0	CF	Carry flag	Set if last op caused unsigned overflow
6	ZF	Zero flag	Set if last op yielded a 0 result
7	SF	Sign flag	Set if last op yielded a negative
8	TF	Trap flag	Used by gdb to stop after one ASM instruction
9	IF	Interrupt flag	1: handle hardware interrupts, 0: ignore them
11	OF	Overflow flag	Set if last op caused signed overflow/underflow

Comparisons and Tests

Set the EFLAGS register by using comparison instructions

Name	Instruction	Examples	Notes
Compare	cmpX B, A	cmpl \$1,%eax	Like if(eax > 1){...}
	Like: A - B	cmpq %rsi,%rdi	Like if(rdi > rsi){...}
Test	testX B, A	testq %rcx,%rdx	Like if(rdx & rcx){...}
	Like: A & B	testl %rax,%rax	Like if(rax){...}

- ▶ Immediates like \$2 must be the first argument B
- ▶ B,A are NOT altered with cmp/test instructions
- ▶ EFLAGS register IS changed by cmp/test to indicate less than, greater than, 0, etc.

EXAMPLES:

```
movl $5, %eax      # 5 = 0b1001
cmpl $1, %eax      # [    ] (No flags set)
cmpl $5, %eax      # [ZF  ] (0 -> zero flag)
cmpl $9, %eax      # [  SF] (negative -> sign flag)
```

```
testq $0b1000, %rax # [    ] (No flags set)
testq $0b0110, %rax # [ZF  ] (0 -> zero flag)
```

Jump Instruction Summary

All control structures implemented using combination of Compare/Test + Jump instructions.

<i>Instruction</i>	<i>Jump Condition</i>	<i>FLAGS</i>
jmp LAB	Unconditional jump	-
je LAB	Equal / zero	ZF
jz LAB		ZF
jne LAB	Not equal / non-zero	!ZF
jnz LAB		!ZF
js LAB	Negative ("signed")	SF
jns LAB	Nonnegative	!SF
jg LAB	Greater-than signed	!SF & !ZF
jge LAB	Greater-than-equal signed	!SF
jl LAB	Less-than signed	SF & !ZF
jle LAB	Less-than-equal signed	SF
ja LAB	Above unsigned	!CF & !ZF
jae LAB	Above-equal unsigned	!CF
jb LAB	Below unsigned	CF & !ZF
jbe LAB	Below-equal unsigned	CF
jmp *OPER	Unconditional jump to variable address	-

Examine: Compiler Comparison Inversion

- ▶ Often compiler inverts comparisons
- ▶ $i < n$ becomes `cmpX / jge` (jump greater/equal)
- ▶ $i == 0$ becomes `cmpX / jne` (jump not equal)
- ▶ This allows “true” case to fall through immediately
- ▶ Depending on structure, may have additional jumps
 - ▶ `if(){ .. }` usually has a single jump
 - ▶ `if(){} else {}` may have a couple

```
## Assembly translation of
## if(rbx >= 2){
##   rdx = 10;
## }
## else{
##   rdx = 5;
## }
## return rdx;
    cmpq  $2,%rbx      # compare: rbx-0
    jnl   .LESSTHAN    # goto less than
    ## if(rbx >= 2){
    movq  $10,%rdx     # greater/equal
    ## }
    jmp   .AFTER
.LESSTHAN:
    ## else{
    movq  $5,%rdx      # less than
    ## }
.AFTER:
    ## rdx is 10 if rbx >= 2
    ## rdx is 5 otherwise
    movq  %rdx,%rax
    ret
```

Exercise: Other Kinds of Conditions

Other Things to Look For

- ▶ `testl %eax,%eax` used to check zero/nonzero
- ▶ Followed by `je` / `jz` / `jne` / `jnz`
- ▶ Also works for NULL checks
- ▶ Negative Values, followed by `js` / `jns` (jump sign / jump no sign)

See `jmp_tests_asm.s`

- ▶ Trace the execution of this code
- ▶ Determine return value in `%eax`

Exercise: Other Kinds of Conditions

```
1  main:
2      movl    $0,%eax
3      movl    $5,%edi
4      movl    $3,%esi
5      movq    $0,%rdx
6      movl    $-4,%ecx
7
8      testl   %edi,%edi
9      jnz     .NONZERO
10     addl    $20,%eax
11
12     .NONZERO:
13         testl   %esi,%esi
14         jz      .FALSEY
15         addl    $30,%eax
16
17     .FALSEY:
18         testq   %rdx,%rdx
19         je      .ISNULL
20         addl    $40,%eax
21
22     .ISNULL:
23         testl   %ecx,%ecx
24         jns     .NONNEGATIVE
25         addl    $50,%eax
26
27     .NEGATIVE:
28         ret
```

Answers: Other Kinds of Conditions

```
1  ### From jmp_tests_asm_commented.s
2  main:
3      movl    $0,%eax        # eax is 0
4
5      movl    $5,%edi        # set initial vals
6      movl    $3,%esi        # for registers to
7      movl    $0,%edx        # use in tests
8      movl    $-4,%ecx
9
10     ## eax=0, edi=5, esi=3, edx=NULL, ecx=-4
11     testl   %edi,%edi       # any bits set?
12     jnz     .NONZERO        # jump on !ZF (zero flag), same as jne
13     ## if(edi == 0){
14     addl    $20,%eax
15     ## }
16     .NONZERO:
17     testl   %esi,%esi       # any bits set?
18     jz      .FALSEY        # jump on ZF same as je
19     ## if(esi){
20     addl    $30,%eax
21     ## }
22     .FALSEY:
23     testq   %rdx,%rdx       # any bits set
24     je      .ISNULL        # same as jz: jump on ZF
25     ## if(rdx != NULL){
26     addl    $40,%eax
27     ## }
28     .ISNULL:
29     testl   %ecx,%ecx       # sign flag set on test to indicate negative results
30     jns     .NEGATIVE      # jump on !SF (not sign flag/pos)
31     ## if(ecx < 0){
32     addl    $50,%eax
33     ## }
34     .NEGATIVE:
35     ret                ## eax is return value
```

cmov Family: Conditional Moves

- ▶ A family of instructions allows conditional movement of data into registers
- ▶ Can limit jumping in simple assignments

```
cmpq    %r8,%r9
cmovge  %r11,%r10  # if(r9 >= r8) { r10 = r11 }
cmovg   %r13,%r12  # if(r9 >  r8) { r12 = r13 }
```

- ▶ Note that condition flags are set on arithmetic operations
- ▶ cmpX is like subQ: both set FLAG bits the same
- ▶ Greater than is based on the SIGN flag indicating subtraction would be negative allowing the following:

```
subq    %r8,%r9    # r9 = r9 - r8
cmovge  %r11,%r10  # if(r9 >= 0) { r10 = r11 }
cmovg   %r13,%r12  # if(r9 >  0) { r12 = r13 }
```

Procedure Calls

Have seen basics so far:

```
main:
    ...
    call my_func # call a function
    ## arguments in %rdi, %rsi, %rdx, etc.
    ## control jumps to my_func, returns here when done
    ...

my_func:
    ## arguments in %rdi, %rsi, %rdx, etc.
    ...
    movl $0,%eax # set up return value
    ret          # return from function
    ## return value in %rax
    ## returns control to wherever it came from
```

Need several additional notions

- ▶ Control Transfer to called function?
- ▶ Return back to calling function?
- ▶ Stack alignment and conventions
- ▶ Register conventions

Procedure Calls Return to Arbitrary Locations

- ▶ call instructions always transfer control to start of `return_seven` at line 4/5, like `jmp` instruction which modifies `%rip`
- ▶ `ret` instruction at line 6 must transfer control to **different locations**
 1. call-ed at line 11
ret to line 12
 2. call-ed at line 17
ret to line 18

`ret` cannot be a normal `jmp`
- ▶ To enable return to multiple places, record a **Return Address** when call-ing, use it when ret-urning

```
1  ### return_seven_asm.s
2  .text
3  .global return_seven
4  return_seven:
5      movl    $7, %eax
6      ret     ## jump to line 12 OR 18??
7  .global main
8  main:
9      subq    $8, %rsp
10
11     call     return_seven  ## to line 5
12     leaq     .FORMAT_1(%rip), %rdi
13     movl     %eax, %esi
14     movl     $0, %eax
15     call     printf@PLT
16
17     call     return_seven  ## to line 5
18     leaq     .FORMAT_2(%rip), %rdi
19     movl     %eax, %esi
20     movl     $0, %eax
21     call     printf@PLT
22
23     addq     $8, %rsp
24     movl     $0, %eax
25     ret
26 .data
27 .FORMAT_1: .asciz "first: %d\n"
28 .FORMAT_2: .asciz "second: %d\n"
```

call / ret with Return Address in Stack

call Instruction

1. Push the “caller” **Return Address** onto the stack
Return address is for instruction after call
2. Change rip to first instruction of the “callee” function

ret Instruction

1. Set rip to Return Address at top of stack
2. Pop the Return Address off the stack shrinking stack




Figure: Bryant/O'Hallaron Fig 3.26 demonstrates call/return in assembly

return_seven_asm.s 1/2: Control Transfer with call

BEFORE CALL

return_seven:

```
0x55555555139 <return_seven>  mov    $0x7,%eax
0x5555555513e <return_seven+5> retq
```

main: ...

```
0x5555555513f <main>          sub    $0x8,%rsp
=> 0x55555555143 <main+4>      callq  0x55555555139 <return_seven>
0x55555555148 <main+9>      lea     0x2ee1(%rip),%rdi
0x5555555514f <main+16>     mov     %eax,%esi
```

(gdb) stepi

```
rsp = 0x7fffffff450 -> call -> 0x7fffffff448    # push on return address
rip = 0x55555555143 -> call -> 0x55555555139    # jump control to procedure
```

AFTER CALL

return_seven:

```
=> 0x55555555139 <return_seven>  mov    $0x7,%eax
0x5555555513e <return_seven+5> retq
```

main: ...

```
0x5555555513f <main>          sub    $0x8,%rsp
0x55555555143 <main+4>      callq  0x55555555139 <return_seven>
0x55555555148 <main+9>      lea     0x2ee1(%rip),%rdi
0x5555555514f <main+16>     mov     %eax,%esi
```

(gdb) x/gx \$rsp

```
0x7fffffff448: 0x000055555555148    # stack grew 8 bytes with call
                                # return address in main on stack
```

return_seven_asm.s 2/2: Control Transfer with ret

BEFORE RET

return_seven:

0x55555555139 <return_seven> mov \$0x7,%eax

=> 0x5555555513e <return_seven+5> retq

main: ...

0x5555555513f <main> sub \$0x8,%rsp

0x55555555143 <main+4> callq 0x55555555139 <return_seven>

0x55555555148 <main+9> lea 0x2ee1(%rip),%rdi

0x5555555514f <main+16> mov %eax,%esi

(gdb) x/gx \$rsp

0x7fffffff448: 0x000055555555148 # return address pointed to by %rsp

(gdb) stepi # EXECUTE RET INSTRUCTION

rsp = 0x7fffffff448 -> ret -> 0x7fffffff450 # pops return address off

rip = 0x5555555513e -> ret -> 0x55555555148 # sets %rip to return address

AFTER RET

return_seven:

0x55555555139 <return_seven> mov \$0x7,%eax

0x5555555513e <return_seven+5> retq

main: ...

0x5555555513f <main> sub \$0x8,%rsp

0x55555555143 <main+4> callq 0x55555555139 <return_seven>

=> 0x55555555148 <main+9> lea 0x2ee1(%rip),%rdi

0x5555555514f <main+16> mov %eax,%esi

(gdb) print \$rsp --> \$3 = 0x7fffffff450

Warning: `%rsp` is important for returns

- ▶ When a function is about to return `%rsp` MUST refer to the memory location of the return address
- ▶ `ret` uses value pointed to `%rsp` as the return address
- ▶ Major problems arise if this is not so
- ▶ Using `pushX` / `subq` instructions to extend stack during a function MUST be coupled with `popX` / `addq` instructions
- ▶ There are computer security issues associated stack-based return value we will discuss later

Stack Alignment

- ▶ According to the strict x86-64 ABI, must align `rsp` (stack pointer) to 16-byte boundaries when calling functions
- ▶ Will often see arbitrary pushes or subtractions to align
 - ▶ Always enter a function with 8-byte Return Address on the stack
 - ▶ Means that it is aligned to 8-byte boundary
- ▶ `rsp` changes must be undone prior to return

```
main:                                # enter with at 8-byte boundary
    subq    $8, %rsp                # align stack for func calls
    ...
    call    sum_range               # call function
    ...
    addq    $8, %rsp                # remove rsp change
    ret
```
- ▶ Failing to align the stack may work but may break
- ▶ Failing to “undo” stack pointer changes will likely result in return to the wrong spot : major problems

x86-64 Register/Procedure Convention

- ▶ Used by Linux/Mac/BSD/General Unix
- ▶ Params and return in registers if possible

Parameters and Return

- ▶ First 6 arguments are put into
 1. rdi / edi / di (arg 1)
 2. rsi / esi / si (arg 2)
 3. rdx / edx / dx (arg 3)
 4. rcx / ecx / cx (arg 4)
 5. r8 / r8d / r8w (arg 5)
 6. r9 / r9d / r9w (arg 6)
- ▶ Additional arguments are pushed onto the stack
- ▶ Return Value in rax / eax / ...

Caller/Callee Save

Caller save registers: alter freely

```
rax rcx rdx rdi rsi  
r8  r9  r10 r11
```

Callee save registers: must restore these on return

```
rbx rbp r12 r13 r14  
r15
```

Careful messing with stack pointer

```
rsp # stack pointer
```

Pushing and Popping the Stack

- ▶ If local variables are needed on the stack, can use push / pop for these
- ▶ pushX %reg: grow rsp (lower value), move value to top of main memory stack,
 - ▶ pushq %rax: grows rsp by 8, puts contents of rax at top
 - ▶ pushl \$25: grows rsp by 4, puts constant 5 at top of stack
- ▶ popX %reg: move value from top of main memory stack to reg, shrink rsp (higher value)
 - ▶ popl %eax: move (%rsp) to eax, shrink rsp by 4

```
main:
    pushq    %rbp                # save register, aligns stack
                                   # like subq $8,%rsp; movq %rbp,(%rsp)
    call     sum_range           # call function
    movl     %eax, %ebp          # save answer
    ...
    call     sum_range           # call function, ebp not affected
    ...
    popq     %rbp                # restore rbp, shrinks stack
                                   # like movq (%rsp),%rbp; addq $8,%rsp
    ret
```

Exercise: Local Variables which need an Address

Compare code in files

- ▶ `swap_pointers.c` : familiar C code for swap via pointers
- ▶ `swap_pointers_asm.s` : hand-coded assembly version

Determine the following

1. Where are local C variables `x`, `y` stored in assembly version?
2. Where does the assembly version “grow” the stack?
3. How are the values in `main()` passed as arguments to `swap_ptr()`?
4. Where does the assembly version “shrink” the stack?

Exercise: Local Variables which need an Address

```
1  #include <stdio.h>
2
3  void swap_ptr(int *a, int *b){
4      int tmp = *a;
5      *a = *b;
6      *b = tmp;
7      return;
8  }
9  int main(int argc, char *argv[]){
10     int x = 19;
11     int y = 31;
12     swap_ptr(&x, &y);
13     printf("%d %d\n", x, y);
14     return 0;
15 }
```

```
1  .text
2  .global swap_ptr
3  swap_ptr:
4      movl    (%rdi), %eax
5      movl    (%rsi), %edx
6      movl    %edx, (%rdi)
7      movl    %eax, (%rsi)
8      ret
9  .global main
10 main:
11      subq    $8, %rsp
12      movl    $19, (%rsp)
13      movl    $31, 4(%rsp)
14      movq    %rsp, %rdi
15      leaq    4(%rsp), %rsi
16      call    swap_ptr
17
18      leaq    .FORMAT(%rip), %rdi
19      movl    (%rsp), %esi
20      movl    4(%rsp), %edx
21      movl    $0, %eax
22      call    printf@PLT
23
24      addq    $8, %rsp
25      movl    $0, %eax
26      ret
27 .data
28 .FORMAT:
29     .asciz  "%d %d\n"
```


Answers: Local Variables which need an Address

1. Where are local C variables x,y stored in assembly version?
2. Where does the assembly version “grow” the stack?
3. How are the values in main() passed as arguments to swap_ptr()?

// C CODE

```
int x = 19, y = 31;
```

```
swap_ptr(&x, &y) // need main mem addresses for x,y
```

ASSEMBLY CODE

main:

main() function

```
    subq    $8, %rsp        # grow stack by 8 bytes
```

```
    movl    $19, (%rsp)     # move 19 to local variable x
```

```
    movl    $31, 4(%rsp)    # move 31 to local variable y
```

```
    movq    %rsp, %rdi      # address of x into rdi, 1st arg to swap_ptr()
```

```
    leaq    4(%rsp), %rsi   # address of y into rsi, 2nd arg to swap_ptr()
```

```
    call    swap_ptr        # call swap function
```

4. Where does the assembly version “shrink” the stack?

```
    addq    $8, %rsp        # shrink stack by 8 bytes
```

```
    movl    $0, %eax        # set return value
```

```
    ret
```

Diagram of Stack Variables

- ▶ Compiler determines if local variables go on stack
- ▶ If so, calculates location as $\text{rsp} + \text{offsets}$

<pre>1 // C Code: locals.c 2 int set_buf(char *b, int *s); 3 int main(){ 4 // locals re-ordered on 5 // stack by compiler 6 int size = -1; 7 char buf[16]; 8 ... 9 int x = set_buf(buf, &size); 10 ... 11 }</pre>	<table border="1" style="border-collapse: collapse; width: 100%; text-align: center;"><thead><tr><th style="padding: 5px;">REG</th><th style="padding: 5px;">VALUE</th><th style="padding: 5px;">Name</th></tr></thead><tbody><tr><td style="padding: 5px;">rsp</td><td style="padding: 5px;">#1024</td><td style="padding: 5px;">top of stack during main</td></tr><tr><td style="padding: 5px;">MEM</td><td style="padding: 5px;"></td><td style="padding: 5px;"></td></tr><tr><td style="padding: 5px;">...</td><td style="padding: 5px;">...</td><td style="padding: 5px;">...</td></tr><tr><td style="padding: 5px;">#1031</td><td style="padding: 5px;">h</td><td style="padding: 5px;">buf[3]</td></tr><tr><td style="padding: 5px;">#1030</td><td style="padding: 5px;">s</td><td style="padding: 5px;">buf[2]</td></tr><tr><td style="padding: 5px;">#1029</td><td style="padding: 5px;">u</td><td style="padding: 5px;">buf[1]</td></tr><tr><td style="padding: 5px;">#1028</td><td style="padding: 5px;">p</td><td style="padding: 5px;">buf[0]</td></tr><tr><td style="padding: 5px;">#1024</td><td style="padding: 5px;">-1</td><td style="padding: 5px;">size</td></tr></tbody></table>	REG	VALUE	Name	rsp	#1024	top of stack during main	MEM			#1031	h	buf[3]	#1030	s	buf[2]	#1029	u	buf[1]	#1028	p	buf[0]	#1024	-1	size
REG	VALUE	Name																										
rsp	#1024	top of stack during main																										
MEM																												
...																										
#1031	h	buf[3]																										
#1030	s	buf[2]																										
#1029	u	buf[1]																										
#1028	p	buf[0]																										
#1024	-1	size																										
<pre>1 ## EQUIVALENT ASSEMBLY 2 main: 3 subq \$24, %rsp # space for buf/size and stack alignment 4 movl \$-1, (%rsp) # old rip already in stack so: 20+4+8 = 32 5 # initialize buf and size: main line 6 6 leaq 0(%rsp), %rdi # address of size arg1 7 leaq 4(%rsp), %rsi # address of buf arg2 8 call set_buf # call function, aligned to 16-byte boundary 9 movl %eax,%r8 # get return value 10 ... 11 addq \$24, %rsp # shrink stack size</pre>																												

Summary of Procedure Calls: `ABC()` calls `XYZ()`

<code>ABC()</code>	Caller	<code>callq XYZ</code>	# <code>ABC</code> to <code>XYZ</code>
<code>XYZ()</code>	Callee	<code>retq</code>	# <code>XYZ</code> to <code>ABC</code>

1. `ABC()` “saves” any Caller Save registers it needs by either copying them into Callee Save registers or pushing them into the stack
2. `ABC()` places up to 6 arguments in `%rsi`, `%rdi`, `%rdx`, ..., remaining arguments in stack
3. `ABC()` ensures that stack is “aligned”: `%rsp` contains an address that is evenly divisible by 16
4. `ABC()` issues the `callq ABC` instruction which (1) grows the stack by subtracting 8 from `%rsp` and copies a return address to that location and (2) changes `%rip` to the starting address of `func`
5. `XYZ()` now has control: `%rip` points to first instruction of `XYZ()`
6. `XYZ()` may issue `pushX val` instructions or `subq N,%rsp` instructions to grow the stack for local variables
7. `XYZ()` may freely change Caller Save registers BUT Callee Save registers it changes must be restored prior to returning.
8. `XYZ()` must shrink the stack to its original position via `popX %reg` or `addq N,%rsp` instructions before returning.
9. `XYZ()` sets `%rax` / `%eax` / `%ax` to its return value if any.
10. `XYZ()` finishes, issues the `retq` instruction which (1) sets the `%rip` to the 8-byte return address at the top of the stack (pointed to by `%rsp`) and (2) shrinks the stack by doing `addq $8,%rsp`
11. `ABC()` function now has control back with `%rip` pointing to instruction after `call XYZ`; may have a return value in `%rax` register
12. `ABC()` must assume all Caller Save registers have changed

Historical Aside: Base Pointer rbp was Important

```
int bar(int, int, int);
int foo(void) {
    int x = callee(1, 2, 3);
    return x+5;
}
```

- ▶ 32-bit x86 / IA32 assembly used rbp as bottom of stack frame, rsp as top.
- ▶ Push all arguments onto the stack when calling changing both rsp and rbp
- ▶ x86-64: default rbp to general purpose register, not used for stack tracking

Old x86 / IA32 calling sequence: set both %esp and %ebp for function call foo:

```
    pushl %ebp                # modifying ebp, save it
## Set up for function call to bar()
    movl %esp,%ebp           # new frame for next function
    pushl 3                   # push all arguments to
    pushl 2                   # function onto stack
    pushl 1                   # no regs used
    call bar                  # call function, return val in %eax
## Tear down for function call bar()
    movl %ebp,%esp           # restore stack top: args popped
## Continue with function foo()
    addl 5,%eax               # add onto answer
    popl %ebp                 # restore previous base pointer
    ret
```