

CSCI 2021: Memory Systems

Chris Kauffman

*Last Updated:
Mon Nov 15 02:07:50 PM CST 2021*

Logistics

Reading Bryant/O'Hallaron

- ▶ Ch 4: Finish / Skim
- ▶ Ch 6: Memory

Goals

- ▶ Finish Arch
- ▶ 2D arrays
- ▶ Timing issues
- ▶ Memory efficient programs
- ▶ Permanent Storage

HW / Lab 10: Timing

- ▶ Observe Superscalar / Pipelined features in processor
- ▶ Observe differences that can result from small code changes
- ▶ Lecture will partially explain **why** timing differences occur

Project 4

Later this week

Architecture Performance

```
// LOOP 1
for(i=0; i<iters; i++){
    retA += delA;
    retB += delB;
}
*start = retA+retB;
```

```
// LOOP 2
for(i=0; i<iters; i++){
    retA += delA;
    retA += delB;
}
*start = retA;
```

- ▶ LOOP1 or LOOP2 faster?
- ▶ Why?

Exercise: 2D Arrays

- ▶ Several ways to construct “2D” arrays in C
- ▶ All must *embed* a 2D construct into 1-dimensional memory
- ▶ Consider the 2 styles below: **how will the picture of memory look different?**

```
// REPEATED MALLOC
// allocate
int rows=100, cols=30;
int **mat =
    malloc(rows * sizeof(int*));

for(int i=0; i<rows; i++){
    mat[i] = malloc(cols*sizeof(int));
}

// do work
mat[i][j] = ...

// free memory
for(int i=0; i<rows; i++){
    free(mat[i]);
}
free(mat);
```

```
// TWO MALLOCs
// allocate
int rows=100, cols=30;
int **mat =
    malloc(rows * sizeof(int*));
int *data =
    malloc(rows*cols*sizeof(int));
for(int i=0; i<rows; i++){
    mat[i] = data+i*cols;
}

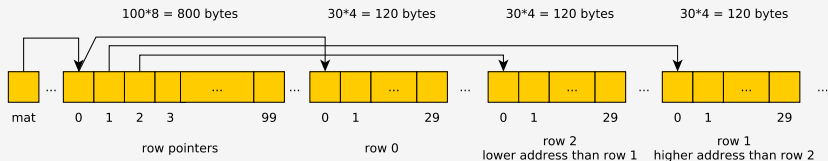
// do work
mat[i][j] = ...

// free memory
free(data);

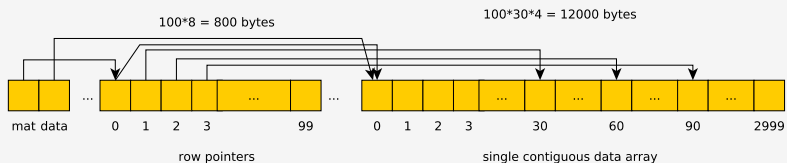
free(mat);
```

Answer: 2D Arrays

Repeated Mallocs



Two Mallocs



Single Malloc Matrices

Somewhat common to use a 1D array as a 2D matrix as in

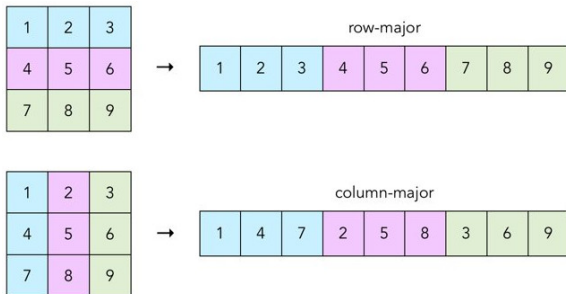
```
int *matrix =  
    malloc(rows*cols*sizeof(int));  
  
int i=5, j=20;  
int elem_ij = matrix[ i*cols + j ]; // retrieve element i,j
```

P4 will use this technique along with some structs and macros to make it more readable:

```
matrix_t mat;  
matrix_init(&mat, rows, cols);  
  
int elij = MGET(mat,i,j);  
// elij = mat.data[ mat.cols*i + j]  
  
MSET(mat,i,j, 55);  
// mat.data[ mat.cols*i + j ] = 55;
```

Aside: Row-Major vs Col-Major Layout

- ▶ Many languages use **Row-Major** order for 2D arrays/lists
 - ▶ C, Java, Python, Ocaml,...
 - ▶ `mat[i]` is a contiguous row, `mat[i][j]` is an element
- ▶ Numerically-oriented languages use **Column-Major** order
 - ▶ Fortran, Matlab/Octave, R, Ocaml (?)...
 - ▶ `mat[j]` is a contiguous **column**, `mat[i][j]` is an element
- ▶ Being aware of language convention can increase efficiency



Source: The Craft of Coding

Exercise: Matrix Summing

- ▶ How are the two codes below different?
- ▶ Are they doing the same number of operations?
- ▶ Which will run faster?

```
int sumR = 0;
for(int i=0; i<rows; i++){
    for(int j=0; j<cols; j++){
        sumR += mat[i][j];
    }
}
```

```
int sumC = 0;
for(int j=0; j<cols; j++){
    for(int i=0; i<rows; i++){
        sumC += mat[i][j];
    }
}
```


Answer: Matrix Summing

- ▶ Show timing in `matrix_timing.c`
- ▶ `sumR` faster than `sumC`: caching effects
- ▶ Discuss timing functions used to determine duration of runs

```
> gcc -Og matrix_timing.c
```

```
> a.out 50000 10000
```

```
sumR: 1711656320 row-wise CPU time: 0.265 sec, Wall time: 0.265
```

```
sumC: 1711656320 col-wise CPU time: 1.307 sec, Wall time: 1.307
```

- ▶ `sumR` runs about 6 times faster than `sumC`
- ▶ Understanding why requires knowledge of the memory hierarchy and cache behavior

Measuring Time in Code

- ▶ Measure CPU time with the standard `clock()` function; measure time difference and convert to seconds
- ▶ Measure Wall (real) time with `gettimeofday()` or related functions; fills struct with info on time of day (duh)

CPU Time

```
#include <time.h>

clock_t begin, end;
begin = clock(); // current cpu moment

do_something();

end = clock(); // later moment

double cpu_time =
    ((double) (end-begin)) / CLOCKS_PER_SEC;
```

Real (Wall) Time

```
#include <sys/time.h>

struct timeval tv1, tv2;
gettimeofday(&tv1, NULL); // early time

do_something();

gettimeofday(&tv2, NULL); // later time

double wall_time =
    ((tv2.tv_sec-tv1.tv_sec)) +
    ((tv2.tv_usec-tv1.tv_usec) / 1000000.0);
```

Tools to Measure Performance: perf

- ▶ The Linux `perf` tool is useful to measure performance of an entire program
- ▶ Shows variety of statistics tracked by the kernel about things like memory performance
- ▶ **Examine** examples involving the `matrix_timing` program: `sumR` vs `sumC`
- ▶ **Determine** statistics that explain the performance gap between these two?

Exercise: perf stats for sumR vs sumC, what's striking?

```
> perf stat $perfopts ./matrix_timing 8000 4000 row    ## RUN sumR ROW SUMMING
sumR: 1227611136 row-wise CPU time: 0.019 sec, Wall time: 0.019
```

```
Performance counter stats for './matrix_timing 8000 4000 row':          %SAMPLED
135,161,407    cycles:u                                              (45.27%)
417,889,646    instructions:u          #  3.09  insn per cycle      (56.22%)
 56,413,529    L1-dcache-loads:u                                           (55.96%)
  3,843,602    L1-dcache-load-misses:u #  6.81% of all L1-dcache hits (50.41%)
28,153,429    L1-dcache-stores:u                                           (47.42%)
   125         L1-icache-load-misses:u                                     (44.77%)
 3,473,211    cache-references:u      # last level of cache      (56.22%)
 1,161,006    cache-misses:u         # 33.427 % of all cache refs (56.22%)
```

```
> perf stat $perfopts ./matrix_timing 8000 4000 col    # RUN sumC COLUMN SUMMING
sumC: 1227611136 col-wise CPU time: 0.086 sec, Wall time: 0.086
```

```
Performance counter stats for './matrix_timing 8000 4000 col':          %SAMPLED
372,203,024    cycles:u                                              (40.60%)
404,821,793    instructions:u          #  1.09  insn per cycle      (57.23%)
 61,990,626    L1-dcache-loads:u                                           (60.21%)
 39,281,370    L1-dcache-load-misses:u # 63.37% of all L1-dcache hits (45.66%)
23,886,332    L1-dcache-stores:u                                           (43.24%)
   2,486       L1-icache-load-misses:u                                     (40.82%)
32,582,656    cache-references:u      # last level of cache      (59.38%)
 1,894,514    cache-misses:u         #  5.814 % of all cache refs (60.38%)
```

Answers: perf stats for sumR vs sumC, what's striking?

Observations

- ▶ Similar number of instructions between row/col versions
- ▶ #cycles lower for row version → higher insn per cycle
- ▶ **L1-dcache-misses**: marked difference between row/col version
- ▶ **Last Level Cache Refs** : many, many more in col version
- ▶ Col version: much time spent waiting for memory system to feed in data to the processor

Notes

- ▶ The right-side percentages like (50.41%) indicate how much of how much of the time this feature is measured; some items can't be monitored all the time.
- ▶ Specific perf invocation is in `10-memory-systems-code/measure-cache.sh`

Exercise: Time and Throughput

Consider the following simple loop to sum elements of an array from `stride_throughput.c`

```
int *data = ...; // global array
int sum_simple(int len, int stride){
    int sum = 0;
    for(int i=0; i<len; i+=stride)
    {
        sum += data[i];
    }
    return sum;
}
```

- ▶ Param `stride` controls step size through loop
- ▶ Interested in two features of the `sum_simple()` function:

1. Total Time to complete

2. **Throughput:**

$$\text{Throughput} = \frac{\# \text{Additions}}{\text{Second}}$$

- ▶ How would one **measure and calculate** these two in a program?
- ▶ As stride increases, **predict how Total Time and Throughput change**

Answers: Time and Throughput

Measuring Time/Throughput

Most interested in CPU time so

```
begin = clock();  
sum_simple(length,stride);  
end = clock();  
cpu_time = ((double) (end-begin))  
            / CLOCKS_PER_SEC;
```

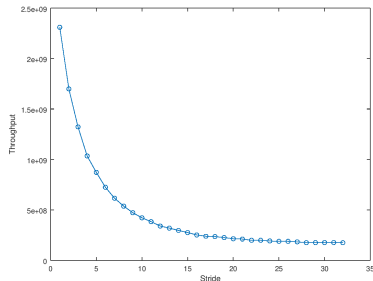
```
throughput = ((double) length) /  
             stride /  
             cpu_time;
```

Time vs Throughput

As stride increases...

- ▶ Time decreases: doing fewer additions (duh)
- ▶ Throughput **decreases**

Plot of Stride vs Throughput




- ▶ Stride = 1: consecutive memory accesses
- ▶ Stride = 16: jumps through memory, more time


Memory Mountains from Bryant/O'Hallaron

- ▶ Varying stride for a fixed length leads to decreasing performance, 2D plot
- ▶ Can also vary length for size of array to get a 3D plot
- ▶ Illustrates features of CPU/memory on a system
- ▶ The “Memory Mountain” on the cover of our textbook
- ▶ What **interesting structure** do you see?

CS:APP3e: Core i5-4440 (2013 Haswell)



Raspberry Pi 3B (2016 ARM Cortex-A53)



Increasing Efficiency

- ▶ Can increase the efficiency of loop summing with tricks
- ▶ B/O'H use multiple *accumulators*: multiple variables for summing
- ▶ Facilitates pipelining / superscalar processor
- ▶ Code is significantly faster BUT much trickier and less readable
- ▶ May be compiler options which enable this but not with defaults in gcc -O3 (try searching [optimization options](#), ~67 pages)

```
// From Bryant/O'Hallaron
int sum_add4(int elems, int stride){
    int i,
        sx1 = stride*1, sx2 = stride*2,
        sx3 = stride*3, sx4 = stride*4,
        acc0 = 0, acc1 = 0,
        acc2 = 0, acc3 = 0;
    int length = elems;
    int limit = length - sx4;

    /* Combine 4 elements at a time */
    for (i = 0; i < limit; i += sx4) {
        acc0 = acc0 + data[i];
        acc1 = acc1 + data[i+sx1];
        acc2 = acc2 + data[i+sx2];
        acc3 = acc3 + data[i+sx3];
    }

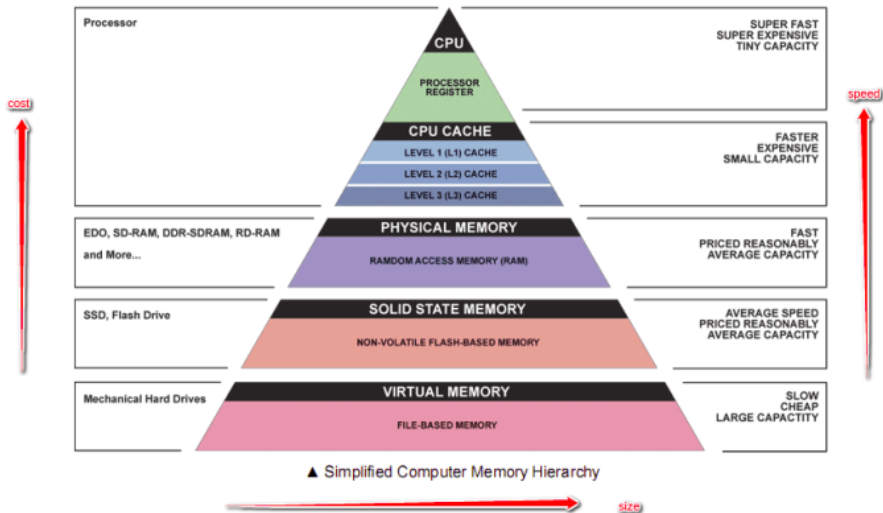
    /* Finish any remaining elements */
    for (; i < length; i += stride) {
        acc0 = acc0 + data[i];
    }
    return acc0+acc1+acc2+acc3;
}
```

Temporal and Spatial Locality

- ▶ In the beginning, there was only CPU and Memory
- ▶ Both ran at about the same speed (same clock frequency)
- ▶ CPUs were easier to make faster, began outpacing speed of memory
- ▶ Hardware folks noticed programmers often write loops like

```
for(int i=0; i<0; i++){  
    sum += array[i];  
}
```
- ▶ Led to development of faster memories exploit Locality
- ▶ **Temporal Locality**: memory recently used likely to be used again soon
- ▶ **Spatial Locality**: memory near to recently used memory likely to be used
- ▶ Register file and Cache were developed to exploit this: faster memory that is automatically managed

The Memory Pyramid



Source

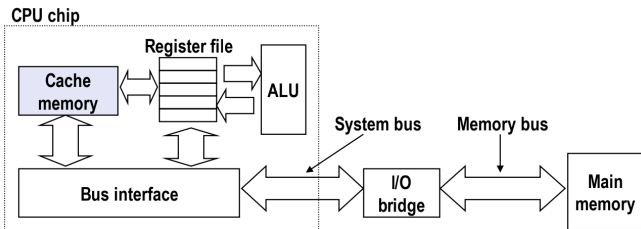
Numbers Everyone Should Know

Edited Excerpt of [Jeff Dean's](#) talk on data centers.


Reference	Time	Analogy
Register	-	Your brain
L1 cache reference	0.5 ns	Your desk
L2 cache reference	7 ns	Neighbor's Desk
Main memory reference	100 ns	This Room
Disk seek	10,000,000 ns	Salt Lake City

Big-O Analysis does NOT capture these; proficient programmers do

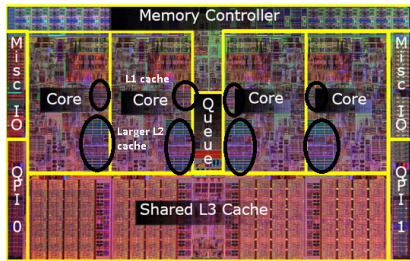
Diagrams of Memory Interface and Cache Levels



Source: Bryant/O'Hallaron CS:APP 3rd Ed.



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition



Source: SO "Where exactly L1, L2 and L3 Caches located in computer?"

Cache Principles: Hits and Misses

CPU-Memory is a Client-Server

- ▶ CPU makes requests
- ▶ Memory system services request as fast as possible

Cache Hit

- ▶ CPU requests memory at address 0xFFFF1234 be loaded into register %rax
- ▶ **Finds** valid data for 0xFFFF1234 in L1 Cache: **L1 Hit**
- ▶ Loads into register fast

Cache Miss

- ▶ CPU requests memory at address 0xFFFF7890 be loaded into register %rax
- ▶ 0xFFFF7890 **not in** L1
Cache: **L1 Miss**
- ▶ Search L2: if found move into L1, then %rax
- ▶ Search L3: if found move into L2, L1, %rax
- ▶ Search main memory: if found, move into caches, if not...

Wait, how could 0xFFFF7890 not be in main memory... ?

Types of Cache Misses

Compulsory “Cold” Miss: Program Getting Started

- ▶ All cache entries start with `valid=0`: cache contains leftover garbage from previous program runs
- ▶ After the cache “warms up” most entries will have `Valid=1`, data for running program

Capacity Miss: Data Too Big to Fit

- ▶ **Working set** is set of memory being frequently accessed in a particular phase of a program run
- ▶ Large working set may exceed the size of a cache causing misses


Conflict Miss: This Stall Occupied

- ▶ Internal **placement policy** of cache dictates where data goes
- ▶ If two needed piece of data both go to the same position in cache, leads to misses as they overwrite each other

Memory Address Determines Location in a Cache

Cache is like a **Hash Table**

- ▶ Cache has a # of **Sets** which can hold a copy of Main Memory
- ▶ Each Main Memory address has some bits indicating
 - ▶ **Set** - where in cache data should go
 - ▶ **Tag** - identifier to track what's in cache
- ▶ Each cache Set can hold 1 or more **Lines** of data with a specific Tag
- ▶ Main Memory divides into cache **Blocks** which share Tag/Set and move in/out of cache together



Address Bits to Cache Location

- ▶ Bits from address determine location for memory in cache
- ▶ Direct-Mapped cache, 4 sets and 16 byte blocks/lines
- ▶ Load address 0x28

```
      0  2  8
0x28 = 00 10 1000
      |  |  |
      |  |  +--> Offset: 4 bits
      |  +--> Set: 2 bits
      +--> Tag: Remaining bits
```

- ▶ 0x20 in the same line, will also be loaded into set #2

Exercises: Anatomy of a Simple CPU Cache

MAIN MEMORY

Addr	Addr Bits			Value
00	00	00	0000	331
08	00	00	1000	332
10	00	01	0000	333
18	00	01	1000	334
20	00	10	0000	335
28	00	10	1000	336
30	00	11	0000	337
38	00	11	1000	338
	..	.		
C0	11	00	0000	551
C8	11	00	1000	552
D0	11	01	0000	553
D8	11	01	1000	554
E0	11	10	0000	555
E8	11	10	1000	556
F0	11	11	0000	557
F8	11	11	1000	558
	Tag	Set	Offset	

CACHE

Set	V	Tag	Blocks/Line	
0-7			0-7	8-15
00	0	-	-	
01	1	00	333	334
10	1	11	555	556
11	1	00	337	338
			0-7	8-15

DIRECT-MAPPED Cache

- Direct-mapped: 1 Line per Set
- 16-byte lines = 4-bit offset
- 4 Sets = 2-bit index
- 8-bit Address = 2-bit tag
- Total Cache Size = 64 bytes
4 sets * 16 bytes

HITS OR MISSES? Show effects

1. Load 0x08
2. Load 0xF0
3. Load 0x18

Answers: Anatomy of a Simple CPU Cache

MAIN MEMORY

Addr	Addr Bits				Value
00	00	00	0000		331
08	00	00	1000		332
10	00	01	0000		333
18	00	01	1000		334
20	00	10	0000		335
28	00	10	1000		336
30	00	11	0000		337
38	00	11	1000		338
	..	.			
C0	11	00	0000		551
C8	11	00	1000		552
D0	11	01	0000		553
D8	11	01	1000		554
E0	11	10	0000		555
E8	11	10	1000		556
F0	11	11	0000		557
F8	11	11	1000		558
-----+-----+-----+-----+-----					
	Tag	Set	Offset		

CACHE

			Blocks/Line	
Set	V	Tag	0-7	8-15
00	1	*00	331	332
01	1	00	333	334
10	1	11	555	556
11	1	*11	557	558
-----+-----+-----+-----+-----				
			0-7	8-15

DIRECT-MAPPED Cache

- Direct-mapped: 1 line per set
- 16-byte lines = 4-bit offset
- 4 Sets = 2-bit index
- 8-bit Address = 2-bit tag
- Total Cache Size = 64 bytes
4 sets * 16 bytes

HITS OR MISSES? Show effects

1. Load 0x08: MISS to set 00
2. Load 0xF0: MISS overwrite set 11
3. Load 0x18: HIT in set 01
no change

Direct vs Associative Caches

Direct Mapped

One line per set

Set	V	Tag	Blocks/Line	
			0-7	8-15
00	0	-	-	-
01	1	00	333	334
10	1	11	555	556
11	1	00	337	338

- ▶ Simple circuitry
- ▶ **Conflict misses** may result: 1 slot for many possible tags
- ▶ **Thrashing:** need memory with overlapping tags

vv
0x10 = 00 01 0000 : in cache
0xD8 = 11 01 1000 : conflict
^^

N-Way Associative Cache

Ex: 2-way = 2 lines per set

Set	V	Tag	Blocks		
			0-7	8-15	
00	0	-	-	-	Line1
	1	11	551	552	Line2
01	1	00	333	334	Line1
	1	11	553	554	Line2
10	1	11	555	556	Line1
	0	-	-	-	Line2
11	1	00	337	338	Line1
	1	11	557	558	Line2

- ▶ Complex circuitry → \$\$
- ▶ Requires an **eviction policy**, usually least recently used

How big is your cache? Check Linux System special Files

lscpu Utility

Handy Linux program that summarizes info on CPU(s)

```
> lscpu
Architecture:      x86_64
CPU op-mode(s):    32-bit, 64-bit
Byte Order:        Little Endian
Address sizes:      36 bits physical,
                    48 bits virtual
CPU(s):            4
Vendor ID:         GenuineIntel
CPU family:        6
Model:             58
Model name:        Intel(R) Core(TM)
                   i7-3667U CPU @ 2.00GHz
...
L1d cache:         64 KiB
L1i cache:         64 KiB
L2 cache:          512 KiB
L3 cache:          4 MiB
Vulnerability Meltdown: Mitigation; ...
Vulnerability Spectre v1: Mitigation; ...
...
```

Detailed Hardware Info

Files under /sys/devices/...
show hardware info (caches)

```
> cd /sys/devices/system/cpu/cpu0/cache/
> ls
index0  index1  index2  index3 ...

> ls index0/
number_of_sets  type  level  size
ways_of_associativity ...

> cd index0
> cat level type number_* ways_* size
1 Data 64 8 32K

> cd ../index1
> cat level type number_* ways_* size
1 Instruction 64 8 32K

> cd ../index3
> cat level type number_* ways_* size
3 Unified 8192 20 10240K
```

Disks: Persistent Block Storage

- ▶ Have discussed a variety of fast memories which are **small**
- ▶ At the bottom of the pyramid are **disks**: slow but **large** memories
- ▶ These are **persistent**: when powered off, they retain information

Using Disk as Main Memory

- ▶ Operating Systems can create the illusion that main memory is larger than it is in reality
- ▶ Ex: 2 GB DRAM + 6 GB of disk space = 8 GB Main Memory
- ▶ Disk file is called **swap** or a **swap file**
- ▶ Naturally much slower than RAM so OS will try to limit its use
- ▶ A **Virtual Memory** system manages RAM/Disk as main memory, will discuss later in the course

Flavors of Permanent Storage

- ▶ Permanent storage often referred to as a “drive”
- ▶ Comes in many variants but these 3 are worth knowing about in the modern era
 1. Rotating Disk Drive
 2. Solid State Drive
 3. Magnetic Tape Drive
- ▶ Surveyed in the slides that follow

Ye Olde Rotating Disk

- ▶ Store bits “permanently” as magnetized areas on special platters
- ▶ Magnetic disks: moving parts → slow
- ▶ Cheap per GB of space

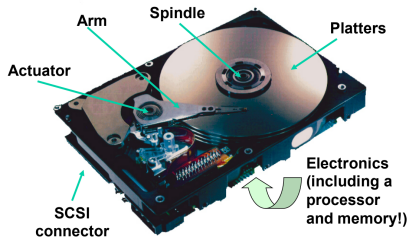
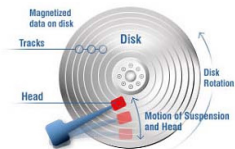


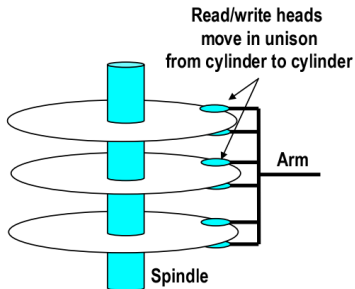
Image courtesy of Seagate Technology

Source: CS:APP Slides

HARD DRIVE DATA READ & WRITE OPERATION MOTION DIAGRAM



Source: Realtechs.net



Rotating Disk Drive Features of Interest

Measures of Quality

- ▶ Capacity: bigger is usually better
- ▶ Seek Time: delay before a head assembly reaches an arbitrary track of the disk that contains data
- ▶ Rotational Latency: time for disk to spin around to correct position; faster rotation → lower Latency
- ▶ Transfer Rate: once correct read/write position is found, how fast data moves between disk and RAM

Sequential vs Random Access

Due to the rotational nature of Magnetic Disks...

- ▶ Sequential reads/writes comparatively FAST
- ▶ Random reads/writes comparatively very SLOW


Solid State Drives

- ▶ No moving parts → speed
- ▶ Most use “flash” memory, non-volatile circuitry
- ▶ Major drawback: limited number of **writes**, disk wears out eventually
- ▶ Reads faster than writes
- ▶ Sequential somewhat faster than random access
- ▶ **Expensive:**
A 1TB internal 2.5-inch hard drive costs between \$40 and

A 1TB internal 2.5-inch hard drive costs between \$40 and \$50, but as of this writing, an SSD of the same capacity and form factor starts at \$250. That translates into

- 4 to 5 cents/GB for HDD
- 25 cents/GB for the SSD.

PC Magazine, "SSD vs HDD" by Tom Brant and Joel Santo Domingo March 26, 2018

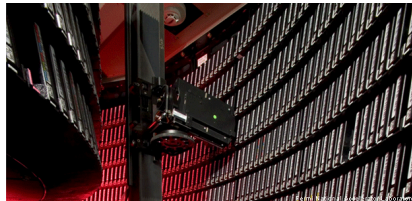


Tape Drives

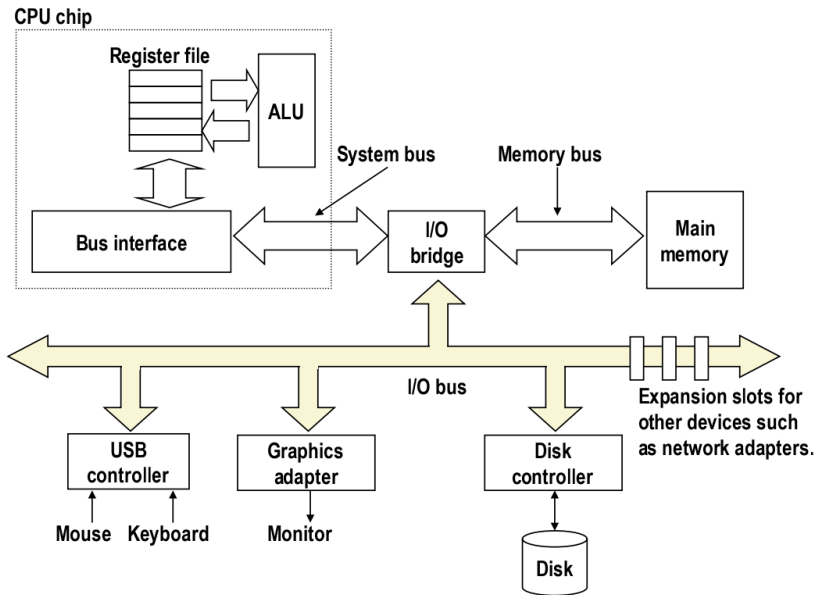
- ▶ Slowest yet: store bits as magnetic field on a piece of “tape” a la 1980’s cassette tape / video recorder



- ▶ Extremely cheap per GB so mostly used in backup systems
- ▶ Ex: CSELabs does nightly backups of home directories, recoverable from tape at request to Operator



The I/O System Connects CPU and Peripherals



Terminology

Bus A collection of wires which allow communication between parts of the computer. May be serial (single wire) or parallel (several wires), must have a communication protocol over it.

Bus Speed Frequency of the clock signal on a particular bus, usually different between components/buses requiring interface chips


CPU Frequency > Memory Bus > I/O Bus

Interface/Bridge Computing chips that manage communications across the bus possibly routing signals to correct part of the computer and adapting to differing speeds of components

Motherboard A printed circuit board connects to connect CPU to RAM chips and peripherals. Has buses present on it to allow communication between parts. *Form factor* dictates which components can be handled.

Memory Mapped I/O

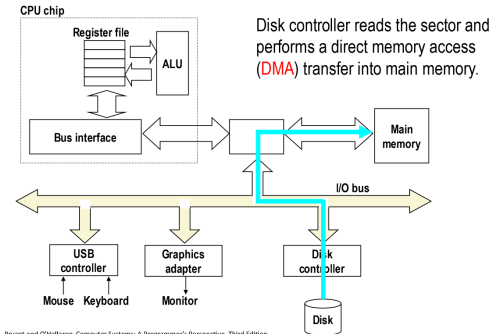
- ▶ Modern systems are a collection of devices and microprocessors
- ▶ CPU usually uses **memory mapped I/O**: read/write certain memory addresses translated to communication with devices on I/O bus



Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Direct Memory Access

- ▶ Communication received by *other* microprocessors like a Disk Controller or Memory Management Unit (MMU)
- ▶ Other controllers may talk: Disk Controller loads data directly into Main Memory via **direct memory access**

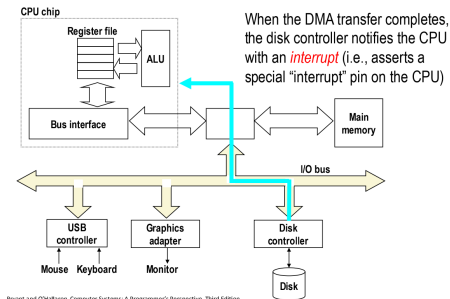


Interrupts and I/O

Recall access times

Place	Time
L1 cache	0.5 ns
RAM	100 ns
Disk	10,000,000 ns

- ▶ While running Program X, CPU reads an int from disk into %rax
- ▶ Communicates to disk controller to read from file
- ▶ Rather than wait, OS puts Program X to “sleep”, starts running program Y



- ▶ When disk controller completes read, signals the CPU via an **interrupt**, electrical signals indicating an event
- ▶ OS handles interrupt, schedules Program X as “ready to run”

Interrupts from Outside and Inside

- ▶ Examples of events that generate interrupts
 - ▶ Integer divide by 0
 - ▶ I/O Operation complete
 - ▶ Memory address not in RAM (Page Fault)
 - ▶ User generated: x86 instruction `int 80`
- ▶ Interrupts are mainly the business of the Operating System
- ▶ Usually cause generating program to immediately transfer control to the OS for handling
- ▶ When building your own OS, must write “interrupt handlers” to deal with above situations
 - ▶ Divide by 0: **signal** program usually terminating it
 - ▶ I/O Complete: schedule requesting program to run
 - ▶ Page Fault: sleep program until page loaded
 - ▶ User generated: perform system call
- ▶ User-level programs will sometimes get a little access to interrupts via **signals**, a topic for CSCI 4061