

CSCI 2021: Practice Final Exam SOLUTION

Fall 2021

University of Minnesota

Exam period: 20 minutes

Points available: 40

Background: Nearby are several C files along with two attempts to compile them on the left. Study these and answer the questions that follow.

```

1 > gcc vf_weak_var.c vf_strong_func.c vf_main.c # COMPILE 1
2 /usr/bin/ld: warning: size of symbol 'foo' changed from 4 to 14
3 /usr/bin/ld: warning: type of symbol 'foo' changed from 1 to 2
4 > file a.out
5 a.out: ELF 64-bit LSB pie executable, x86-64, version
6 > ./a.out
7 -573193927
8 4
9 > rm a.out
10
11 > gcc vf_strong_var.c vf_strong_func.c vf_main.c # COMPILE 2
12 /usr/bin/ld: multiple definition of 'foo';
13 collect2: error: ld returned 1 exit status
14 > file a.out
15 a.out: cannot open 'a.out' (No such file or directory)

```

```

1 // FILE: vf_main.c
2 #include <stdio.h>
3 int foo(int x);
4 int main(){
5     printf("%d\n",foo);
6     printf("%d\n",foo(2));
7     return 0;
8 }
9
10 // FILE: vf_strong_func.c
11 int foo(int x){
12     return 2*x;
13 }
14
15 // FILE: vf_strong_var.c
16 int foo = 0;
17
18 // FILE: vf_weak_var.c
19 int foo;

```

Problem 1 (10 pts): Why does COMPILE 1 succeed while COMPILE 2 fails? Mention pertinent properties of ELF files in your answer.

SOLUTION: COMPILE 1 succeeds because the integer `int foo` is uninitialized and therefore weak. It is overridden by the strong symbol `int foo(int x)` so the resulting ELF file has only the function version. COMPILE 2 fails as the C file initializes `int foo=0`; making both definitions strong. Two strong symbols with the same name cannot exist in an ELF file causing linking to fail.

Problem 2 (10 pts): Nearby is the output of `pmap` showing page table virtual memory mapping information for a running program called `memory_parts`. Answer the following questions about this output.

(A) The mapped memory references something called `libc-2.26.so`. Describe this entity and what kind of information you would expect to find at the mapped locations.

SOLUTION: This is the C standard library. It is a shared object with the `.so` extension and is likely to contain binary assembly instructions standard C functions like `printf()` and `malloc()`.

(B) Why does `pmap` only show a limited number of virtual addresses? What would happen if the program attempted to access an address not listed in the output? Example: address `0x00` is not in the listing.

*SOLUTION: The page table only contains mapped pages for program. These mapped addresses are what is shown. The large number of other addresses are unmapped. Attempting to access these unmapped addresses will result in errors such as **segmentation faults**; this usually causes the program to be immediately terminated.*

```

> pmap 7986
7986: ./memory_parts
00005579a4abd000      4K r-x-- memory_parts
00005579a4cbd000      4K r---- memory_parts
00005579a4cbe000      4K rw--- memory_parts
00005579a4cbf000      4K rw--- [ anon ]
00005579a53aa000     132K rw--- [ heap ]
00007f441f2e1000    1720K r-x-- libc-2.26.so
00007f441f48f000    2044K ----- libc-2.26.so
00007f441f68e000      16K r---- libc-2.26.so
00007f441f692000       8K rw--- libc-2.26.so
00007f441f694000      16K rw--- [ anon ]
00007f441f698000     148K r-x-- ld-2.26.so
00007f441f88f000       8K rw--- [ anon ]
00007f441f8bb000       4K r---- gettysburg.txt
00007f441f8bc000       4K r---- ld-2.26.so
00007f441f8bd000       4K rw--- ld-2.26.so
00007f441f8be000       4K rw--- [ anon ]
00007fff96ae1000     132K rw--- [ stack ]
00007fff96b48000      12K r---- [ anon ]
00007fff96b4b000       8K r-x-- [ anon ]
total                    4276K

```

Problem 3 (10 pts): We have seen that a common use of `mmap()` is to map files into the virtual memory space of a program to make it easy for them to be processed. However, this is only one of the uses for `mmap()` which is a fundamental tool for programs to interact with the Operating System and hardware. The Loader is the program which will take the disk image of a program like `a.out` and load it into memory to run. Discuss how `mmap()` can be used by the loader to place the required sections of ELF files in memory and establish areas such as the Stack and Heap for that program. In this, mention what important data structure about a program `mmap()` manipulates.

SOLUTION: `mmap()` manipulates the Page Table for a program, the OS data structure that maps virtual addresses that the program uses to physical addresses. This allows programs to directly map more memory into their address space. When loading a program like `a.out` to start it running, the Loader would likely locate the `.text` section of an `a.out` and `mmap()` it into the address space of a program as Executable to allow its code to be executed. Similarly, the Loader would map the `.data` section into pages marked Read/Write allowing those variables to change over the course of the program. In order for the program to have a Stack and Heap, the Loader would likely use `mmap()` to map Read/Write pages for these. Unlike the `.text` / `.data` sections, the Stack and Heap don't correspond to any part of the `a.out` ELF file: the Stack and Heap are only used at run time so aren't saved on disk. To that end, the space that is `mmap()`'d is "anonymous" and not associated with any file area.

Problem 4 (10 pts): New programmers are often surprised to learn that once an array is allocated, its size cannot be extended. In C code, this is easily observable as calling `malloc(16)` will yield a block of 16 bytes but there are no simple calls to expand this block of memory. Consider a proposed function for EL Malloc called `int el_expand_block(el_blockhead_t *boc)` which would expand a given block.

- (A) What conditions need to occur to for the function to succeed?
- (B) Why is it impossible to expand a block in some cases?

SOLUTION: Expanding an in-use block (one previously granted by a call to `el_malloc()` / `malloc()`) could occur if the block "above" (next highest in memory address) was still available. In that case, the block above could be shrunk to a smaller size allowing the existing in-use block to expand. This would be prevented by either the expansion size being larger than the size of the block above OR the block above not being available. Since in-use blocks cannot be moved around without invalidating pointers to them, there would be no way to expand an existing block in those cases and it is likely to arise as heap blocks are usually packed tightly together.

SIDEBAR: The `realloc()` function attempts to do just this: on reallocating to a larger size, it will expand the existing block if it is possible but move it if it is not possible to expand to the requested size; it is worth knowing about).