

# CSCI 4061: Inter-Process Communication

Chris Kauffman

*Last Updated:  
Tue Mar 30 11:43:49 AM CDT 2021*

# Logistics

## Reading

- ▶ Stevens/Rago Ch 15.6-12
- ▶ [Wikip: Dining Philosophers](#)

## Goals

- ▶ Project Plans
- ▶ File Append Problem
- ▶ Semaphore Basics
- ▶ Shared Memory
- ▶ Message Queues
- ▶ Dining Philosophers

Date	Event
Wed 3/31	IPC ShMem IPC MsgQ
Mon 4/5	Spring Break No Class
Mon 4/12	Review
Wed 4/14	Exam 2

## Lab 11

- ▶ Email lookup server/client
- ▶ Use of FIFO to communicate
- ▶ Difficult to write tests for it
  - sorry for any Gradescope problems
- ▶ How did it go?

# Project Plans

- ▶ Don't have time for 3 projects anymore which is Kauffman's fault  
*I apologize for this mistake. I have experienced some personal problems which have interfered with my ability to adequately prepare a solid Version Control project. I regret that I was not able to provide a project that puts the topics we have discussed into practical use.*
- ▶ P2: release after Exam 2
- ▶ Focus on Interprocess Communication: a local Chat Server/Client
- ▶ Same size as P1, Worth 20% of grade
- ▶ Opportunities for some Makeup Credit

## Exercise: Forms of IPC we've seen

- ▶ Identify as many forms of **inter-process communication** that we have studied as you can
- ▶ For each, identify **restrictions**
  - ▶ Must processes be related?
  - ▶ What must processes know about each other to communicate?
- ▶ You should be able to name at least 3-4 such mechanisms

## Answers: Forms of IPC we've seen

- ▶ Pipes
- ▶ FIFOs
- ▶ Signals
- ▶ Files
- ▶ Maybe `mmap()`'ed files

# Inter-Process Communication Libraries (IPC)

- ▶ Signals/FIFOs allow info transfer between unrelated processes
- ▶ Neither provides much
  - ▶ Communication synchronization between entities
  - ▶ Structure to data being communicated
  - ▶ Flexibility over access
- ▶ **Inter-Process Communication Libraries (IPC)** provide alternatives
  1. Semaphores: atomic counter + wait queue for coordination
  2. Message queues: direct-ish communication between processes
  3. Shared memory: array of bytes accessible to multiple processes

Two broad flavors of IPC that provide semaphores, message queues, shared memory...

# Which Flavor of IPC?

## System V IPC (XSI IPC)

- ▶ Most of systems have System V IPC but it's kind of strange, has its own *namespace* to identify shared things
- ▶ Part of Unix standards, referred to as **XSI IPC** and may be listed as optional
- ▶ Most textbooks/online sources discuss some System V IPC. Example:
  - ▶ Stevens/Rago 15.8 (semaphores)
  - ▶ Robbins/Robbins 15.2 (semaphore sets)
  - ▶ [Beej's Guide to IPC](#)

## POSIX IPC

- ▶ POSIX IPC little more regular, uses filesystem to identify IPC objects
- ▶ Originated as optional POSIX/SUS extension, now required for compliant Unix
- ▶ Covered in our textbooks partially. Example:
  - ▶ Stevens/Rago 15.10 POSIX Semaphores
  - ▶ Robbins/Robbins 14.3-5 POSIX Semaphores
- ▶ [Additional differences on StackOverflow](#)

**We will favor POSIX**

## Exercise: Concurrent Appends to a File

C code to append to a file some number of times.

```
1 // append_loop.c
2 int main(int argc, char *argv[]){
3     char *filename = argv[1];
4     int count = atoi(argv[2]);
5     int key = atoi(argv[3]);
6     int fd = open(filename,
7                   O_CREAT | O_RDWR ,
8                   S_IRUSR | S_IWUSR);
9
10    char line[128];
11    sprintf(line,"%04d\n",key);
12    int len = strlen(line);
13
14    for(int i=0; i<count; i++){
15        lseek(fd, 0, SEEK_END);
16        write(fd, line, len);
17    }
18    close(fd);
19    return 0;
20 }
```

Shell code demos its use. **What's wrong** with the last count?

```
> ./a.out
usage: ./a.out <filename> <count> <key>
> ./a.out thefile.txt 100 5555
> wc -l thefile.txt
100 thefile.txt
> ./a.out thefile.txt 100 7777
> wc -l thefile.txt
200 thefile.txt
> sort thefile.txt | uniq -c
    100 5555
    100 7777

> rm thefile.txt
> for i in $(seq 10); do
    ./a.out thefile.txt 100 $i &
done
> wc -l thefile.txt
732 thefile.txt
```



# Concurrency Principles

## Atomic Action

- ▶ Cannot be divided; will run completely before any other action taken. Some system calls are atomic like ...
- ▶ `nbytes = write(fd, data, len);` is atomic: `nbytes` of data written in sequence, data from other `write()` calls before/after but NOT in the middle
- ▶ `lseek()` is atomic: modifies file position in kernel data structure

## Race Condition

- ▶ Outcome depends on the ordering of unpredictable events such as the OS scheduler interrupting a process
- ▶ Race Conditions are **bad**: unlucky timing causes unpredictable behavior, bugs that only occasionally occur

## Race Condition in append\_loop.c 1 / 2

FILE	PROC1 key=5555	PROC2 key=7777
len=15		
5555		
5555	lseek(fd, 0, SEEK_END);	
7777	// pos = 15	
	<-----write(fd, line, len);	
len=20		
5555		
5555		
7777		lseek(fd, 0, SEEK_END);
5555		// pos = 20
	<-----	write(fd, line, len);

All appears well BUT cannot guarantee that `lseek()` / `write()` happen uninterrupted

- ▶ Individually atomic
- ▶ Combination is not

## Race Condition in append\_loop.c 1 / 2

FILE	PROC1 key=5555	PROC2 key=7777
len=25		
5555		lseek(fd, 0, SEEK_END);
5555		// pos = 25
7777	lseek(fd, 0, SEEK_END);	
7777	// pos = 25	
	<-----write(fd, line, len);	
len=30		
5555		
5555		
7777		
7777		// pos = 25
5555	<-----	write(fd, line, len);
len=30		
5555		
5555		
7777		
7777		
7777	# Overwritten	

Result: 1 line is lost as the `lseek()` between process is not coordinated

## Exercise: Solve this with Current IPC

Suggest a way to solve this problem with current IPC mechanisms

*Start an arbitrary number of processes. Each repeatedly appends a given key to a given file. All keys must be present at the end.*

- ▶ Describe new / old processes
- ▶ Describe new / old code and IPC to be used

*Hint: where have we recently seen a bunch of entities that all want access to data? How were these requests coordinated?*

# Answers: Solve this with Current IPC

*Use a FIFO to coordinate multiple writers*

## Manager Process

- ▶ Only the manager writes to `thefile.txt`
- ▶ Starting the manager creates a FIFO; manager `read()`'s from the FIFO, appends text to the end of the file

## Writer Processes

- ▶ Writer processes write into the FIFO (not `thefile.txt`)
- ▶ FIFOs automatically serialize data: no chance for loss as OS controls the singular read/write positions

## Familiar but Unsatisfactory

- ▶ Similar to `em_server` / `em_client` from Lab/HW
- ▶ Works and requires now new IPC mechanisms BUT...
- ▶ Dissatisfying: **must split code into manager/writer**

Would like a more straight-forward solution if possible

# Locking the Critical Region

## Critical Region

- ▶ Code sequence `lseek(); write()` is a **Critical Region**: not atomic, unsafe to have multiple entities in it at the same time
- ▶ Typically protect these with a coordination mechanism, a **lock** for the critical region

## OS Locking Mechanisms

- ▶ **Semaphore**: general purpose locking mechanism associated with multi-process programming
- ▶ **Mutex**: locking mechanism associated with threaded programming
- ▶ **File Locks**: lock all or portions of a file, always

# Semaphore History



Source: Wikipedia Railway Semaphore Signal

## **Semaphore: *noun***

A system of sending messages by holding the arms or two flags or poles in certain positions...

– Oxford Dictionary

## **Semaphore: (*computing*)**

In computer science, a semaphore is a variable or abstract data type used to control access to a common resource by multiple processes and avoid critical section problems in a concurrent system such as a multitasking operating system.

The semaphore concept was invented by Dutch computer scientist [Edsger Dijkstra](#)...

– Wikipedia

# Semaphore Basics: 3 Parts

## Counter Variable variable

Semaphores have an integer value indicating how much of a resource is available

- ▶  $S=0$ : none left
- ▶  $S>0$ : some available

Most common case is  $S=1$  (available) or  $S=0$  (in-use)

## Atomic Operations

- ▶ **Acquire**: If  $S>0$ , decrement; Else, enter wait-queue and block
- ▶ **Release**: Increment  $S$ , notify wait-queue of availability

## Wait Queue

Modern semaphores include a wait-queue. If  $S==0$ , **Acquire** will cause an entity (process) to enter the wait-queue and **block**.



# Posix Implementation of Semaphores

```
sem_t *sem =
    sem_open("/the_sem", O_CREAT, S_IRUSR | S_IWUSR);
// abstract type sem_t representing semaphores
// file-like semantics with open, semaphore name, flags, permissions

// Note: "the_sem" may or may not appear in the file system somewhere
// Under Linux, will be at /dev/shm/the_sem

sem_init(sem, 1, 1);    // Initialize the semaphore value
//                      | +-----> Initial counter value = 1
//                      +-----> Share among Processes (1: Processes, 0: Threads)

sem_wait(sem);
// ACQUIRE the semaphore; block and queue up if not available

// CRITICAL REGION

sem_post(sem);
// RELEASE the semaphore; notifies any queued processes of availability

sem_close(sem);
// file-like semantics: close when process is finished using it

sem_unlink("/the_sem");
// POSIX named semaphores have kernel persistence: if not removed by
// sem_unlink(), a semaphore will exist until the system is shut down.
```

## Examine: `append_file_sem.c`

Examine and experiment with `append_file_sem.c` which solves coordinates appends using a POSIX semaphore.

Look for use of semaphore functions like

- ▶ Opening
- ▶ Unlinking, initializing
- ▶ Acquiring / Releasing
- ▶ How the critical region is protected

```
> gcc -g append_loop_sem.c -lpthread
> ./a.out -init 1 1
initializing

> for i in $(seq 10); do
    ./a.out thefile.txt 100 $i &
done

> wc -l thefile.txt
1000 thefile.txt      # ALL THERE!

> sort thefile.txt |uniq -c
    100 0001          # ALL KEYS
    100 0002          # FROM ALL
    100 0003          # PROCESSES
    100 0004
    100 0005
    100 0006
    100 0007
    100 0008
    100 0009
    100 0010

> ./a.out -unlink 1 1
unlinking
```