

# CSCI 4061: Inter-Process Communication

Chris Kauffman

*Last Updated:  
Wed Mar 24 03:26:48 PM CDT 2021*

## Exercise: Forms of IPC we've seen

- ▶ Identify as many forms of **inter-process communication** that we have studied as you can
- ▶ For each, identify **restrictions**
  - ▶ Must processes be related?
  - ▶ What must processes know about each other to communicate?
- ▶ You should be able to name at least 3-4 such mechanisms

## Answers: Forms of IPC we've seen

- ▶ Pipes
- ▶ FIFOs
- ▶ Signals
- ▶ Files
- ▶ Maybe `mmap()`'ed files

# Inter-Process Communication Libraries (IPC)

- ▶ Signals/FIFOs allow info transfer between unrelated processes
- ▶ Neither provides much
  - ▶ Communication synchronization between entities
  - ▶ Structure to data being communicated
  - ▶ Flexibility over access
- ▶ **Inter-Process Communication Libraries (IPC)** provide alternatives
  1. Semaphores: atomic counter + wait queue for coordination
  2. Message queues: direct-ish communication between processes
  3. Shared memory: array of bytes accessible to multiple processes

Two broad flavors of IPC that provide semaphores, message queues, shared memory...

# Which Flavor of IPC?

## System V IPC (XSI IPC)

- ▶ Most of systems have System V IPC but it's kind of strange, has its own *namespace* to identify shared things
- ▶ Part of Unix standards, referred to as **XSI IPC** and may be listed as optional
- ▶ Most textbooks/online sources discuss some System V IPC. Example:
  - ▶ Stevens/Rago 15.8 (semaphores)
  - ▶ Robbins/Robbins 15.2 (semaphore sets)
  - ▶ [Beej's Guide to IPC](#)


## POSIX IPC

- ▶ POSIX IPC little more regular, uses filesystem to identify IPC objects
- ▶ Originated as optional POSIX/SUS extension, now required for compliant Unix
- ▶ Covered in our textbooks partially. Example:
  - ▶ Stevens/Rago 15.10 POSIX Semaphores
  - ▶ Robbins/Robbins 14.3-5 POSIX Semaphores
- ▶ [Additional differences on StackOverflow](#)

**We will favor POSIX**

## Model Problem: Dining “Philosophers”



- ▶ Each Swansons will only eat with two forks
- ▶ JJ's only has 5 forks, must share
- ▶ After acquiring 2 forks, a Swanson eats an egg, then puts both forks back to consider how awesome he is
- ▶ Algorithms that don't share forks will lead to injury



Source: Aditya Y. Bhargava,  
Originally: Dustin D'Arnault

## Exercise: Protocol for Dining “Philosophers”

- ▶ Each Swansons will only eat with two forks
- ▶ JJ's only has 5 forks, must share
- ▶ Swanson's pick up one fork at a time
- ▶ After acquiring 2 forks, a Swanson eats an egg
- ▶ After eating an egg a Swanson puts both forks considers how awesome he is, repeats
- ▶ Swanson leaves after eating sufficient eggs
- ▶ Is there any potential for **deadlock**?  
How can this be avoided?
- ▶ Is there any chance for **starvation**?



# Answers: Protocol for Dining “Philosophers”

Deadlock

Deadlock: All try for left fork first



- ▶ Each Swanson acquires left fork: cycle
- ▶ Each Waits forever for right fork

Dijkstra: One Swanson goes Right first

- ▶ Breaks the cycle so no deadlock possible
- ▶ Generalization establishes a **partial ordering** for each process to acquire resources, can prove lack of deadlocks

Starvation?

- ▶ A Swanson may wait indefinitely to get both forks, resource **starvation**
- ▶ Requires introduction of **priority** and communication to fix ([Chandy/Misra Solution](#))





# Semaphore



Source: [Wikipedia Railway Semaphore Signal](#)

- ▶ A counter variable with atomic operations
- ▶ **Atomic operation:** not divisible, all or none, no partial completion possible
- ▶ Used to coordinate access to shared resources such as shared memory, files, connections
- ▶ Typically allocate one semaphore per resource and acquire all that are needed

## Activity: Dining “Philosophers” with Semaphores

Examine the dining philosophers code here:

[http://www.cs.umn.edu/~kauffman/4061/philosophers\\_posix.c](http://www.cs.umn.edu/~kauffman/4061/philosophers_posix.c)

Use the Man Pages here:

[http://man7.org/linux/man-pages/man7/sem\\_overview.7.html](http://man7.org/linux/man-pages/man7/sem_overview.7.html)

Find out how the following are done:

1. What does a POSIX semaphore look like?
2. How does one create a POSIX semaphore?
3. What calls are used to “acquire” and “release” a POSIX semaphore?
4. What happens when multiple processes modify the same semaphore?
5. How are semaphores used to coordinate use of forks?
6. How is deadlock avoided in the code?

# Lessons Learned from philosophers\_posix.c (1/2)

## Posix Semaphores

- ▶ POSIX semaphores are single integer values with atomic operations
- ▶ Semaphore operations are guaranteed to be **atomic**: only one process can increment/decrement at a time, function as efficient **locks**
- ▶ `sem_t *sem = sem_open(name, ...)`; is used to obtain a semaphore from the operating system. Uses *named semaphores* which are managed by OS, shared between processes
- ▶ `sem_wait(sem)`; wait until semaphore is non-zero, then atomically decrement/lock it
- ▶ `sem_post(sem)`; increment/unlock a semaphore and schedule a process waiting on it to run
- ▶ `sem_close(sem)`; to stop using a semaphore

## Lessons Learned from philosophers\_posix.c (2/2)

### Dining Philosophers

- ▶ In Dining Philosophers, one semaphore per utensil
- ▶ Philosophers acquire both adjacent utensils to eat
- ▶ Circular Deadlock avoided via one Philosopher acquiring in a different order
  - ▶ Philosopher N: Get right utensil, then left utensil
  - ▶ Other Philosopher: Get left utensil, the right utensil

## Alternative: System V Semaphores

- ▶ File `philosophers_sysv.c` implements same problem with System V semaphores which look stranger than POSIX
  - ▶ Always come in an array of multiple semaphores
  - ▶ Operate atomically on the array: can incr/decr multiple semaphores at once
  - ▶ Requires use of structs to perform operations
  - ▶ Provide some other forms of synchronization such as waiting until a semaphore reaches 0
  - ▶ C calls such as `semget()`, `semctl()`, `semop()`
- ▶ Net effect is the same: each Philosopher locks a utensil by atomically decrementing and incrementing semaphores

# The Nature of a Semaphore

- ▶ As seen, semaphores have several component parts that the OS manages
  1. An integer, usually quantity of resources available, often 1 or 0
  2. A locking mechanism allowing atomic operations on the value
  3. A **wait queue** of processes (or threads) that are blocked until the semaphore becomes non-zero
- ▶ Semaphores typically used to **lock** a resource or protect a critical region in code, efficiently blocks competing processes
- ▶ Later will discuss each component separately in context of threads
  - ▶ Locks are **Mutexes**
  - ▶ Wait queues are **Condition variables**
  - ▶ Integers can be any variable in memory
  - ▶ Can build our own semaphores as with these
  - ▶ [SO: cucufrog on Condition Variables vs Semaphores](#)

## Linux shows Posix IPC objects under /dev/shm

```
> gcc -o philosophers philosophers_posix.c -lpthread
> ./philosophers
Swanson 0: wants utensils 0 and 1
Swanson 2: wants utensils 2 and 3
Swanson 1: wants utensils 1 and 2
...
Swanson 3 (egg 10/10): leaving the diner
pausing prior to cleanup/exit (press enter to continue)
while you're waiting, have a look in /dev/shm
  C-z
[1]+  Stopped                  ./philosophers

> ls -l /dev/shm
total 20K
-rw----- 1 kauffman kauffman 32 Apr  1 21:36 sem.utensil_0
-rw----- 1 kauffman kauffman 32 Apr  1 21:36 sem.utensil_1
-rw----- 1 kauffman kauffman 32 Apr  1 21:36 sem.utensil_2
-rw----- 1 kauffman kauffman 32 Apr  1 21:36 sem.utensil_3
-rw----- 1 kauffman kauffman 32 Apr  1 21:36 sem.utensil_4

> fg
./philosophers

> ls -l /dev/shm
total 0
```

/dev/shm is a Linux convention, shard memory under as well,  
message queues under /dev/mqueue

## Exercise: Concurrent Appends to a File

- ▶ C code to the right appends a string to `the_file.txt`
- ▶ Shell code below launches 100 processes to append
- ▶ What's with the results of `wc`?

```
> gcc -o append_clobber append_clobber.c
> for i in $(seq 100); do
    append_clobber $i &
done
[3] 18724
[4] 18725
...
[99]- Done append_clobber $i
[100]+ Done append_clobber $i
> wc the_file.txt
 97  96 281 the_file.txt
```

```
1 // append_clobber.c
2 int main(int argc, char *argv[]){
3     if(argc < 2){
4         printf("usage: %s <word>\n");
5         return 1;
6     }
7
8     char *word = argv[1];
9     char *filename = "the_file.txt";
10
11     int fd = open(filename,
12                   O_CREAT | O_RDWR ,
13                   S_IRUSR | S_IWUSR);
14     lseek(fd, 0, SEEK_END);
15
16     int n = strlen(word);
17     word[n] = '\n';
18     write(fd, word, n+1);
19     close(fd);
20
21     return 0;
22 }
```



## Exercise 2: Concurrent Appends to a File

- ▶ `append_clobber.c` does not coordinate writes to the end of `the_file.txt` leading to some data to be overwritten and lost
- ▶ Fix this using **semaphores**
- ▶ Research alternatives that also allow safe appends to files

```
1 // append_clobber.c
2 int main(int argc, char *argv[]){
3     if(argc < 2){
4         printf("usage: %s <word>\n");
5         return 1;
6     }
7
8     char *word = argv[1];
9     char *filename = "the_file.txt";
10
11     int fd = open(filename,
12                   O_CREAT | O_RDWR ,
13                   S_IRUSR | S_IWUSR);
14     lseek(fd, 0, SEEK_END);
15
16     int n = strlen(word);
17     word[n] = '\n';
18     write(fd, word, n+1);
19     close(fd);
20
21     return 0;
22 }
```

# Answers: Concurrent Appends to a File

## Semaphore Solution

- ▶ `append_sem.c` uses a POSIX semaphore named `/the_lock` to lock `the_file.txt` prior to working on it
- ▶ Must lock prior to the `lseek()`, unlock after the `write()`

## Alternatives

- ▶ `append_lockf.c`: Use the `lockf()` function to get an exclusive lock on the file, identical to semaphore solution but tied to individual files
- ▶ `append_os.c`: Open with `O_APPEND`, OS guarantees all `write()`'s are appended to the end

```
1 // append_sem.c
2 int main(int argc, char *argv[]){
3     if(argc < 2){
4         printf("usage: %s <word>\n");
5         return 1;
6     }
7
8     char *word = argv[1];
9     char *sem_name = "/the_lock";
10    char *filename = "the_file.txt";
11
12    int fd = open(filename,
13                  O_CREAT | O_RDWR ,
14                  S_IRUSR | S_IWUSR);
15    sem_t *sem = sem_open(sem_name,
16                          O_CREAT,
17                          S_IRUSR | S_IWUSR);
18    sem_wait(sem);    // wait/lock semaphore
19    lseek(fd, 0, SEEK_END);
20
21    int n = strlen(word);
22    word[n] = '\n';
23    write(fd, word,n+1);
24    close(fd);
25
26    sem_post(sem);    // unlock semaphore
27    sem_close(sem);   // done with semaphore
28
29    return 0;
30 }
```


## Semaphore Application in Lab

This week's Lab will illustrate a more useful application of semaphores to coordinate multiple unrelated processes

- ▶ Process a "jobs" file with a list of shell commands to run
  - seq 100000
  - gcc --version
  - du . -h
  - ls
  - ls -l
  - date
  - ...
- ▶ Start multiple 'runners' which execute lines from the jobs file
  - > runner jobs.txt & runner jobs.txt &
  - # starts 2 runners to work on jobs.txt
- ▶ Runners read file lines, execute jobs, mark as done
  - D seq 100000
  - D gcc --version
  - R du . -h
  - D ls
  - R ls -l
  - date
  - ...
- ▶ To prevent duplication of job running, use a Semaphore to lock the file while making changes

# Shared Memory Segments

- ▶ The ultimate in flexibility is to get a segment of raw bytes that can be shared between processes
- ▶ POSIX shared memory outlives a process
- ▶ **Examine** `shmdemo_posix.c` to see how this looks
- ▶ Importantly, this program creates shared memory that outlives the program: must clean it up at some point



Source: SoftPrayog System V IPC

## Exercise: Shared Memory Coordination

- ▶ Creating shared memory is relatively easy
- ▶ Like files, Coordinating shared memory is not automatic
- ▶ Consider `shared_flip.c`
  - ▶ Shared memory of all “00000” or “1111”
  - ▶ `shared_flip -flip` flips all characters ( $0 \rightarrow 1$ ,  $1 \rightarrow 0$ )
- ▶ What happens if multiple programs simultaneously try to flip bits?
- ▶ How does one prevent “corruption” of that data?
- ▶ Experiment noting that a command like  
`for i in $(seq 100); do ./shared_flip -flip & done`  
will start 100 identical processes as background jobs

# Answers: Shared Memory Coordination

```
1 // No Coordination: Errors
2
3
4
5 printf("flipping bits\n");
6 for(int i=0; i<SHM_SIZE-1; i++){
7     if(shared_bytes[i] == '0'){
8         shared_bytes[i] = '1';
9     }
10    else if(shared_bytes[i] == '1'){
11        shared_bytes[i] = '0';
12    }
13 }
14
15
```

```
1 // Coordinate via Semaphore
2 sem_t *sem =
3     sem_open(sem_name, O_CREAT, S_IRUSR|S_IWUSR);
4 sem_wait(sem); // lock semaphore
5 printf("flipping bits\n");
6 for(int i=0; i<SHM_SIZE-1; i++){
7     if(shared_bytes[i] == '0'){
8         shared_bytes[i] = '1';
9     }
10    else if(shared_bytes[i] == '1'){
11        shared_bytes[i] = '0';
12    }
13 }
14 sem_post(sem); // unlock sem
15 sem_close(sem);
```

# Shared Memory vs mmap'd Files

- ▶ Recall Memory Mapped files give direct access of OS buffer for disk files
- ▶ Changes to file are done in RAM and occasionally `sync()`'d to disk (permanent storage)
- ▶ POSIX Shared Memory segment cut out the disk entirely: an OS buffer that looks like a file but has no permanent backing storage
- ▶ Which to pick?
  - ▶ Shared Memory when data does not need to be saved permanently and/or syncing would costly
  - ▶ Memory Mapped File when data should be saved permanently
- ▶ Related concept: [RAM Disk](#), a main memory file system, high performance with no permanence

## Exercise: Email lookup with Shared Memory

- ▶ In HW, worked on a simple email lookup “server” or database
- ▶ Clients connected to server, server gave back emails based on name
- ▶ Shared memory makes server/client less relevant
- ▶ Propose how to use shared memory for email lookups AND alterations
- ▶ How might multiple processes coordinate use of shared memory?

```
// structure to store a lookup_t of  
// name-to-email association
```

```
typedef struct {  
    char name [STRSIZE];  
    char email[STRSIZE];  
} lookup_t;
```

```
lookup_t original_data[NRECS] = {  
    {"Chris Kauffman"      ,"kauffman@umn.edu"},  
    {"Christopher Jonathan","jonat003@umn.edu"},  
    {"Amy Larson"          ,"larson@cs.umn.edu"},  
    {"Chris Dovolis"       ,"dovolis@cs.umn.edu"},  
    {"Dan Knights"         ,"knights@cs.umn.edu"},  
    {"George Karypis"      ,"karypis@cs.umn.edu"},  
    ...
```

```
# Sample of potential use
```

```
> email_db lookup 'Chris Kauffman'
```

```
Looking up Chris Kauffman
```

```
Found: kauffman@umn.edu
```

```
> email_db lookup 'Rick Sanchez'
```

```
Looking up Rick Sanchez
```

```
Not found
```

```
> email_db change 'Chris Kauffman' 'kman@kauffmoney.com'
```

```
Changing Chris Kauffman to kman@kauffmoney.com
```

```
Alteration complete
```

```
> email_db lookup 'Chris Kauffman'
```

```
Looking up Chris Kauffman
```

```
Found: kman@kauffmoney.com
```



## Answer: Email lookup with Shared Memory

- ▶ Store entire array of name/email in a piece of shared memory with a known name/file
- ▶ Likely want database of saved so a **memory mapped file** is probably best
- ▶ Processes open shared memory/file, scan through looking
- ▶ Updates can be done by altering the shared memory
- ▶ **Danger** multiple processes writing may corrupt the data
- ▶ Use semaphores to control access for reading/writing, would need to establish a **protocol** for this
  - ▶ Lock entire database: easy but only one lookup at a time
  - ▶ Lock individual name/emails: more complex but potential for more throughput
  - ▶ Differentiate reading locks and writing locks

# Posix Message Queues

- ▶ Implements basic send/receive functionality through shared memory
- ▶ Message Queues share much with FIFOs
  - ▶ `mq_send()` is similar to `write()` to a FIFO
  - ▶ `mq_receive()` is similar to `read()` from a FIFO
  - ▶ Known global name of a message queue ~ name of FIFO file
- ▶ Differences from FIFOs
  - ▶ FIFOs/Pipes have a fixed total size (64K)
  - ▶ FIFOs allow `read()/write()` of arbitrary # of bytes
  - ▶ Message Queues limit #messages and max size of messages on queue
  - ▶ Message Queues send/receive individual messages

# Kirk and Spock: Talking Across Interprocess Space

- ▶ Demo the following pair of simple communication codes which use Posix Message Queues.
- ▶ Examine source code to figure out how they work.



See `msg_kirk_posix.c` and `msg_spock_posix.c`

# Email Lookup with Message Queues

- ▶ Recent HWs build an email lookup server using FIFOs
- ▶ Another HW compare it to an approach that uses Message Queues
- ▶ Worth of study to see the many similarities between FIFOs/Message Queues and a few of the differences between them
- ▶ Such contrast between IPC mechanisms make for good Exam questions

# More Resources IPC

## System V IPC

- ▶ <http://beej.us/guide/bgipc/>
- ▶ <http://www.tldp.org/LDP/tlk/ipc/ipc.html>

## General Overview

- ▶ [http://man7.org/conf/lca2013/IPC\\_Overview-LCA-2013-printable.pdf](http://man7.org/conf/lca2013/IPC_Overview-LCA-2013-printable.pdf)