

Introduction to Parallel Computing

Chris Kauffman

*Last Updated:
Tue Jan 18 09:38:11 AM CST 2022*

Logistics

Reading

Grama Ch 1

Goals

- ▶ Motivate: Parallel Programming
- ▶ Overview concepts a bit
- ▶ Discuss course mechanics

Assignment 1

- ▶ Post Tomorrow
- ▶ Due at end of next week
- ▶ Basic theory / terminology

Registered or Not?

If you are **not registered** for the course but want to be...


- ▶ Come to the first week of Lecture so you don't fall behind
- ▶ Write on a piece of paper the following information
 1. Name, UMN Email address, Student ID Number
 2. Which Lecture and Lab section you want to register for
 3. 2-3 sentences about why you absolutely must take 2021 this semester, consequences if you do not
- ▶ Give me that sheet of paper
- ▶ Wait and hope: very limited space + waitlists for full labs

How Small are Transistors?

- ▶ Recent Intel CPUs uses a 14 nanometer manufacturing process
- ▶ Distance between memory units in the processor is about 28 nanometers
- ▶ A hard sphere radius of a hydrogen Atom is about 0.11 nanometers
- ▶ About 255 **atoms** apart - weird things start happening at that scale
- ▶ The term “X-nanometer process” is not a true standard but the point is manufacturers continue increase Transistor Density

Moore's Law Alive and Well

48 Years of Microprocessor Trend Data




Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2019 by K. Rupp

Source: [karlrupp / microprocessor-trend-data on Github](#)

- ▶ Still get more transistors (yay!)
- ▶ Clock Frequency & Single CPU Speed is stalling (boo!)

Similar Plot on Processor Speed



Source: Danowitz et al

- ▶ CPU speed isn't getting faster these days
- ▶ Fastest Dell Speed I found was 4.8 GhZ on "Turbo" (my laptop caps at 3.4 GhZ)


Today's Processors

Mini-Parallel Computer

- ▶ 1 CPU comes with several “Cores” - independent identical processing units
- ▶ Cores share some resources like Cache and Comm. hardware
- ▶ Cores can run multiple programs simultaneously
- ▶ Cores can **cooperate** within on a single program to (possibly) make that program faster



Performance is Moving to Parallel



Source: There's plenty of room at the Top: What will drive computer performance after Moore's law? by Leiserson et. al, Science 2020

Multitasking

multitask (verb)

To use the restroom and brush your teeth at the same time.


I was late for work today so I had to multitask!

- ▶ [Urban Dictionary Definition 2 by sasm](#)

Questions

- ▶ Do humans multitask and if so how?
- ▶ What kinds of things can humans multitask?
- ▶ Are you a good multitasker?
- ▶ How do humans get a big job done faster?

Focus



The students in the first experiment who were asked to multitask averaged 11 per cent lower on their quiz.

The students in the second experiment who were surrounded by laptops scored 17 per cent lower.

Laptop use lowers student grades, experiment shows, The Canadian Press, 8-14-2013
Original Paper by Sana et al

Computers and Multitasking

- ▶ How do computers multitask?
- ▶ Can a single CPU computer multitask?
- ▶ What are the advantages/disadvantages of this?
- ▶ What might drive one to write a parallel program?

Parallelism Shows Up a Lot in Computing

Low-level/hardware parallelism

- ▶ CPU Instructions,
- ▶ VLIW: Very Long Instruction Word
- ▶ Multi-media CPU instructions
- ▶ Graphics and GPU instructions
- ▶ CPU Pipelines
- ▶ Memory subsystem
- ▶ I/O controller

Some of these can be controlled via assembly language choices, others are implicit in the hardware and only indirectly controllable by a programmer.

- ▶ We will touch at times on **Vector Instructions** that to perform the same calculation on several pieces of data
- ▶ We will focus on **Programmatic Parallelism**: writing code that is parallel at a higher level

Primary Motivation for Parallel Computing

Use multiple processors to...

- ▶ **Speedup**: solve the same size problem in less time
- ▶ **Sizeup**: solve a larger problem in the same time

EXAMPLE: A *serial* program takes 100 seconds to perform a calculation on a 10MB input image file.

Speedup

- ▶ 2 processors should finish the 10MB calculation in 50 seconds
- ▶ Right? right?
- ▶ Just like two people can bake a cake twice as fast as a single person...

Sizeup

- ▶ 2 processors should finish a 20MB file calculation in 100 seconds
- ▶ Right? right?
- ▶ Just like 2 ovens can handle a double-sized cake...

Exercise: Scan for Max Serial vs Parallel

- ▶ Have an array of large size N with numbers in it
- ▶ Want to find the maximum number

```
# Find max serially
# ON CPU 0
my_max = arr[0]
for i=1 to N-1:
    if my_max < arr[i]:
        my_max = arr[i]
done
return my_max
```

If you were allowed both CPU 0 **and** CPU 1, how could you make this go faster?

State any assumptions you make about how CPU 0 & 1 interact / access data in `arr[]`.

Answers: Scan for Max

Assumption: both CPUs can access all of arr[]

```
# ON CPU 0
```

```
my_max = arr[0]
```

```
for i=1 to N/2-1:
```

```
    if my_max < arr[i]:
```

```
        my_max = arr[i]
```

```
done
```

```
other_max = receive number
```

```
            from CPU 1
```

```
if other_max > my_max:
```

```
    my_max = other_max
```

```
return my_max
```

```
# ON CPU 1
```

```
my_max = arr[N/2]
```

```
for i=N/2+1 to N-1:
```

```
    if my_max < arr[i]:
```

```
        my_max = arr[i]
```

```
done
```

```
send my_max to CPU 0
```

Final comparison steps cannot be parallelized - creates a limit (albeit mild) on how much can be done in parallel.

Amdahl's Law and the Cost of Coordination

Likely that Parallel Max version completes in around 50% of time of the Serial Max, but never quite a perfect 2X speedup. This has a formalization.

Amdahl's Law

Speedup is limited by the portion of the program that can be parallelized and the degree to which that portion can be parallelized

- ▶ Amdahl's law states that can get close to linear speedup but it's never perfect
- ▶ All algorithms will contain some serial portion that cannot be parallelized
- ▶ All parallel algorithms feature some communication overhead that is not present in serial versions

Will study this and algorithm in more detail later

Concurrency and Parallelism

StackOverflow Questions

Concurrency vs Parallelism - What is the difference?

Accepted Answer (RichieHindle)

Concurrency is when two or more tasks can start, run, and complete in overlapping time periods. It doesn't necessarily mean they'll ever both be running at the same instant. Eg. multitasking on a single-core machine.

Parallelism is when tasks literally run at the same time, eg. on a multicore processor.

Implications

Concurrent $\not\Rightarrow$ *Parallel*: A concurrent program does not need to be run in parallel (but it often is).

Parallel \Rightarrow *Concurrent*: A parallel program must have concurrent parts and had better deal with concurrency issues.

Speed to Completion Matters

By 2015 the GFS model had fallen behind the accuracy of other global weather models. This was most notable in the GFS model incorrectly predicting Hurricane Sandy turning out to sea until four days before landfall, while the European Centre for Medium-Range Weather Forecasts' model predicted landfall correctly at 7 days. Much of this was suggested to be due to limits in computational resources within the National Weather Service. In response, the NWS purchased new supercomputers, increasing processing power from 776 teraflops to 5.78 petaflops. In 2018, the processing power was increased again to 8.4 petaflops.

– [*Wikip: Global Forecast System*](#)

Getting correct predictions in time for action is important in many disciplines - parallel computing provides the mechanism for this.

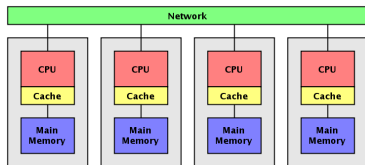
What We Will Discuss

- ▶ Low latency parallel computers
- ▶ General hardware/network architectures of parallel computers
- ▶ Distributed memory parallel computers
 - ▶ Message Passing Interface (MPI)
- ▶ Shared Memory parallel computers
 - ▶ Interprocess Communication
 - ▶ OpenMP
 - ▶ POSIX Threads
- ▶ Co-processor Parallelism
 - ▶ GPGPU: General Purpose Graphics Processing Units
 - ▶ CUDA for NVIDIA GPU programming
- ▶ Analyzing Parallel Algorithms
- ▶ Parallelizing Classic Algorithms
 - ▶ Matrix Ops
 - ▶ Sorting
- ▶ *Scientific computing angle*

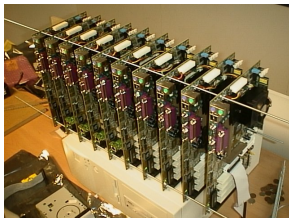
General Arrangement of Topics

Weeks 1-7

- ▶ Basics + Theory
- ▶ Distributed Memory



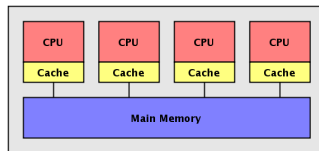
Source: Kaminsky/Parallel Java



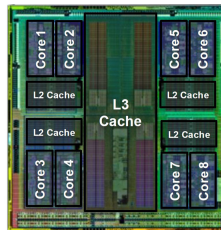
Source: ClusterComputer.com

Weeks 8-12

- ▶ Shared Memory
- ▶ GPGPU / CUDA Programming



Source: Kaminsky/Parallel Java



What We Won't Discuss

- ▶ General Networking (CSCI 4211/5211)
- ▶ Server/Client Model for Web Programming (CSCI 4131)
- ▶ Deep Dive into Parallel Hardware (CSCI 5204)
- ▶ Graphics applications for GPUs (CSCI 4611 / 5607+8)
- ▶ Distributed/Grid Computing (CSCI 5105 / 5751)

Course Mechanics

Mull over the syllabus linked to our Canvas Page

Look for

- ▶ Contact info
- ▶ GTA
- ▶ Textbook
- ▶ Weights / Grading Scheme
- ▶ Hot Seats / Bonus Cards

Schedule of topics and lecture materials:

[https:](https://www-users.cse.umn.edu/~kauffman/5451/schedule.html)

[//www-users.cse.umn.edu/~kauffman/5451/schedule.html](https://www-users.cse.umn.edu/~kauffman/5451/schedule.html)

Assignments and Hardware

- ▶ Assignments are 50% of your grade
- ▶ Combination of writing and programming
- ▶ Programming mostly in **C in a Linux environment**
 - ▶ If you have never programmed in C before, begin studying it **intensively** or drop the course
 - ▶ If syntax like `malloc()` or `int *p = &s.field;` is unfamiliar, begin reviewing **intensively**
 - ▶ If unfamiliar with `make` and command line, **get familiar quick**
 - ▶ Consult me during office hours if in doubt
- ▶ Getting a parallel cluster set up for you to use on CSE Labs
- ▶ Assignments will allow work in Pairs of students
- ▶ 5 Assignments roughly on
 1. Basics + Theory
 2. MPI Coding + Analysis
 3. PThreads Coding + (Maybe IPC as well)
 4. OpenMP Shared Memory Coding
 5. CUDA / GPGPU coding

Exams Small and Large

Mini-Exams: 3 scheduled, 30% of grade

- ▶ Between a quiz and a midterm exam in length
- ▶ 30 minutes during lecture
- ▶ 1 page, front and back
- ▶ Topics will cover stuff from lecture, assignments, readings
- ▶ Will work practice problems in class preceding them

Final Exam: 2 hours, 20% of grade

Open Resource Exams

Unless otherwise specified, exams are open resource: use notes, compiler, code, slides, textbook. No googling, browsing, communication, cheating


A Word on Safety

Please wear your mask during lecture



- ▶ For your safety, my safety, the safety of the class, and the safety of all the old, young, and immuno-compromised loved ones that we see but do not want to hurt, Mask Up.
- ▶ Refrain from eating/drinking during lecture
- ▶ Keep your mask on the whole time
- ▶ If you feel sick, stay home, watch the videos, notify me if the illness is prolonged and we will make arrangements

Reading For Next Time



Reading: Grama Ch 2

- ▶ **Focus on 2.3-5**, material pertaining to distributed memory
- ▶ We will return to shared memory arch later in the course
- ▶ Cache Coherence, PRAM models, False Sharing, Memory Bus are all shared memory topics
- ▶ Sections 2.1 and 2.2 optional, deeper architectures
- ▶ Sections 2.6 and 2.7 encouraged, deeper on networks