

MPI Basics

Chris Kauffman

*Last Updated:
Thu Feb 3 09:45:08 AM CST 2022*

Logistics

Reading: Grama Ch 6 + 4

- ▶ Ch 6: MPI basics
- ▶ Ch 4: Communication patterns

Assignments

- ▶ A1 Due Soon
 - ▶ On-time submission Wed 2/2
 - ▶ Late Submission Fri 2/4
- ▶ **Questions?**
- ▶ A2 up later this week: MPI Programming

Today

- ▶ Primitives for Distributed Memory Computing
- ▶ MPI Programming

Next Week

- ▶ Comm Patterns
- ▶ Thu 2/10: Mini-Exam 1

Generic Send and Receive

Minimum required functionality to do distributed memory parallel computing:

```
send(void *sendbuf, int nelems, int dest)
```

```
receive(void *recvbuf, int nelems, int source)
```

Sample Use

	P0	P1
1		
2		
3	a = 100;	receive(&a, 1, 0)
4	send(&a, 1, 1);	printf("%d\n", a);
5	a=0;	

- ▶ Proc 0 sends a single integer to Proc 1
- ▶ Proc 0 then 0s that integer
- ▶ Proc 1 receives and prints the integer

More typical appearance

Will typically write this as a single program which every processor runs.

```
void exchange(){  
    int a = 100;  
    int my_proc = get_processor_number();  
    if(my_proc == 0){  
        send(&a, 1, 1);  
        a=0;  
    }  
    else if(my_proc == 1){  
        receive(&a, 1, 0);  
        printf("%d\n", a);  
    }  
}
```

- ▶ Function to identify proc number
- ▶ Branching on proc number to take different actions

Flavors Send/Receive

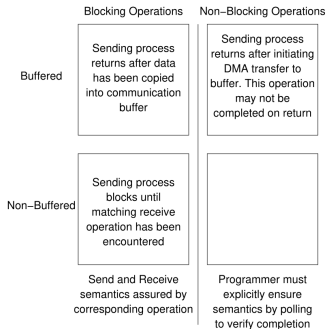


Figure 6.3 Space of possible protocols for send and receive operations.

- ▶ Hardware/OS support for buffered communication tends to make things run faster
- ▶ Usually have function calls available to do `send()` (blocking) and `send_nonblocking()` but must have some hardware support for it

Blocked and Unbuffered

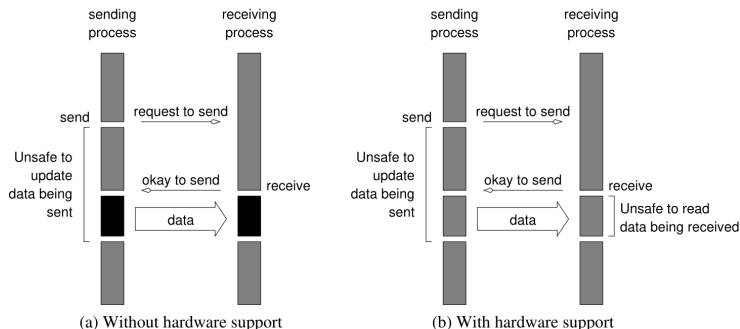


Figure 6.4 Non-blocking non-buffered send and receive operations (a) in absence of communication hardware; (b) in presence of communication hardware.

Blocking/Unbuffered: no extra buffer available to hold pending sends/receives so must wait, wait until message is sent to proceed
Blocked processors are idle, do no work, which cuts into speedup

Ordering of Send Receive

1	P0	P1
2		
3	<code>send(&a, 1, 1);</code>	<code>send(&a, 1, 0);</code>
4	<code>receive(&b, 1, 1);</code>	<code>receive(&b, 1, 0);</code>

Assuming send/receive blocked/unbuffered, what's wrong with the above code?

Blocking with Buffers

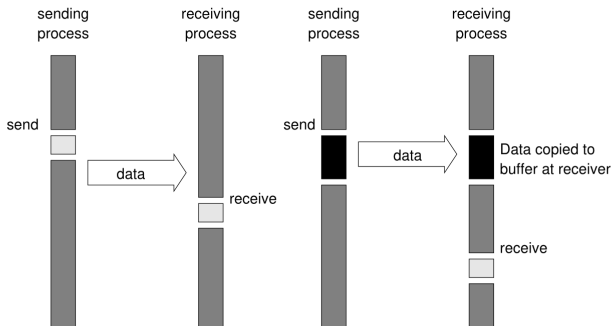


Figure 6.2 Blocking buffered transfer protocols: (a) in the presence of communication hardware with buffers at send and receive ends; and (b) in the absence of communication hardware, sender interrupts receiver and deposits data in buffer at receiver end.

Hardware buffer support, sender and receiver have a memory minion

No buffer support: sender interrupts receiver

The Danger Continues

1	P0	P1
2		
3	<code>receive(&a, 1, 1);</code>	<code>receive(&a, 1, 0);</code>
4	<code>send(&b, 1, 1);</code>	<code>send(&b, 1, 0);</code>

- ▶ `receive()` always blocks until message is obtained
- ▶ Does the above code work even in the buffered setting?

Non-blocking Communication

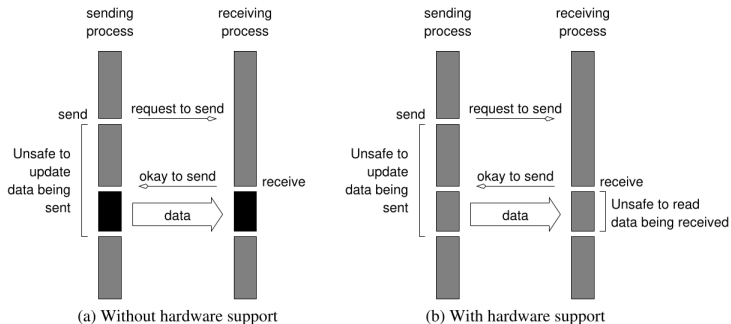
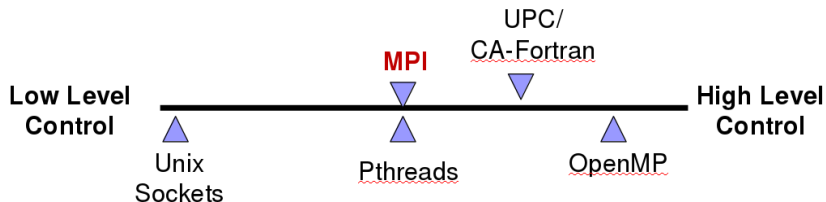


Figure 6.4 Non-blocking non-buffered send and receive operations (a) in absence of communication hardware; (b) in presence of communication hardware.

- ▶ Takes a bit more work on the programming side
- ▶ Must explicitly ensure that transaction completes with function calls
- ▶ `isend(data, dest, status)`: send w/o waiting
- ▶ `ireceive(data, dest, status)`: receive w/o waiting
- ▶ `wait(status)`: wait until a message has been sent or

MPI: Message Passing Interface

- ▶ Standardized library of functions for C/C++/Fortran
- ▶ Communicate between processors in a distributed memory machine
- ▶ Open source implementations: MPICH, Open MPI
- ▶ Proprietary: Intel, Platform, IBM, Platform, Cray
- ▶ Typically geared for particular architecture
- ▶ May exploit specifics of a particular machine



MPI In a Nutshell: 6 Essential Functions

```
// Initializes MPI.
```

```
int MPI_Init(int *argc, char ***argv) ;
```

```
// Terminates MPI.
```

```
int MPI_Finalize() ;
```

```
// Determines the number of processes.
```

```
int MPI_Comm_size(MPI_Comm comm, int *size);
```

```
// Determines the label of the calling process.
```

```
int MPI_Comm_rank(MPI_Comm comm, int *rank);
```

```
// Sends a message.
```

```
int MPI_Send(void *buf, int count, MPI_Datatype datatype,  
             int dest, int tag, MPI_Comm comm);
```

```
// Receives a message.
```

```
int MPI_Recv(void *buf, int count, MPI_Datatype datatype,  
            int source, int tag, MPI_Comm comm,  
            MPI_Status *status);
```

MPI Setup: Hello World

```
#include <stdio.h>
#include <mpi.h>

int main (int argc, char *argv[]){
    int rank, size;
    MPI_Init (&argc, &argv);           // starts MPI
    MPI_Comm_rank (MPI_COMM_WORLD, &rank); // get current process id
    MPI_Comm_size (MPI_COMM_WORLD, &size); // get number of processes
    for(int i=0; i<1; i++){
        printf("Hello world from process %d of %d\n",
               rank, size);
    }
    MPI_Finalize();
    return 0;
}
```

- ▶ Note the use of `MPI_COMM_WORLD` which is a predefined constant corresponding to all processors.
- ▶ Can also set up other communicators that correspond to subsets of processors

Compilation and Running

- ▶ Demo using openmpi implementation
- ▶ mpirun for interactive running
- ▶ mpirun -np 4
progr sets number of
“processors” to 4

```
lila [test-code]% mpicc -o hello hello.c
```

```
lila [test-code]% ./hello  
Hello world from process 0 of 1
```

```
lila [test-code]% mpirun hello  
Hello world from process 0 of 4  
Hello world from process 1 of 4  
Hello world from process 2 of 4  
Hello world from process 3 of 4
```

```
lila [test-code]% mpirun -np 2 hello  
Hello world from process 0 of 2  
Hello world from process 1 of 2
```

```
lila [test-code]% mpirun -np 8 hello  
Hello world from process 7 of 8  
Hello world from process 0 of 8  
Hello world from process 2 of 8  
Hello world from process 3 of 8  
Hello world from process 4 of 8  
Hello world from process 6 of 8  
Hello world from process 1 of 8  
Hello world from process 5 of 8
```

MPI Implementations and OpenMPI Warnings

- ▶ Several Implementations of MPI:
 - ▶ [OpenMPI](#) and [MPICH](#) are free, open source, widely available
 - ▶ HPC Vendors like IBM and Cray provide their own tailored MPI versions
- ▶ Recent Versions of OpenMPI can complain a LOT about various items missing
- ▶ The CSE Labs machines with MPI are not all configured perfectly leading to additional errors
 - ▶ Example: `--mca btl_base_warn_component_unused 0` to warn about missing HPC network components during `mpirun`
 - ▶ Example: `--mca opal_warn_on_missing_libcuda 0` if not intending to use GPU libraries
- ▶ Exact nature of warnings/errors varies a lot, look at messages which often dictate how to disable them

Warning Suppression in OpenMPI

```
csel-plate02 [examples]% mpirun -np 4 ./mpi_hello
```

```
[8230,1]: A high-performance Open MPI point-to-point messaging module
was unable to find any relevant network interfaces:
```

```
Module: OpenFabrics (openib)
Host: csel-plate02
```

Another transport will be used instead, although this may result in lower performance.

NOTE: You can disable this warning by setting the MCA parameter `bt1 base warn component unused` to 0.

```
P 0: Hello world from process 0 of 4 (host: csel-plate02)
Hello from the root processor 0 of 4 (host: csel-plate02)
P 1: Hello world from process 1 of 4 (host: csel-plate02)
P 2: Hello world from process 2 of 4 (host: csel-plate02)
P 3: Hello world from process 3 of 4 (host: csel-plate02)
[csel-plate02:235926] 3 more processes have sent help message help-mpi-btl-base.txt / btl:
[csel-plate02:235926] Set MCA parameter "orte_base_help_aggregate" to 0 to see all help / e
csel-plate02 [examples]% mpirun --btl_base_warn_component_unused 0 -np 4 ./mpi_hello
mpirun: Error: unknown option "--btl_base_warn_component_unused"
Type 'mpirun --help' for usage.
```

```
#
cshel-plate02 [examples]% mpirun --mca btl_base_warn_component_unused 0 -np 4 ./mpi_hello
P 0: Hello world from process 0 of 4 (host: cshel-plate02)
Hello from the root processor 0 of 4 (host: cshel-plate02)
P 1: Hello world from process 1 of 4 (host: cshel-plate02)
P 2: Hello world from process 2 of 4 (host: cshel-plate02)
```


MPI Hostfile

Default OpenMPI config is to use all processors on a single machine then start failing

```
val [04-mpi-code]% mpirun -np 2 ./mpi_hello_plus
P 0: Hello world from process 0 of 2 (host: val)
Hello from the root processor 0 of 2 (host: val)
P 1: Hello world from process 1 of 2 (host: val)
```

```
val [04-mpi-code]% mpirun -np 8 ./mpi_hello_plus
```

There are not enough slots available in the system to satisfy the 8 slots that were requested by the application:

```
./mpi_hello_plus
```

On some systems, like our lab machines, will can use a `hostfile.txt` which gives additional machines to harness - true distributed computation

```
csel-plate02 [examples]% mpirun -np 400 --hostfile hostfile-plate-ip.txt ./mpi_hello_plus
P18: Hello world from process 18 of 400 (host: csel-plate02)
P21: Hello world from process 21 of 400 (host: csel-plate02)
P141: Hello world from process 141 of 400 (host: csel-plate03)
P310: Hello world from process 310 of 400 (host: csel-plate04)
P149: Hello world from process 149 of 400 (host: csel-plate03)
...
```

MPI Systems on CSE Labs

3 primary systems we will use for coding assignments

- ▶ PLATE

- ▶ `ssh csel-plate02.cselabs.umn.edu`
- ▶ 3 nodes with LOTS 72 cores per, 288 HW threads
- ▶ Likely use for MPI coding

- ▶ VEGGIE

- ▶ `ssh csel-carrot.cselabs.umn.edu`
- ▶ 5 nodes with multiple Cores + GPU

- ▶ PHI

- ▶ `ssh phi01.cselabs.umn.edu`
- ▶ Older hardware soon to retire
- ▶ Several cores per node, 8-10 nodes

MPI Send and Recieve

```
int a[10], b[10];
int partner = 1;
...

// Send contents of a to partner proc with tag=1
MPI_Send(a, 10, MPI_INT, partner, 1, MPI_COMM_WORLD);

// Receive message into b from partner proc with tag=1,
// ignore status of receipt
MPI_Recv(b, 10, MPI_INT, partner, 1, MPI_COMM_WORLD, MPI_STATUS_IGNORE)
```

- ▶ Analyze the program `send_receive_test.c`
- ▶ **Compare** with `send_bugs.c` which demos stall problems
- ▶ Note `MPI_ANY_SOURCE` may be used for `recv`'s source

Tags Make Messages Unique

```
int a[10], b[10], myrank;
MPI_Status status;
...
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);
if (myrank == 0) {
    MPI_Send(a, 10, MPI_INT, 1, 1, MPI_COMM_WORLD);
    MPI_Send(b, 10, MPI_INT, 1, 2, MPI_COMM_WORLD);
}
else if (myrank == 1) {
    MPI_Recv(b, 10, MPI_INT, 0, 2, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
    MPI_Recv(a, 10, MPI_INT, 0, 1, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
}
```

- ▶ Tags must be honored on receive
- ▶ Above code may deadlock if not buffered due to the misordering of tags
- ▶ Mostly we will use tag=1 for simplicity, alt MPI_ANY_TAG

MPI Data Types Supported

```
// Sends a message.
```

```
int MPI_Send(void *buf, int count, MPI_Datatype datatype,  
             int dest, int tag, MPI_Comm comm);
```

```
// Receives a message.
```

```
int MPI_Recv(void *buf, int count, MPI_Datatype datatype,  
             int source, int tag, MPI_Comm comm,  
             MPI_Status *status);
```

- ▶ Type of buffer is always untyped (void* buf)
- ▶ To try to get at slightly better safety, MPI has standard datatypes

MPI_CHAR	signed char
MPI_INT	signed int
MPI_LONG	signed long int
MPI_FLOAT	float
MPI_DOUBLE	double
MPI_BYTE	Last two used for sending
MPI_PACKED	structure arrays

Unsigned types also available

Exercise: Heat Transfer in MPI

- ▶ Discuss conversion of the following A1 code to an MPI version
- ▶ How is data in H divided up?
- ▶ Is communication required?
- ▶ How would one arrange MPI_Send / MPI_Recv calls?
- ▶ How much data needs to be transferred and between who?
- ▶ When the computation is finished, how can all data be displayed?

```
// Simulate the temperature changes for internal cells
for(t=0; t<max_time-1; t++){
    for(p=1; p<width-1; p++){
        double left_diff  = H[t][p] - H[t][p-1];
        double right_diff = H[t][p] - H[t][p+1];
        double delta = -k*( left_diff + right_diff );
        H[t+1][p] = H[t][p] + delta;
    }
}
```

Some Patterns that occur in the problem

- ▶ Pair exchange of items: made easier with `MPI_sendrecv`
- ▶ Collecting final output for display: `MPI_Gather`
 - ▶ Previewed here
 - ▶ Discussed in following lectures

Exchange: Sendrecv for exchanging data between pairs

```
{
    double send[10], recv[10]; int partner;
    if(procid % 2 == 1){ // odd procs send left, receive left
        partner = procid-1;
        MPI_Send(send, 10, MPI_DOUBLE, partner, 1, MPI_COMM_WORLD);
        MPI_Recv(recv, 10, MPI_DOUBLE, partner, 1, MPI_COMM_WORLD,
                 MPI_STATUS_IGNORE);
    }
    else{ // even procs receive right, send right
        partner = procid+1;
        MPI_Recv(recv, 10, MPI_DOUBLE, partner, 1, MPI_COMM_WORLD,
                 MPI_STATUS_IGNORE);
        MPI_Send(send, 10, MPI_DOUBLE, partner, 1, MPI_COMM_WORLD);
    }
}

{ // Sendrecv simplifies this pattern
    double send[10], recv[10]; int partner;
    partner = (procid % 2 == 1) ? procid-1 : procid+1;
    MPI_Sendrecv(send, 10, MPI_DOUBLE, partner, 1,
                 recv, 10, MPI_DOUBLE, partner, 1,
                 MPI_COMM_WORLD, MPI_STATUS_IGNORE);
}
```



Take Care: Pair exchange can hang



```
{  
    double send[10], recv[10]; int partner;  
    partner = (procid % 2 == 1) ? procid-1 : procid+1;  
    MPI_Sendrecv(send, 10, MPI_DOUBLE, partner, 1,  
                 recv, 10, MPI_DOUBLE, partner, 1,  
                 MPI_COMM_WORLD, MPI_STATUS_IGNORE);  
}
```

- ▶ With 9 processors, logic is broken
- ▶ Proc 8 will wait to communicate with a partner that doesn't exist
- ▶ Program never terminates

Gather Preview



Source: Shun Yan Cheung Notes on MPI

- ▶ Every processor has computed columns
- ▶ One processor (usually `procid 0`) needs to gather all of the data
- ▶ Everyone calls `MPI_Gather()`

MPI_Gather Sample

Use of Gather

```
// Preamble for any code
MPI_Comm comm = MPI_COMM_WORLD;
int sendarray[100];
int procid, total_procs, *rbuf;
...;
// Only proc 0 needs space for all
// data
if(procid == 0) {
    rbuf = malloc(total_procs*100*
                  sizeof(int));
}

// Everyone calls gather
// proc 0 gets all data eventually
MPI_Gather(sendarray, 100, MPI_INT,
           rbuf, 100, MPI_INT,
           0, comm);
```

Equivalent Non-Gather Code

```
if(rank == 0){
    for(i=0; i<100; i++){
        rbuf[i] = sendarray[i];
    }
    for(i=1; i<total_procs; i++){
        int *rloc = &rbuf[i*100];
        MPI_Recv(rloc, 100,
                 MPI_INT, i,
                 tag, MPI_COMM_WORLD,
                 MPI_STATUS_IGNORE);
    }
}
else{
    MPI_Send(sendarray, 100,
             MPI_INT, 0,
             tag, MPI_COMM_WORLD);
}
```

Collective Communication Patterns Next

- ▶ gather is an example of a class of **Collective Communication Patterns**
- ▶ Will study more of these in subsequent lectures
- ▶ Using built-in collective comm. patterns simplifies programs and allows MPI implementation to exploit network as much as possible

Sending Structs

Sending structs can be done via the MPI_BYTE type

```
{  
    // from send_structs.c  
    typedef struct {  
        double x;  
        int a, b;  
    } dint_t;  
    ...;  
    dint_t mine[10];  
    // calculate data sizes "manually" just as is done in a malloc()  
    MPI_Send(mine, 10*sizeof(dint_t), MPI_BYTE,  
            partner, 1, MPI_COMM_WORLD);  
}
```

- ▶ Simple and effective if all compute nodes **use the same binary layout**
- ▶ MPI also provides a (complex) method for situations where struct layout differs between nodes
- ▶ Must Dictate # of struct fields, types, and ordering into a MPI_Datatype and use MPI_Type_create_struct()
- ▶ Likely hurts performance if struct layout differs so will not discuss in detail