

HPC Linear Algebra

Chris Kauffman

*Last Updated:
Thu Feb 10 09:37:39 AM CST 2022*

Logistics

Assignments

- ▶ A1 grading in progress, look for results in the next day
- ▶ A2 will go up Friday and feature MPI Coding

MSI Accounts

- ▶ MSI has set up class accounts for students in 5451
- ▶ Should allow you to ssh `X500@mangi.msi.umn.edu`
- ▶ Will need to have [UMN VPN](#) enabled for this, [read UMN VPN Docs](#) to get set up
- ▶ More details next week

Today

- ▶ Complete Discussion of MPI Collective Communication
- ▶ Overview of some Linear Algebra Libraries
- ▶ Mini-Exam 1 - last 35min of lecture

Hand-Coded Matrix Algebra

- ▶ Very common for students to learn how to code up some basic linear algebra routines
- ▶ In reality, prototype and production code benefits from use of mature libraries for these
- ▶ Existing libraries for Linear Algebra already exist, are reliable and **fast**, both important in HPC / Parallel Computing

```
void matmult(  
    int arows, int bcols, int midim,  
    double A[][], // arows * midim  
    double B[][], // midim * bcols  
    double C[][]) // arows * bcols  
{  
    for (int i=0 ; i < arows ; ++i ){  
        for (int j=0 ; j < bcols ; ++j ){  
            C[i][j] = 0.0;  
            for (int k=0 ; k < midim ; ++k ){  
                C[i][j] += A[i][k] * B[k][j];  
            }  
        }  
    }  
}  
// try dgemm() instead
```

BLAS: Basic Linear Algebra Subroutines

- ▶ Started in the 1970's and now WIDELY deployed
- ▶ Defines a set of numerical operations in 3 Levels
 1. Vector/Scalar operations (add constant onto all vector elements) and Vector/Vector operations (dot product)
 2. Matrix/Vector operations (mat-vec multiply)
 3. Matrix/Matrix operations (mat-mat multiply)
- ▶ Interestingly these are all mostly $O(N)$, $O(N^2)$, $O(N^3)$ operations respectively
- ▶ The names for the function **suck** and take significant study to understand and use effectively

`axpy()`? `ddot()`? `sgemm()`? Are these rappers, hacker handles, or did someone just punch the keyboard repeatedly to name all the functions?

According to legend, all the function names are 5 letters or less as this was the limit imposed by the Fortran compiler which originally compiled them.

BLAS Introductory Example

dgemm() : Multiply 2 double precision, general format matrices

```
dgemm(opa, opb,           // transpose A,B or not
      arows, bcols, midim, // matrix dimensions
      alpha,              // scaling factor for product
      A, lda, B, ldb,     // A and B matrix + cols
      beta, C, ldc)       // answer scaling + storage + dim
```

$C := \alpha * opa(A) * opb(B) + \beta * C$

Super transparent, excellent software engineering...

- ▶ Targets Fortran77: different calling conventions than C
- ▶ Complex due to flexibility: 4 variants based on opa, opb

$$C \leftarrow \alpha AB + \beta C \quad C \leftarrow \alpha A^T B + \beta C$$

$$C \leftarrow \alpha AB^T + \beta C \quad C \leftarrow \alpha A^T B^T + \beta C$$

- ▶ Allows for scaling with alpha, beta but both often are 1
- ▶ Naming Convention: d ge mm ()
 - ▶ d: double precision real
 - ▶ ge: general matrix, not symmetric or banded
 - ▶ mm: matrix multiply

C BLAS Example

- ▶ cblas are C language bindings to BLAS routines
- ▶ Slightly easier to understand, uses symbolic names for some (extra) arguments
- ▶ Accounts for C being Row-Major vs Fortran being Column-Major

```
// dgemm_example.c
// A : arows * midim matrix
// B : midim * bcols matrix
// B : arows * bcols matrix
// C <- A*B
cblas_dgemm(CblasRowMajor, CblasNoTrans, CblasNoTrans,
            arows, bcols, midim,
            alpha,
            A, midim, B, bcols,
            beta, C, bcols);
```

LAPACK: Linear Algebra Package

- ▶ Basic Operations like Matrix Multiply are covered in BLAS
- ▶ Many Linear Algebra problems come up in HPC
 - ▶ Solve a Linear System: $Ax = b$, find x give A, b
 - ▶ Determine eigenvectors / eigenvalues for matrix A
 - ▶ Calculate Singular Value Decomposition on A
- ▶ LAPACK builds on BLAS to provide algorithms for all of these
- ▶ Has many of the same properties as BLAS
 - ▶ Netlib version Written in Fortran77
 - ▶ Has bindings for C in LAPACKE
 - ▶ Naming conventions are difficult: `dgesv()`
 - d: double precision real
 - ge: general format matrix
 - sv: “solve” linear system via a LUP decomposition
- ▶ LAPACK / BLAS often packaged together in single libraries

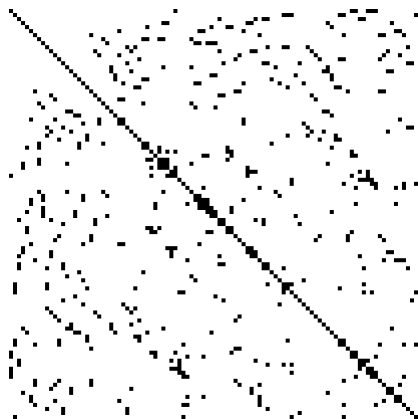
Implementations of BLAS+LAPACK

Package	Notes
Netlib	The official LAPACK. (BSD/Open Source)
ATLAS	Automatically Tuned Linear Algebra Software (BSD/Open Source)
GSL	GNU Scientific Library (GNU / Open Source)
Intel MKL	Intel's Math routines for their x86 CPUs. (Freeware/Closed Source)
ARM PL	ARM Processor Performance Libraries (Freeware/Closed Source)
NVBLAS	NVIDIA BLAS, optimized for CUDA / GPU execution
cuSOLVER	NVIDIA LAPACK, optimized for CUDA / GPU execution
...	

- ▶ Note vendor implementations for specific processors / architectures: target efficient operation on these chips
- ▶ ATLAS is notable as on install, runs a series of benchmarks to set parameters in BLAS that give the best performance (Automatically Tuned)

Sparse Matrices

- ▶ Many problems that arise in HPC involve matrices with many Zero elements
- ▶ Referred to as **sparse matrices** especially when stored in data structures that reduce their size
- ▶ Contrast with **dense matrices** which we have assumed so far
- ▶ Example: LINKS matrix in Page Rank could benefit a lot from sparse storage
- ▶ BLAS / LAPACK deal with dense matrices, Sparse Matrices are their own beast



Example of a sparsity pattern in a large matrix: **black** indicates non-zero element, white is zero

Data Structures to Store Sparse Matrices

- ▶ Dense Matrices
 - ▶ Use $NROW \times NCOL$ space
 - ▶ Provide $O(1)$ lookup for element (i, j)
- ▶ Sparse Matrix formats
 - ▶ Use $O(NNZ)$ storage: NNZ is the **Number of NonZeros**
 - ▶ Provide worse than $O(1)$ lookup for element (i, j)
 - ▶ Store only elements that are nonzero, assume if an index is not present that it is zero
- ▶ Storage savings for sparse formats can be significant when matrix is mostly zeros ($NNZ \ll NROW \times NCOL$)

Octave Example of Sparse Matrix Storage

- ▶ Octave is an open-source scientific computing environment, mostly compatible with Matlab
- ▶ Has built-in support for sparse matrices
- ▶ Uses the Compressed Sparse Column format (CSC) internally
- ▶ Makes it easy to show space savings

```
octave:2> A=[
  [10  20  0  0  0  0]
  [ 0  30  0  40  0  0]
  [ 0  0  50  60  70  0]
  [ 0  0  0  0  0  80]];
```

```
octave:3> As = sparse(A);
```

```
octave:4> B = [ A zeros(4,94);
               zeros(96, 100)];
```

```
octave:5> Bs=sparse(B);
```

```
octave:6> whos
```

Variables visible from the current scope:

variables in scope: top scope

Attr	Name	Size	Bytes	Class
====	====	====	=====	=====
	A	4x6	192	double
	As	4x6	184	double
	B	100x100	80000	double
	Bs	100x100	936	double
	ans	1x1	8	double

Total is 20049 elements using 81320 bytes 11

Coordinate Format (COO)

- ▶ Store (row,col,val) for all non-zero elements
- ▶ Values/Indices stored in separate arrays
- ▶ **Justify** operational complexities:
 1. Space requirement is $3*NNZ$
 2. Finding element (i,j) is $O(\log(NNZ))$
 3. Transpose is $O(NNZ)$

	0	1	2	3	4	5	
0	[10	20	0	0	0	0]	SAMPLE DENSE MATRIX
1	[0	30	0	40	0	0]	
2	[0	0	50	60	70	0]	
3	[0	0	0	0	0	80]	

NNZ = 8, NROW=4, NCOL=6

		0	1	2	3	4	5	6	7		
VALUES	= [10	20	30	40	50	60	70	80]	COO DATA ARRAYS
ROW_INDEX	= [0	0	1	1	2	2	2	3]	
COL_INDEX	= [0	1	1	3	2	3	4	5]	

Compressed Sparse Row Format (CSR)

- ▶ Save space by “compressing” rows : store only row start positions
- ▶ Length of Row I is $\text{ROW_START}[I+1] - \text{ROW_START}[I]$
- ▶ **Justify** operational complexities:
 1. Space requirement is $2 * \text{NNZ} + \text{NROW} + 1$
 2. Finding element (i, j) is close to $O(1)$
 3. Transpose is $O(\text{NNZ})$

	0	1	2	3	4	5	
0	[10	20	0	0	0	0]	SAMPLE DENSE MATRIX
1	[0	30	0	40	0	0]	
2	[0	0	50	60	70	0]	
3	[0	0	0	0	0	80]	

$\text{NNZ} = 8, \text{NROW}=4, \text{NCOL}=6$

		0	1	2	3	4	5	6	7			
VALUES	=	[10	20	30	40	50	60	70	80]	CSR DATA ARRAYS
COL_INDEX	=	[0	1	1	3	2	3	4	5]	
ROW_START	=	[0	2	4	7	8]				

^ Extra element = NNZ

Algorithms for Sparse Matrices

- ▶ BLAS / LAPACK do NOT work for sparse matrices
- ▶ Must utilize different algorithms
- ▶ Less standardization around sparse matrices but libraries / vectors exist
- ▶ Sparse BLAS spec exists but fewer implementations
- ▶ Factorization like in LAPACK require significant algorithm changes to work for sparse matrices
- ▶ Prof. Yousef Saad is our local expert in this area and is worth chatting up if you want to learn more

Related Materials

Stephen Boyd's EE364 Linear Algebra Overview

[https://stanford.edu/class/ee364b/lectures/
num-lin-alg-software.pdf](https://stanford.edu/class/ee364b/lectures/num-lin-alg-software.pdf)

- ▶ Discusses many of the same items we talked about here
- ▶ Focus is on optimization problems like linear programming
- ▶ Very similar considerations

CSCI 5304 - Computational Aspects of Matrix Theory

- ▶ Great course on doing linear algebra for scientific problems
- ▶ Some coverage of BLAS/LAPACK and sparse matrices
- ▶ Often taught by Prof Saad who is a resident expert on all things Matrix/Sparse