

Parallel Dense Matrix Algorithms

Chris Kauffman

*Last Updated:
Sun Oct 3 10:02:31 PM CDT 2021*

Logistics

Assignments

- ▶ A1 grading in progress, look for results in the next day
- ▶ A2 will go up around ~~Friday~~ Thursday and feature MPI Coding

Today

- ▶ Overview of some Linear Algebra Libraries
- ▶ Mini-Exam 1

Recall Matrix Transpose

- ▶ Common operation on matrices is a **transpose** notated A^T
- ▶ Interchanges rows/columns of A :
 $a_{ij} \rightarrow a_{ji}$
- ▶ Diagonal elements stay the same
- ▶ Algorithms that perform operations on A can often be performed on A^T without re-arranging A - **how?**
Hint: consider summing rows of A vs summing rows of A^T

Original matrix A

0	5	10	15
20	25	30	35
40	45	50	55
60	65	70	75

transpose(A)

0	20	40	60
5	25	45	65
10	30	50	70
15	35	55	75

Exercise: Matrix Partitioning Across Processors

Row Partition				Column Partition				Block Partition				Proc Location	
00	01	02	03	00	01	02	03	00	01	02	03		P0 / P00
10	11	12	13	10	11	12	13	10	11	12	13		P1 / P01
20	21	22	23	20	21	22	23	20	21	22	23		P2 / P10
30	31	32	33	30	31	32	33	30	31	32	33		P3 / P11

- ▶ Recall several ways to partition matrices across processors
- ▶ Diagram shows these
 - ▶ Entry ij may be an individual element OR...
 - ▶ Entry ij may be a **Block**: ex. Block (2,3) is the submatrix from rows 200-299 and cols 300-399
- ▶ Assume **square** matrices : $\#rows = \#cols$
- ▶ Common to multiply to compute product: $C = A \times B$
- ▶ Ideal partitioning for A and B ?
- ▶ Ideal partitioning for $C = A^T \times B$
- ▶ Ideal partitioning for $C = A \times B^T$

Answers: Matrix Partitioning Across Processors

Row Partition				Column Partition				Block Partition				Proc Location	
00	01	02	03	00	01	02	03	00	01	02	03		P0 / P00
10	11	12	13	10	11	12	13	10	11	12	13		P1 / P01
20	21	22	23	20	21	22	23	20	21	22	23		P2 / P10
30	31	32	33	30	31	32	33	30	31	32	33		P3 / P11

- ▶ $C = A \times B$
 - ▶ Ideally A is row-partitioned, B is column partitioned
 - ▶ Then block-partitioned C could be computed w/o communication
- ▶ $C = A^T \times B$
 - ▶ Ideally A and B column-partitioned
- ▶ $C = A \times B^T$
 - ▶ Ideally A and B row-partitioned
- ▶ Block-partitioning often used: not ideal for any version but less communication required when both A and A^T will be used


Naive Parallel Dense Multiplication: Overview

- ▶ Most applications start with block-partitioned matrices
- ▶ To compute Matrix-Matrix multiply, procs must have (eventually) multiply full rows by full columns to compute an output block
- ▶ Naive method: each Proc stores full rows/columns needed for it to independently compute output block which it stores

Naive Parallel Dense Multiplication: Demo


1. Initial data layout:
each Proc holds a block
of A, B, and C
respectively. Processors
are arranged in a logical
grid that reflects their
initial data.

In subsequent steps,
received / computed
data is bolded.




3. Each proc participates
in an All-to-All sharing of
data for the Column it is
in.

This leaves each row
with complete columns
as well.




2. Each proc participates
in an All-to-All sharing of
data for the Row it is in.

This leaves each proc
with entire rows of A.



4. Each Proc now has a
unique set of complete
Rows and Columns and
can independently
compute a block of
output matrix C through
block multiplication.



Exercise: Analysis of Naive Dense Mult.

Assumptions

- ▶ Matrices A and B are size $N \times N$ so N^2 elements
- ▶ P processors with block partitioning: initially N^2/\sqrt{P} elements of A, B on each proc (assume P is a perfect square)
- ▶ Simplified communication cost for All-to-All on a Ring:

$$t_{comm} = (p - 1)(t_s + t_w M)$$

with p #procs in ring, t_s comm startup time, t_w per word transfer rate, M message size.

Questions

1. What is **communication cost** of this algorithm?
2. How much **time** does the final **block matrix multiply** take?
3. What is the **memory requirement for each proc**?
4. What do you see as the **biggest disadvantage** for this algorithm?

Answers: Analysis of Naive Dense Mult.

1. What is **communication cost** of this algorithm?

- ▶ #Procs in rows/cols is $\sqrt{P} \sim$ ring size
- ▶ $M = N^2/\sqrt{P}$: message size is num elements on each proc
- ▶ 2 All-to-All shares : 1 for rows, 1 for cols

$$t_{comm} = 2(\sqrt{P} - 1) \times (t_s + t_w(N^2/\sqrt{P}))$$

2. What is the **memory requirement for each proc**? E.g. how many submatrices of A,B are on each proc?

- ▶ Full rows/cols on each proc
- ▶ Requires \sqrt{P} submatrices for each Proc

3. How much **time** does the final **block matrix multiply** take?

- ▶ Each proc has \sqrt{P} submats of A,B to multiply
- ▶ Each submat is size N/\sqrt{P} with size s requiring $O(s^3)$ opts

$$t_{mult} = O((\sqrt{P}) \times ((N/\sqrt{P})^3)) = O(N^3/P)$$

4. What do you see as the **biggest disadvantage** for this algorithm?

- ▶ Major: The need to store \sqrt{P} sub matrices on all procs may be prohibitive
- ▶ Minor: Not much chance to overlap communication / computation in the algorithm

Cannon's Algorithm

- ▶ Proposed in Lynn Elliot Carter's 1969 thesis
- ▶ Target was very small parallel machines implementing a **Kalman Filter** algorithm in hardware
- ▶ "Communication" happening between small Procs with data in registers
- ▶ Scales nicely to large machines and overcomes the large memory requirement of the Naive Mat-Mult Algorithm

A CELLULAR COMPUTER TO IMPLEMENT

THE KALMAN FILTER ALGORITHM

by

LYNN ELLIOT CANNON

By the conventional definition of matrix product, if A is multiplied by B, the result, call it C, is given by

$$C = AxB = \begin{bmatrix} a_1b_1+a_2b_2+a_3b_3 & a_1b_4+a_2b_5+a_3b_6 & a_1b_7+a_2b_8+a_3b_9 \\ a_4b_1+a_5b_2+a_6b_3 & a_4b_4+a_5b_5+a_6b_6 & a_4b_7+a_5b_8+a_6b_9 \\ a_7b_1+a_8b_2+a_9b_3 & a_7b_4+a_8b_5+a_9b_6 & a_7b_7+a_8b_8+a_9b_9 \end{bmatrix}.$$

The symmetry of this product can be seen by comparing the ij^{th} element with the ji^{th} element and noticing that one is obtained from

-24-

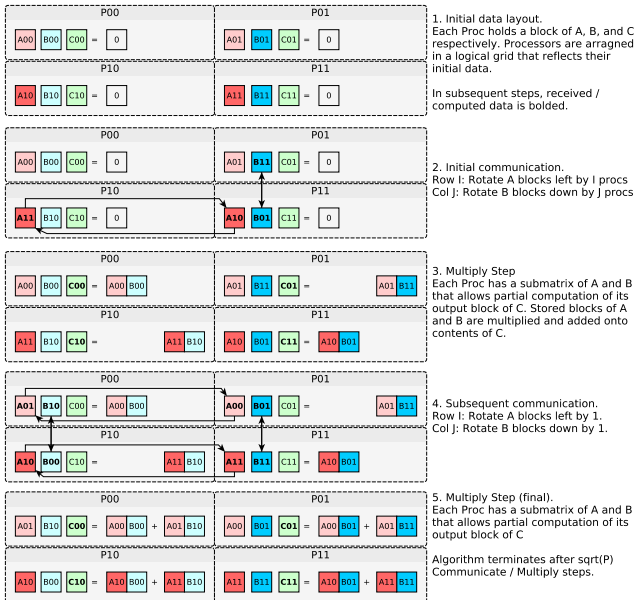
- A.
 1. The first row of A is left alone.
 2. The second row of A is shifted left one column.
 3. The third row of A is shifted left two columns.
 (Note, in general the i^{th} row of A is shifted left $i-1$ columns for $i = 1, \dots, n$).
- B.
 1. The first column of B is left alone.
 2. The second column of B is shifted up one row.
 3. The third column of B is shifted up two rows.
 (Note, in general the j^{th} column of B is shifted up $j-1$ rows for $j = 1, \dots, n$).

Once the registers have been shifted the multiplication pr

Demo

Cannon's Algorithm for Parallel Matrix Multiply: Demo for 2x2 block arrangement

$$\begin{bmatrix} C00 & C01 \\ C10 & C11 \end{bmatrix} = \begin{bmatrix} A00 & A01 \\ A10 & A11 \end{bmatrix} \times \begin{bmatrix} B00 & B01 \\ B10 & B11 \end{bmatrix}$$



Cannon's Algorithm Pseudocode

```

procedure Cannon(i , j, N){
  PE(i,j) has blocks A1=A(i,j) and B1=B(i,j)
  Allocate space A2, B2, Cij sized as A1

```

```

  rsend_col = (j-i+N % N)
  rrecv_col = (j+i+N % N)
  doboth send A1 to PE(i,rsend_col)
    recv A2 from PE(i,rsend_col)

```

```

  csend_row = (i-j+N % N)
  crecv_row = (i+j+N % N)
  doboth send B1 to PE(csend_row,j)
    recv B2 from PE(csend_row,j)


```

```


for(i=1 to N){
  swap A1,A2 and B1,B2
  Cij += A1 * B1

  doboth send A1 to PE(i, j-1+N % N)
    recv A2 from PE(i, j+1+N % N)
  doboth send B1 to PE(i-1+N % N, j)
    recv B2 from PE(i+1+N % N, j)
}
}


```




(a) Initial alignment of A




(b) Initial alignment of B




(c) A and B after initial alignment



(d) Submatrix locations after first shift



(e) Submatrix locations after second shift



(f) Submatrix locations after third shift

Figure 8.3 The communication steps in Cannon's algorithm on 16 processes.