

Parallel Sorting

Chris Kauffman

*Last Updated:
Wed Oct 20 09:32:40 AM CDT 2021*

Logistics

Today

- ▶ Parallel Sorting:
Quicksort
- ▶ Wed: Mini-exam 2

Reading: Grama Ch 9

- ▶ Sorting
- ▶ Focus on 9.4: Quicksort

Quick Review

- ▶ What is Amdahl's law? What does it say about the speedup achievable by parallel programs?
- ▶ How does one calculate the following for a parallel algorithm
 - ▶ S : Speedup
 - ▶ E : Efficiency
 - ▶ C : Cost
- ▶ How does the Efficiency of a parallel usually change if the number of processors P increases but the problem size stays the same? If number of procs stays the same but problem size increases?
- ▶ What was the major benefit that Cannon's Algorithm provided over a naive implementation of parallel matrix multiply?
- ▶ In broad strokes, how was the LU factorization parallelized?

Sorting

- ▶ Much loved computation problem
- ▶ What is the best complexity of general purpose (comparison-based) sorting algorithms?
- ▶ What are some algorithms which have this complexity?
- ▶ What are some other sorting algorithms which aren't so hot?
- ▶ What issues need to be addressed to parallelize any sorting algorithm?

Partition and Quicksort

- ▶ Quicksort has $O(N \log N)$ average complexity
- ▶ In-place, low overhead sorting, recursive

Partition

- ▶ Partition: select pivot value
- ▶ Rearrange elements so
 - ▶ Left array is \leq pivot
 - ▶ Right array is $>$ pivot
 - ▶ pivot is in “middle”

```
// A is an array, lo/hi are
// inclusive boundaries
algorithm partition(A, lo, hi) is
    pivot := A[hi]
    boundary := lo
    for j := lo to hi do
        if A[j] <= pivot then
            swap A[boundary] with A[j]
            boundary++
    swap A[boundary] with A[hi]
    return boundary
```

Quicksort

- ▶ Partition into two parts
- ▶ Recurse on both halves
- ▶ Bail out when boundaries
lo/hi cross

```
algorithm quicksort(A, lo, hi) is
    if lo < hi then
        p := partition(A, lo, hi)
        quicksort(A, lo, p - 1)
        quicksort(A, p + 1, hi)
```

Practical Parallel Sorting Setup

- ▶ Input array A of size N is already spread across P processors (no need to scatter)

P0: $A[] = \{ 84 \ 31 \ 21 \ 28 \}$

P1: $A[] = \{ 17 \ 20 \ 24 \ 84 \}$

P2: $A[] = \{ 24 \ 11 \ 31 \ 99 \}$

P3: $A[] = \{ 13 \ 32 \ 26 \ 75 \}$

- ▶ Goal: Numbers sorted across processors. Smallest on P0, next smallest on P1, etc.

P0: $A[] = \{ 11 \ 13 \ 17 \ 20 \}$

P1: $A[] = \{ 21 \ 24 \ 24 \ 26 \}$

P2: $A[] = \{ 28 \ 31 \ 32 \ 33 \}$

P3: $A[] = \{ 75 \ 84 \ 84 \ 99 \}$

- ▶ Want to use P processors as effectively as possible
- ▶ Favor bulk communication over many small messages

Exercise: Parallel Quicksort

- ▶ Find a way to parallelize quicksort
- ▶ **Hint:** The last step is each processor sorting its own data using a serial algorithm. Try to arrange data so this is possible.

START:

P0: A[] = { 84 32 21 28 }

P1: A[] = { 17 20 25 85 }

P2: A[] = { 24 11 31 99 }

P3: A[] = { 13 32 26 75 }

GOAL

P0: A[] = { 11 13 17 20 }

P1: A[] = { 21 24 25 26 }

P2: A[] = { 28 31 32 33 }

P3: A[] = { 75 84 85 99 }

SERIAL ALGORITHM

```
algorithm quicksort(A, lo, hi) is
  if lo < hi then
    p := partition(A, lo, hi)
    quicksort(A, lo, p - 1)
    quicksort(A, p + 1, hi)
```

```
algorithm partition(A, lo, hi) is
  pivot := A[hi]
  boundary := lo
  for j := lo to hi - 1 do
    if A[j] <= pivot then
      swap A[boundary] with A[j]
      boundary++
  swap A[i] with A[hi]
  return boundary
```

Answers: Parallel Quicksort Ideas 1 / 2

A[] = { 84 32 21 11 | 17 20 25 85 | 24 28 31 99 | 13 33 26 75 }
P0 P1 P2 P3

Partition(pivot=26) on each processor

A[] = { 21 11 84 32 | 17 20 25 85 | 24 28 31 99 | 13 26 33 75 }
Boundary: ^ ^ ^ ^
Counts: P0: 2 P1: 3 P2: 1 P3: 2

Calculate which data goes where

A[] = { 21 11 84 32 | 17 20 25 85 | 24 28 31 99 | 13 26 33 75 }
P0 P0 P2 P2 P0 P0 P1 P2 P1 P2 P3 P3 P1 P1 P3 P3

Re-arrange so values ≤ 26 on P0 and P1, > 26 on P2 and P3

A[] = { 21 11 17 20 | 25 24 13 26 | 84 32 85 28 | 31 99 33 75 }
P0 P1 P2 P3

Split the world: 2 groups

A[] = { 21 11 17 20 | 25 24 13 26 } | { 84 32 85 28 | 31 99 33 75 }
P0 P1 P2 P3

Answers: Parallel Quicksort Ideas 2 / 2

Each half partitions on different pivot value

P0-P1: Partition(pivot=20) P2-P3: Partition(pivot=33)

A[] = { 11 17 20 21 | 13 25 24 26 } | { 28 32 84 85 | 31 33 99 75 }

Boundary: ^ ^ ^ ^

Counts: P0: 3 P1: 1 P2: 2 P3: 2

Calculate which data goes where

A[] = { 11 17 20 21 | 13 25 24 26 } | { 28 32 84 85 | 31 33 99 75 }

P0 P0 P0 P1 P0 P1 P1 P1 P2 P2 P3 P3 P2 P2 P3 P3

Re-arrange values to proper processors

A[] = { 11 17 20 13 | 21 25 24 25 } | { 28 32 31 33 | 84 85 99 75 }

P0 P1 P2 P3

Split the world: 4 groups

A[] = { 11 17 20 13 } | { 21 25 24 25 } | { 28 32 31 33 } | { 84 85 99 75 }

P0 P1 P2 P3

4 groups == 4 processors, all processors sort locally

A[] = { 11 13 17 20 } | { 21 24 25 25 } | { 28 31 32 33 } | { 75 84 85 99 }

P0 P1 P2 P3

Quicksort Difficulties

Communication

- ▶ Determine which data go to which processors, how many send/receives are required
- ▶ Opportunity for **all-to-all communications** in MPI

Recurring

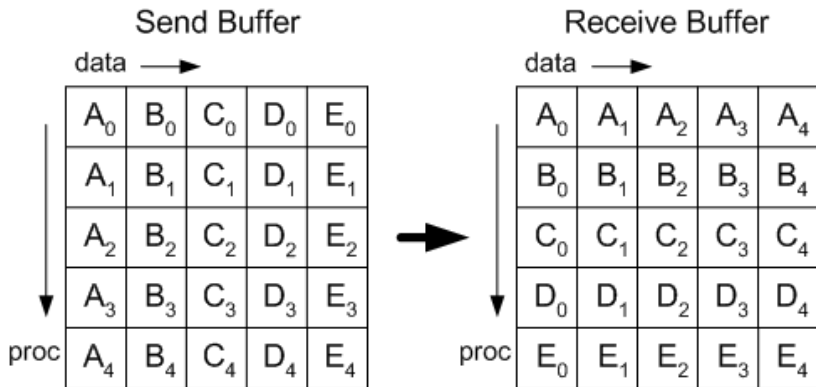
- ▶ Recursive step of algorithm requires smaller “worlds”
- ▶ Use MPI's **communicator splitting** capability

Pivot Value Selection

- ▶ In example, pivot values were cherry-picked to get even distribution of data among processors
- ▶ A bad pivot splits data unevenly, is annoying for serial Quicksort, shaves off processors in parallel quicksort destroying efficiency

All-to-All Personalized Communication

All-to-all personalized communication: like every processor scattering to every other processor.



Source: Cornell University Center for Advanced Computing

MPI_Alltoall

- ▶ Standard version: every processor gets a slice of sendbuf, same sized data
- ▶ Vector version allows different sized slices (appropriate for quicksort)

```
int MPI_Alltoall(  
    void *sendbuf, int sendcount, MPI_Datatype sendtype,  
    void *recvbuf, int recvcount, MPI_Datatype recvtype,  
    MPI_Comm comm);
```

```
int MPI_Alltoallv(  
    void *sendbuf, int sendcounts[], int sdispls[], MPI_Datatype sendtype,  
    void *recvbuf, int recvcounts[], int rdispls[], MPI_Datatype recvtype,  
    MPI_Comm comm);
```

Exercise: Redistribution during Quicksort

- ▶ After partition, procs must redistribute data
- ▶ Use All-Gather so that all procs have the following table

Element Count vs Pivot			
Proc	<=	>	
P0	2	2	
P1	3	1	
P2	1	3	
P3	2	2	

Want each Proc to calculate its own Count/Displ in this table:

P#		P0	P1	P2	P3		P0	P1	P2	P3	P#
P0	RecvCount	2	2	0	0	SendCount	2	0	2	0	P0
P1		0	1	1	2		2	1	1	0	P1
P2		2	1	1	0		0	1	1	2	P2
P3		0	0	2	2		0	2	0	2	P3
P0	RecvDispl	0	2	4	4	SendDispl	0	0	2	0	P0
P1		0	0	1	2		0	2	3	4	P1
P2		0	2	3	4		0	0	1	2	P2
P3		0	0	0	2		0	0	2	2	P3

- ▶ Describe the process of calculating RecvCount
- ▶ Given RecvCount, how can one calculate RecvDispl

Answers: Redistribution during Quicksort


- ▶ RecvCount can be calculate through an iterative process
- ▶ Compute the **prefix sum** of for the table

Element Count vs Pivot				PS[]: PREFIX SUMS			
Proc	<=		>	Proc	<=		>
-----+-----+---				-----+-----+---			
P0	2		2	P0	2		2
P1	3		1	P1	5		3
P2	1		3	P2	6		6
P3	2		2	P3	8		8

- ▶ Know each proc must send/receive $N / P = 4$ elements
- ▶ Procs receiving \leq pivot, proc # i , scan column 0 for
 - ▶ First partner is proc F where $PS[F,0] \leq 4*i$
 - ▶ Last partner is proc L where $PS[L,0] \geq 4*(i+1)$
- ▶ Procs receiving $>$ pivot, proc # i , scan column 1 for
 - ▶ First partner is proc F where $PS[F,1] \leq 4*(i-2)$
 - ▶ Last partner is proc L where $PS[L,1] \geq 4*(i-2+1)$
- ▶ Actual code will need to do some arithmetic (e.g. P1 receives 1 element from itself)
- ▶ RecvDispl is the **prefix sum** of RecvCount

Prefix Sums / Scan

Prefix Sums or Prefix Scans are supported in parallel via MPI



```
int MPI_Scan(const void *sendbuf, void *recvbuf, int count,  
             MPI_Datatype datatype, MPI_Op op, MPI_Comm comm)
```

- ▶ Similar to reduction but only add on values from procs \leq `proc_id`
- ▶ `op` can be sum/max/min/etc.
- ▶ In simple Quicksort implementations, **don't use parallel prefix scan** as this does not yield enough info to calculate send/receive partners

Overall Flow

1. Pivot selection (open question how to do this right)
2. Broadcast of pivot value
3. Each processor partition's its data
4. All-gather to get element/pivot counts
5. Calculate send/receives
6. Redistribute data via `MPI_Alltoallv()`
7. And then...

Splitting the World

```
int MPI_Comm_split(MPI_Comm comm, int color, int key,  
                  MPI_Comm *newcomm);
```

- ▶ `comm` is the old communicator (start with `MPI_COMM_WORLD`)
- ▶ `color` is which sub-comm to go into
 - ▶ Colors 0,1 splits into 2 communicators
 - ▶ Colors 0,1,2,3 splits into 4 communicators
 - ▶ Etc.
- ▶ `key` establishes rank in new sub-comm, usually `proc_id`
- ▶ `newcomm` is filled in with a new communicator
- ▶ Examine `04-mpi-code/comm_split.c`
- ▶ In Quicksort, new comm is different for lower/upper procs

Exercise: Pivot Selection

- ▶ So far have assumed a “good” pivot can be found
- ▶ Pivot evenly splits $N/2$ data, half to lower # processors, half to upper

Discuss the following questions with a neighbor

1. What if the pivot is poorly selected? E.g. $1/4$ below pivot, $3/4$ above? Could the algorithm adapt?
2. How could one avoid a bad pivot? Discuss a some strategies
3. Is there a way to avoid recusing entirely?

Answers: Pivot Selection 1/2

Discuss the following questions with a neighbor

1. What if the pivot is poorly selected? E.g. $1/4$ below pivot, $3/4$ above? Could the algorithm adapt?

With some additional computation, can split the world unevenly: $1/4$ procs assigned to “low” numbers, $3/4$ to “high” numbers. Still broken if a tiny fraction of the array is lower/higher than the pivot: should just try another pivot at that point or use a scheme that prevents poor pivot selection.

2. How could one avoid a bad pivot? Discuss a some strategies
Lots of these exist, some mentioned in the textbook such as having a randomly selected processor compute its median and broadcast it as the pivot (main text of Grama) or have processors sample random elements, perform All-Gather, then compute a median as the common pivot (Grama Exercise 9.21).

Answers: Pivot Selection 1/2

3. Is there a way to avoid recusing entirely, e.g. single multiway pivot?

Gramma Exercise 9.20 explores this:

- ▶ *Each proc samples elements, often around $\log(N)$ elements, and procs perform an All-Gather*
- ▶ *All procs use common sample to select $P - 1$ common pivots.*
- ▶ *Elements between pivots are sent directly to final destination procs in an All-to-All communication.*
- ▶ *Local sorting commences.*

Exercise: Odd-Even Sort

- ▶ Variant of bubble sort which splits bubbling into odd/even phases
- ▶ $O(N^2)$ complexity of serial algorithm
- ▶ There is potential for parallelism here: **what is it?**
 - ▶ Consider simple case where each $P = N$: each proc hold a single number
 - ▶ What can be parallelized and how?

```
ODD_EVEN_SORT(A[]) {  
    N = length(A[])  
    if(r is even){  
        for(i=0; i<N-1; i+=2){  
            compare_exchange(A, i, i+1);  
        }  
    }  
    if(r is odd){  
        for(i=1; i<N-1; i+=2){  
            compare_exchange(A, i, i+1);  
        }  
    }  
}  
  
COMPARE_EXCHANGE(A[], i, j){  
    if(A[i] > A[j]){  
        temp = A[i]  
        A[i] = A[j]  
        A[j] = temp  
    }  
}
```

Answers: Odd-Even Sort

- ▶ There is potential for parallelism here: **what is it?**
- ▶ Consider simple case where each $P = N$: each proc hold a single number
- ▶ What can be parallelized and how?
 - ▶ *The inner loops of `compare_exchange()` can be executed in parallel as it involves communication between 2 procs to potentially exchange elements but only with a single partner.*
 - ▶ *Even iterations, lower evens exchange with higher odds*
 - ▶ *Odd iterations lower odds exchange with higher evens*
 - ▶ *Exchange can be done via a Send/Receive of elements and then “keeping” the appropriate element, min on lower proc, max on higher proc*

Odd-Even Sort in Practice

- ▶ As before, unrealistic to have $P = N$, rather each proc holds N/P elements of the array $A[]$
- ▶ `COMPARE_EXCHANGE()` becomes `COMPARE_SPLIT()`




Figure 9.2 A compare-split operation. Each process sends its block of size n/p to the other process. Each process merges the received block with its own block and retains only the appropriate half of the merged block. In this example, process P_i retains the smaller elements and process P_j retains the larger elements.

Analysis of Odd-Even Sort

- ▶ Initially all procs sort their local array: $O(\frac{N}{P} \log \frac{N}{P})$
- ▶ Conduct P Outer iterations of ODD_EVEN_SORT()
- ▶ Each odd/even inner loop is done in parallel by all procs communicating with a neighbor
- ▶ Neighbor procs exchange arrays: $O(t_s + t_w \frac{N}{P})$
- ▶ Each proc then performs a compare/split: $O(\frac{2N}{P})$
- ▶ Overall complexity of parallel algorithm:

$$\begin{aligned} T_{par} &= O\left(\frac{N}{P} \log \frac{N}{P}\right) + P \times \left(O\left(t_s + t_w \frac{N}{P}\right) + O\left(\frac{2N}{P}\right)\right) \\ &= O\left(\frac{N}{P} \log \frac{N}{P}\right) + O(N) + O(N) \end{aligned}$$

Isoefficiency? : Reported in textbook as $O(P2^P)$, linear increase in P requires exponential increase in problem size to maintain efficiency. Verifying this is a good exercise.

Sorting Extras

Odd-Even Sort to Shell Sort

- ▶ Allowing bigger “moves” in odd-even sort can improve practical efficiency of algorithm
- ▶ Shell Sort provides a mechanism for this: neighbors selected according to a “gap” scheme, less known sort with yet mysterious complexity analysis

Sorting Hardware

- ▶ Grama Ch 9.1 discusses Sorting networks, specialized hardware which can implement sorting
- ▶ With N processors, can implement Bitonic Sort in a sorting network and achieve $T_{par} = O(\log^2 N)$
- ▶ Hardware that implements sorting networks is not common but...
- ▶ Can do this in GPUs and may revisit this algorithm when we do CUDA programming