Shared Memory Architectures

Chris Kauffman

Last Updated:
Thu Mar 17 09:03:15 AM CDT 2022

Assignments + Grading

- ▶ Will combine A2 and A3 to do an MPI and OpenMP versions of K-Means clustering
- ▶ Post later this week - again, apologies for the delay
- ▶ Mini-Exam 1 and 2 grades will be up by next week

Today

- ▶ Shared Memory Architecture Theory + Practicalities
- ▶ Cache Performance Effects
- ▶ Next Week: PThreads + OpenMP for shared memory machines

Reading

- ▶ Grama 2.4.1 (PRAM), 2.4.6 (cache)
- ▶ Gram 7.1-9 (PThreads)
- ▶ Grama 7.10 (OpenMP)
- ▶ OpenMP Tutorial at Laurence Livermore

PRAM: Parallel Random Access *Machine*, Grama Ch 2.4.1

RAM: Random Access *Machine*

- ▶ An unfortunate name as every other use of "RAM" is random access memory
- ▶ Single CPU attached to random access memory
- ▶ Simplistic model for a real machine: CPU reads memory, performs operations in registers, writes to memory, repeats

PRAM: Parallel Extension to RAM

- ▶ Again, theoretical model for a real parallel machine
- ▶ Multiple CPUs attached to memory, share clock but can execute different instructions
- ▶ Must clarify behavior of PRAM machine that is not possible in single PE situation: how are conflicts between processors resolve for **simultaneous memory access**

Theoretical Flavors of PRAM

Exclusive-Read, Exclusive-Write (EREW) PRAM

Multiple CPUs cannot touch same memory at all. No parallelism possible for reads / writes of the same memory location.

Concurrent-Read, Exclusive-Write (CREW) PRAM

Multiple CPUs may read same memory location at same time. Writes to same location must be resolved.

Exclusive-Read, Concurrent-Write (ERCW) PRAM

Multiple write accesses are allowed to a memory location, but multiple read accesses are serialized. (This is just weird)

Concurrent-Read, Concurrent-Write (CRCW) PRAM

Simultaenous Read AND Write of the same memory location: the most "powerful" PRAM model for some definition of "power". *Q: What else must be specified for the xxCW models?*

Resolution Schemes for Concurrent Reads/Writes

A: How concurrent writes resolve.

- ▶ **Common**: concurrent writes are allowed if all the values that the processors are attempting to write are identical.
- ▶ **Arbitrary:** an arbitrary processor is allowed to proceed with the write operation and the rest fail.
- ▶ **Priority:** procs are organized in a predefined prioritized list; proc with the highest priority writes its value while others fail.
- ▶ **Sum:** the sum of all the quantities is written

Above categories do not resolve concurrent read+write such as:

```
MEM[#1024] is 10
P0 reads MEM[#1024] into REG1
P1 writes 20 to MEM[#1024]
```

Proper treatments of PRAM specify results for this such as

- ▶ All Reads resolve first, then Writes resolve
- ▶ Concurrent Reads/Writes occur in arbitrary order
- ▶ etc.

However, we'll proceed with more practical matters

Pros and Cons of PRAM

Why the PRAM Model?

- ▶ It's simple
- ▶ Much literature devoted to studying benefits of algorithms under different variants *(e.g. Parallel Array Sum with CRCW + summing on concurrent write)*
- ▶ Has significant theoretical importance

Why Not PRAM

- ▶ No real, general machines currently implement any PRAM models
- ▶ Seen some references that GPUs might sort of implement but would require some more work
- ▶ Conclusions one might draw about "good" algorithms is skewed *(e.g. multi-core machines do NOT behave as a CRCW-summing machine; far from it)*

Exercise: Recall the Memory Cache

- Parallel programs are driven towards performance
- Optimize serial performance first: requires understanding of the memory hierarchy

Questions

From your computer architecture experience…

- Describe a **memory cache** and why most CPUs have several layers of them
- Give an example of "strange" cache effects where similar algorithms have very different performance
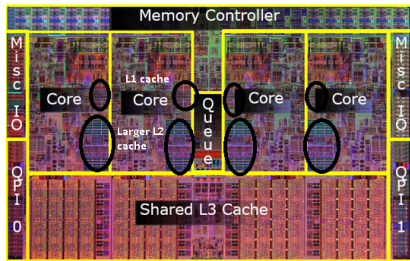
Diagrams of Memory Interface and Cache Levels
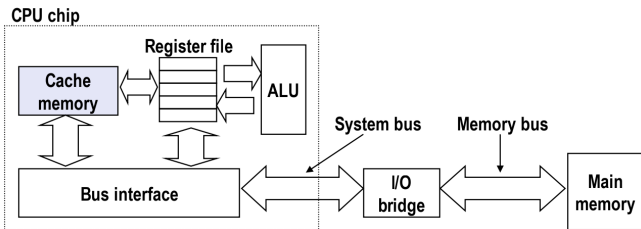
Source: Bryant/O'Hallaron CS:APP 3rd Ed.
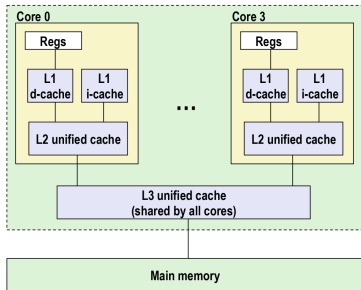
Bryant and O'Hallaron, Computer Systems: A Programmer's Perspective, Third Edition

Source: SO "Where exactly L1, L2 and L3 Caches located in computer?"

Numbers Everyone (in Computing) Should Know

Edited Excerpt of Jeff Dean's talk on data centers.

Operation	Time (ns)
L1 cache reference	0.5
Branch mispredict	5
L2 cache reference	7
Mutex lock/unlock	100
Main memory reference	100
Compress 1K bytes with Zippy	10,000
Send 2K bytes over 1 Gbps network	20,000
Read 1 MB sequentially from memory	250,000
Round trip within same datacenter	500,000
Disk seek	10,000,000
Read 1 MB sequentially from network	10,000,000
Read 1 MB sequentially from disk	30,000,000
Send packet CA->Netherlands->CA	150,000,000

Numbers are likely out of date now but scales are worth knowing
and explain why Cache is useful

Matrix Summing Examples

Sum R

```
double X[N][N]; // N by N mat
...
sum = 0;
for(i=0; i<N; i++){
  for(j=0; j<N; j++){
    sum += X[i][j]
  }
}
```

Sum C

```
double X[N][N]; // N by N mat
...
sum = 0;
for(j=0; j<N; j++){
  for(i=0; i<N; i++){
    sum += X[i][j]
  }
}
```

▶ What's the Big O complexity of each?

▶ What happens with cache?

▶ Will one be faster than the other?

Cache Affects Performance

As measured by hardware counters using linux's `perf` on

```
model name : Intel(R) Core(TM) i5-2400 CPU @ 3.10GHz
cache size : 6144 KB
```

with

```
perf stat $opts java MatrixSums 8000 4000 row
perf stat $opts java MatrixSums 8000 4000 col
```

Measurement	row	col
cycles	3,507,364,715	5,605,621,966
instructions	2,353,887,029	2,543,165,478
L1-dcache-loads	527,694,054	561,540,169
L1-dcache-load-misses	25,638,014	122,663,199
Runtime (seconds)	1.001	1.620

L1 data cache load misses

- ▶ Row: 25K/548K = 4% main memory access
- ▶ Col: 122/585K = 20% main memory access

Cache Issues in Shard Memory Machines
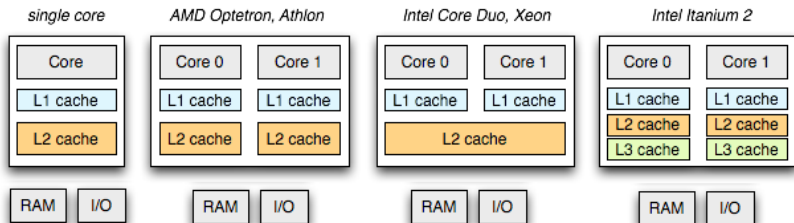
Source: Multi-core, Threads & Message Passing by Ilya Grigorik

Consider the following sequence of operations:

```
                            // MEM[#1024] has value 5
P0: load R1 MEM[#1024]      // slow, populates cache
P0: load R2 MEM[#1024]      // fast, from cache
P0: ADD R1 R1 R2            // R1 is 10
P0: store R1 MEM[#1024]     // cache dirty, MEM[#1024] unchanged
  a short time later
P1: load R3 MEM[#1024]      // read 5 or 10?
```

Illustrates **Cache Coherence** Problem: how do multiple PEs maintain the illusion of a single block of shared memory?

Cache Coherence Protocols

Grama 2.4.6 offers two theoretical protocols to maintain coherence

- **Invalidate**: shared memory written by one PE is marked as invalid in cache of others
- **Update**: shared memory written in by one PE is updated in other PEs

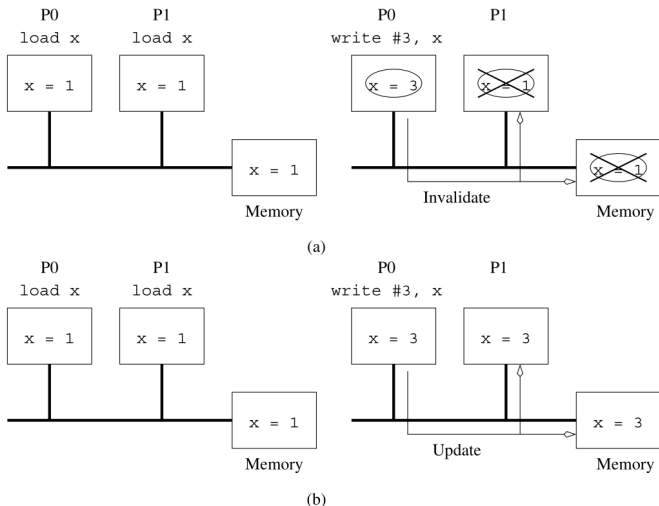# Invalidate and Update Protocol Diagrams

Figure 2.21 Cache coherence in multiprocessor systems: (a) Invalidate protocol; (b) Update protocol for shared variables.

Cache State and Coherence

Shared (S) valid for reads, write changes state

Dirty (D) written by me, must eventually flush to main memory

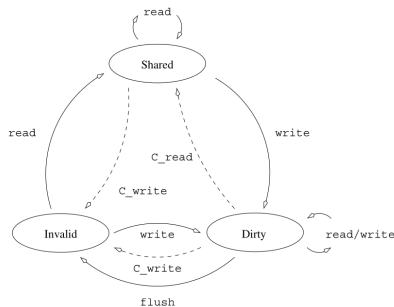Invalid (I) another proc altered it, trigger flush + reload on reading

Figure 2.22 State diagram of a simple three-state coherence protocol.

▶ Each Unit of memory in Cache and DRAM have a S / D / I state, often associated with Cache Lines

▶ NOTE: This is a **simplified scheme** with actual hardware implementations often having several more states and more complexity.

Demonstration

Time	Instruction at Processor 0	Instruction at Processor 1	Variables and their states at Processor 0	Variables and their states at Processor 1	Variables and their states in Global mem.
					x = 5, D y = 12, D
read x			x = 5, S		x = 5, S
	read y			y = 12, S	y = 12, S
x = x + 1			x = 6, D		x = 5, I
	y = y + 1			y = 13, D	y = 12, I
read y			y = 13, S	y = 13, S	y = 13, S
	read x		x = 6, S	x = 6, S	x = 6, S
x = x + y			x = 19, D	x = 6, I	x = 6, I
	y = x + y		y = 13, I	y = 19, D	y = 13, I
x = x + 1			x = 20, D		x = 6, I
	y = y + 1			y = 20, D	y = 13, I

Figure 2.23 Example of parallel program execution with the simple three-state coherence protocol discussed in Section 2.4.6.

The Memory Bus

- ▶ Cache coherence protocols involve communication between procs, obtaining information about changes to memory
- ▶ The **Memory Bus** is the communication channel that enables Procs/Memory chips to "talk" to each other
- ▶ Hardware construct to move data around, usually across wires connected to each Proc and DRAM chip
- ▶ Memory Buses use a communication protocol which includes **device identifiers** so messages about changes are directed to individual PEs or DRAM
- ▶ Bus can get "crowded" if lots of Procs make memory requests
- ▶ All hardware can "see" messages on the bus: allows **snooping** of messages intended for others
- ▶ Example: PE1 sees that PE0 read address #1024 so PE0 knows it may share #1024 now

Diagram of Typical Memory Buses
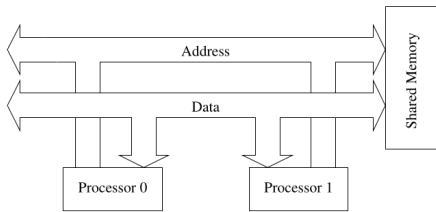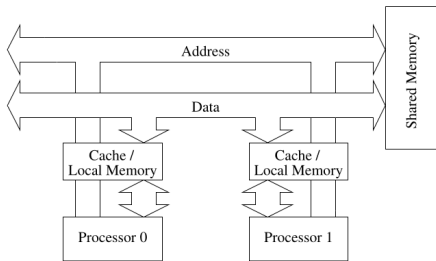
Figure 2.7 Bus-based interconnects (a) with no local caches; (b) with local memory/caches.
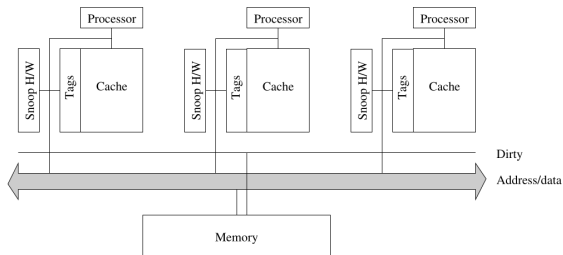
Snoopy Cache

Figure 2.24 A simple snoopy bus based cache coherence system.

Basics

► Additional hardware watches messages on the bus

► Writing to cache invalidates global memory

► Message pertaining to a dirty memory address cause flush, state back to shared

Example

► x in P0 cache Dirty

► x in Global mem Invalid

► P1 reads x
  ► P0 "snoops" request
  ► Flushes x to global mem
  ► P1 can read x from global

► x is now Shared

Different Variable but Same Cache Line $\rightarrow$ Collisions

- ▶ Performance problem: two processors grinding on different but close variables
- ▶ Consider the following program: x,y are adjacent in main memory, likely to share same cache line
- ▶ Proc0 and Proc1 each have own cache, will interfere with one another despite working on different variables
- ▶ Often referred to as **False Sharing** between procs / threads

```
int x=42;
int y=31;
void collide(){
  if(proc_id == 0){
    int i;
    for(i=0; i<1000; i++){
      x = (x+1)*(x+3)/x;
    }
  }
  else{
    int i;
    for(i=0; i<1000; i++){
      y = y/2;
      y = y+2*y;
    }
  }
}
```

Cache Coherence Overall

- ▶ Caches speed up individual processor execution in most cases
- ▶ Coordinating caches across several PEs is complex
- ▶ Requires additional hardware such for Snooping, alternatively Directory-based approach (textbook)
- ▶ Hardware manages most of this but uses techniques that are smack of distributed memory systems
- ▶ To eek out more performance, programmers should be aware of these things when using Thread-based programs