

# OpenMP: Open Multi-Processing

Chris Kauffman

*Last Updated:  
Thu Mar 24 09:33:03 AM CDT 2022*

# Logistics

## Today

OpenMP for shared memory machines

## Reading

- ▶ Grama 7.10 (OpenMP)
- ▶ [OpenMP Tutorial at Laurence Livermore](#)

# OpenMP: High-level Shared Memory Parallelism

- ▶ OpenMP = Open Multi-Processing
- ▶ A standard, implemented by various folks, compiler-makers
- ▶ Targeted at shared memory machines: multiple processing elements sharing memory
- ▶ Specify parallelism in code with
  - ▶ Some function calls: *which thread number am I?*
  - ▶ Directives: *do this loop using multiple threads/processors*
- ▶ Can orient program to work without need of additional processors - direct serial execution
- ▶ OpenMP targets multiple processors, new relative OpenACC which targets “accelerators” like GPUs with same ideas
- ▶ The *easiest* parallelism you'll likely get in C / C++ / Fortran

## #pragma in C

*The '#pragma' directive is the method specified by the C standard for providing additional information to the compiler, beyond what is conveyed in the language itself.*

– *[GCC Manual](#)*

- ▶ Similar in to Java's annotations (@Override)

- ▶ Indicate meta-info about code

```
printf("Normal execution\n");
```

```
#pragma do something special below  
normal_code(x,y,z);
```

- ▶ Several other pragmas supported by gcc including
  - ▶ once: include a header file once only
  - ▶ poison: if a poisoned identifier is used, cause an error
  - ▶ dependency: warn if another file is newer than this one

# OpenMP Basics

```
#pragma omp parallel  
single_parallel_line();
```

```
#pragma omp parallel  
{  
    parallel_block();  
    with_multiple(statements);  
    done_in_parallel();  
}
```

- ▶ Pragas indicate a single line or block should be done in parallel.
- ▶ Examine `omp_basics.c`

# Compiler Support for OpenMP

- ▶ Most other modern compilers have support for OpenMP
- ▶ GCC, CLang/LLVM, Intel C/C++ Compiler, MS Visual Studio, Portland Group / NVidia tools - all support OpenMP in various ways
- ▶ GCC supports OpenMP with appropriate options

```
>> gcc omp_basics.c                # no parallelism
>> gcc omp_basics.c -fopenmp        # enable parallelism
```
- ▶ OpenMP was introduced in the mid 90's and has expanded and added features which are available depending on platform

---

|                |     |     |     |     |     |     |
|----------------|-----|-----|-----|-----|-----|-----|
| GCC Version    | 4.2 | 4.4 | 4.7 | 4.9 | 6.0 | 9.0 |
| OpenMP Version | 2.5 | 3.0 | 3.1 | 4.0 | 4.5 | 5.0 |

---

# Hints at OpenMP Implementation

- ▶ OpenMP  $\approx$  coarse-grained parallelism
- ▶ PThreads  $\approx$  fine-grained parallelism
- ▶ From [libGOMP Documentation](#) (OMP library in GCC)

OMP CODE

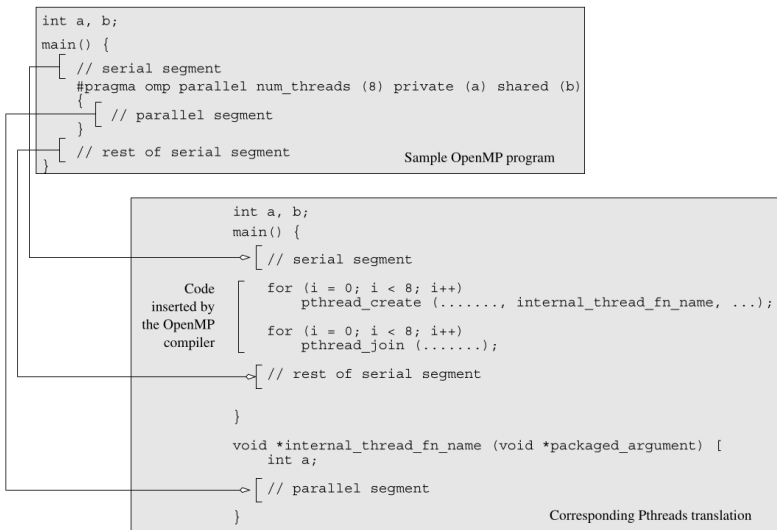
```
#pragma omp parallel
{
    body;
}
```

BECOMES

```
void subfunction (void *data){
    use data;
    body;
}
setup data;
GOMP_parallel_start (subfunction, &data, num_threads);
subfunction (&data);
GOMP_parallel_end ();
```

[Not exactly a source transformation](#), but OpenMP can leverage many existing pieces of Posix Threads libraries.

# Grama Sample Translation: OpenMP → PThreads



**Figure 7.4** A sample OpenMP program along with its Pthreads translation that might be performed by an OpenMP compiler.



# OpenMP Thread Identification

- ▶ OpenMP divides computation into *threads*
- ▶ Nearly identical model to PThreads approach BUT not always implemented via PThreads (icc may use [Intel Thread Building Blocks](#))
- ▶ Threads execute concurrently / in parallel, can have private data, shared data
- ▶ OpenMP provides basic id / environment functions for threads and synchronization constructs

```
#pragma omp parallel
{
    int thread_id = omp_get_thread_num();
    int num_threads = omp_get_num_threads();
    int work_per_thread = total_work / num_threads;
    ...;
}
```

# Specifying Number of Threads

```
#pragma omp parallel                                // Default # threads based on system config
{
    run_with_max_num_threads();
}

if (argc > 1) {                                     // Number of threads based on command line
    omp_set_num_threads( atoi(argv[1]) );
}
#pragma omp parallel
{
    run_with_current_num_threads();
}

#pragma omp parallel num_threads(2) // Number of threads as part of pragma
{
    run_with_two_threads();
}

int NT = 4;                                         // Number of threads from program variable
#pragma omp parallel num_threads(NT)
{
    run_with_four_threads();
}

>> OMP_NUM_THREADS=4 ./a.out                      // Set default via environment variable
```

# Tricky Memory Issues Abound

## Program Fragment

```
// omp_shared_variables.c

int id_shared=-1;
int numThreads=0;

#pragma omp parallel
{
    id_shared = omp_get_thread_num();
    numThreads = omp_get_num_threads();
    printf("A: Hey from thread %d / %d\n",
           id_shared, numThreads);
}

printf("\n");

#pragma omp parallel
{
    int id_private = omp_get_thread_num();
    numThreads = omp_get_num_threads();
    printf("D: Hey from thread %d / %d\n",
           id_private, numThreads);
}
```

## Possible Output

```
A: Hello from thread 2 of 4
A: Hello from thread 3 of 4
A: Hello from thread 0 of 4
A: Hello from thread 0 of 4

D: Hello from thread 1 of 4
D: Hello from thread 3 of 4
D: Hello from thread 0 of 4
D: Hello from thread 2 of 4
```

# Lessons

- ▶ OpenMP Threads share memory just like PThreads including heap, globals, any stack vars in master thread
- ▶ Threads share any stack variables NOT in parallel blocks
- ▶ Thread variables are **private** if declared inside parallel blocks
- ▶ Pragmas can be used to create private copies of otherwise shared variables
- ▶ Take care with shared variables: easy to accidentally share variables as programming language scope does not make sharing as obvious

## Exercise: Pi Calc via OpenMP

Examine:

[https://cs.umn.edu/~kauffman/5451/picalc\\_omp\\_reduction.c](https://cs.umn.edu/~kauffman/5451/picalc_omp_reduction.c)

### Questions

- ▶ Contrast the structure of the program with PThreads version
- ▶ How is the number of threads used to run determined?
- ▶ What is the business with `reduction(+: total_hits)`?
- ▶ Can variables like `points_per_thread` be moved out of the parallel block?
- ▶ Do you expect speedup for this computation?

## Answers: Pi Calc via OpenMP

- ▶ Contrast the structure of the program with PThreads version  
*Shorter and sweeter, no need for auxiliary function, casting, loops to create/join threads.*
- ▶ How is the number of threads used to run determined?  
*From the command line and set via the function `omp_set_num_threads()`*
- ▶ What is the business with `reduction(+: total_hits)`?  
*Performs a reduction on shared variable `total_hits`: correct results + performance; more in a moment...*
- ▶ Can variables like `points_per_thread` be moved out of the parallel block?  
*`points_per_thread` and `num_threads` can be shared; `thread_id` and `state` should NOT be shared.*
- ▶ Do you expect speedup for this computation?  
*Yes - get nearly linear speedup and correct results with less effort than PThreads version.*

## Exercise: Placement of Variables vs Runtime

Analyze these two examples and explain the **timing difference**

```
// (A) picalc_omp_reduction.c
#pragma omp parallel ...
{
    unsigned int state =
        123456789 * thread_id;
    ...
    double x =
        ((double) rand_r(&state))...
```

TIMING

```
>> time a.out 75000000 4
npoints: 75000000
hits:    58910475
pi_est:  3.141892
```

```
real    0m0.291s
user    0m1.125s
sys     0m0.004s
```

```
// (B) picalc_omp_rand_contention.c:
unsigned int state =
    123456789;
#pragma omp parallel...
{
    ...
    double x =
        ((double) rand_r(&state))...
```

TIMING

```
>> time -p a.out 75000000 4
npoints: 75000000
hits:    58910901
pi_est:  3.141915
```

```
real    0m1.200s
user    0m4.285s
sys     0m0.001s
```

## Answers: Placement of Variables vs Runtime

- ▶ (A) `picalc_omp_reduction.c` places the state variable within the parallel region - becomes **thread private**
- ▶ (B) `picalc_omp_rand_contention.c` places its state outside so it is a **shared variable** among threads
- ▶ Each call to `rand_r()` must alter state so there is memory contention around it



## Note on rand()

- ▶ `rand_r()` is **reentrant** and **thread-safe**
  - ▶ When programming in multi-threaded contexts look for these qualities
  - ▶ *Note:* When calling `rand_r()` in multiple threads with the *same* state variable, likely to lose “randomness”
- ▶ `rand()` is another matter...
  - ▶ Generates random numbers a la `int r = rand();`
  - ▶ Uses a “hidden” global variable to track generator state
  - ▶ For many moons, was NOT thread safe
  - ▶ Most Linux / GLIBC implementations are thread safe, but...
  - ▶ Likely use a mutex to protect the state variable slowing things down considerably...

```
>> time ./picalc_omp_rand 75000000 1
```

```
...  
real    0m1.439s  
...
```

```
>> time ./picalc_omp_rand 75000000 4
```

```
...  
real    1m3.403s  
...
```

# Reductions in OpenMP

omp\_picalc.c used a reduction() clause

```
//          operation  ---+  +--- variable
//          V          V
#pragma omp parallel reduction(+: total_hits)
{
    ...;
    total_hits++;
}
```

- ▶ Shared var total\_hits is updated “properly” and reasonably efficiently
  - ▶ May exploit the fact that addition is transitive - can be done in any order
  - ▶ Likely to introduce a private version of reduction variable for each thread then reduce over threads at the end
  - ▶ Alternatively may utilize a mutex or hardware atomic ops
- ▶ Most other arithmetic ops available
- ▶ Statement of **policy** rather than **mechanism**

# OpenMP Atomic Pragmas

```
#pragma omp parallel
{
    ...;
    #pragma omp atomic
    total_hits++;
}
```

- ▶ Use atomic hardware instruction available
- ▶ Restricted to single operations, usually arithmetic
- ▶ No hardware support → compilation problem

```
#pragma omp atomic
printf("woot"); // compile error
```

## Alternative: Critical Block

```
#pragma omp parallel
{
    ...;
    #pragma omp critical
    {
        total_hits++;
    }
}
```

- ▶ Not restricted to hardware supported ops
- ▶ Uses locks to restrict access to a single thread

## Reduction vs. Atomic vs. Critical

- ▶ `omp_picalc_alt.c` has commented out versions of for each of reduction, atomic, and critical
- ▶ Examine timing differences between the three choices

```
lila [openmp-code]% gcc omp_picalc_alt.c -fopenmp
lila [openmp-code]% time -p a.out 100000000 4
npoints: 100000000
hits:      78541717
pi_est:    3.141669
```

```
real ??? - Elapsed (wall) time
user ??? - Total user cpu time
sys  ??? - Total system time
```

| Time      | Threads | real | user  | sys  |
|-----------|---------|------|-------|------|
| Serial    | 1       | 1.80 | 1.80  | 0.00 |
| Reduction | 4       | 0.52 | 2.00  | 0.00 |
| Atomic    | 4       | 2.62 | 9.98  | 0.00 |
| Critical  | 4       | 9.02 | 34.46 | 0.00 |

## Exercise: No Reduction for You

```
int total_hits=0;
#pragma omp parallel reduction(+: total_hits)
{
    int num_threads = omp_get_num_threads();
    int thread_id = omp_get_thread_num();
    int points_per_thread = npoints / num_threads;
    unsigned int state = 123456789 * thread_id;
    int i;
    for (i = 0; i < points_per_thread; i++) {
        double x = ((double) rand_r(&state)) / ((double) RAND_MAX);
        double y ~ ((double) rand_r(&state)) / ((double) RAND_MAX);
        if (x*x + y*y <~ 1.0){
            total_hits++;
        }
    }
}
```

- ▶ Alter picalc to NOT use reduction clause
- ▶ Use alternative like atomic or critical
- ▶ **Goal:** achieve same/better speed as reduction version

# Answers: No Reduction for You

```
// picalc_omp_atomic.c:
#pragma omp parallel
{
    int num_threads = omp_get_num_threads();
    int thread_id = omp_get_thread_num();
    int points_per_thread = npoints / num_threads;
    int my_hits = 0;           // private count
    unsigned int state = 123456789 * thread_id;
    int i;
    for (i = 0; i < points_per_thread; i++) {
        double x = ((double) rand_r(&state)) / ((double) RAND_MAX);
        double y = ((double) rand_r(&state)) / ((double) RAND_MAX);
        if (x*x + y*y <= 1.0){
            my_hits++;
        }
    }
    #pragma omp atomic
    total_hits += my_hits;      // lock total_hits before updating
}
```

## Thread Variable Declarations

Pragmas can specify that variables are either shared or private. See `omp_private_variables.c`

```
tid = -1;
// #pragma omp parallel
#pragma omp parallel shared(tid)
{
    tid = omp_get_thread_num();
    printf("Hello World from thread = %d\n", tid);
}
```

```
tid = -1;
#pragma omp parallel private(tid)
{
    tid = omp_get_thread_num();
    printf("Hello World from thread = %d\n", tid);
}
```

Also available

- ▶ `shared` which is mostly redundant
- ▶ `firstprivate` guarantees initialization with shared value
- ▶ All of these are subsumed by lexical scoping in C



# Parallel Loops

```
#pragma omp parallel for
for (int i = 0; i < 16; i++) {
    int id = omp_get_thread_num();
    printf("Thread %d doing iter %d\n",
           id, i);
}
```

## OUTPUT

```
Thread 0 doing iter 0
Thread 0 doing iter 1
Thread 0 doing iter 2
Thread 0 doing iter 3
Thread 2 doing iter 8
Thread 2 doing iter 9
Thread 2 doing iter 10
Thread 2 doing iter 11
Thread 1 doing iter 4
Thread 1 doing iter 5
...
```

- ▶ OpenMP supports parallelism for independent loop iterations
- ▶ Note variable `i` is declared in loop scope
- ▶ Iterations **automatically divided** between threads in a blocked fashion
- ▶ **Assumption:** Loop iterations are independent

## Exercise: OpenMP Matrix Vector Multiply

```
// matvec_serial.c: Matrix/vector multiply demo
for(i=0; i<rows; i++){
    for(j=0; j<cols; j++){
        result[i] += matrix[i][j] * vector[j];
    }
}
```

- ▶ Describe 3 ways one might parallelize this operation
- ▶ Write OpenMP #pragmas for each
- ▶ Note: reduction on an array variables varies based on OpenMP version

# Answers: OpenMP Matrix Vector Multiply

```
// Outer for loop multiplication
#pragma omp parallel for
for(int i=0; i<rows; i++){
    for(int j=0; j<cols; j++){
        result[i] += matrix[i][j] * vector[j];
    }
}

// Inner for loop multiplication: reduction
// on result[i] added in recent OpenMP
for(int i=0; i<rows; i++){
    #pragma omp parallel for reduction(+:result[i])
    for(int j=0; j<cols; j++){
        result[i] += matrix[i][j] * vector[j];
    }
}

// Outer and Inner for loop multiplication
#pragma omp parallel for
for(int i=0; i<rows; i++){
    #pragma omp parallel for reduction(+:result[i])
    for(int j=0; j<cols; j++){
        result[i] += matrix[i][j] * vector[j];
    }
}
```

# Timing Differences

## Circa 2017

# Desktop

```
>> gcc omp_matvec_timing.c -fopenmp
```

# SKINNY

```
>> a.out 20000 10000
```

```
Outer :    0.2851
```

```
Inner :    0.2022
```

```
Both  :    0.2191
```

# FAT

```
> a.out 10000 20000
```

```
Outer :    0.2486
```

```
Inner :    0.1911
```

```
Both  :    0.2118
```

```
> export OMP_NESTED=true
```

```
> a.out 20000 10000
```

```
Outer :    0.2967
```

```
Inner :    0.2027
```

```
Both  :    1.1783
```

## Today

# Laptop

```
>> gcc matvec_omp.c -O3 -fopenmp
```

# SKINNY -

```
>> OMP_NESTED=false ./a.out 20000 10000
```

```
Outer :    0.1568
```

```
Inner :    0.1888
```

```
Both   :    0.1515
```

# FAT

```
>> OMP_NESTED=false ./a.out 10000 20000
```

```
Outer :    0.1490
```

```
Inner :    0.1869
```

```
Both   :    0.1484
```

# OMP\_NESTED=true is default

```
>> ./a.out 20000 10000
```

```
Outer :    0.1559
```

```
Inner :    0.1935
```

```
Both   :    3.5133
```

# Nested Parallelism Control

- ▶ By default nested parallelism is
  - ▶ Enabled in most recent GCC
  - ▶ Disabled in older GCC versions
- ▶ Like other aspects of OpenMP, can control nested parallelism via function calls like

```
omp_set_nested(1); // ON  
omp_set_nested(0); // OFF
```

- ▶ Can also be specified via environment variables

```
export OMP_NESTED=true  
export OMP_NESTED=false  
export OMP_NUM_THREADS=4
```

- ▶ Env. Vars are handy for experimentation
- ▶ Other Features such as loop scheduling are controllable via directives, function calls, or environment variables

# Syntax Note

```
#pragma omp parallel
{
    #pragma omp for
    for (int i = 0; i < REPS; i++) {
        int id = omp_get_thread_num();
        printf("Thread %d did iter %d\n",
              id, i);
    }
}
printf("\n");
```

// ABOVE AND BELOW IDENTICAL

```
#pragma omp parallel for
for (int i = 0; i < REPS; i++) {
    int id = omp_get_thread_num();
    printf("Thread %d did iter %d\n",
          id, i);
}
printf("\n");
```

- ▶ Directives for OpenMP can be separate or coalesced
- ▶ Code on top and bottom are parallelized the same way
- ▶ In top code, removing first `#pragma` removes parallelism

# Loop Scheduling - 4 Types

## Static

- ▶ So far only done static scheduling with fixed size chunks
- ▶ Threads get fixed size chunks in rotating fashion
- ▶ Great if each iteration has same work load

## Dynamic

- ▶ Threads get fixed chunks but when done, request another chunk
- ▶ Incurs more overhead but balances uneven load better

## Guided

- ▶ Hybrid between static/dynamic, start with each thread taking a “big” chunk
- ▶ When a thread finishes, requests a “smaller” chunk, next request is smaller

## Runtime

- ▶ Environment variables (OMP\_SCHEDULE) used to select one of the others
- ▶ Flexible but requires user awareness

# Basic Loop Scheduling

```
// omp_loop_scheduling.c, assumes OMP_NUM_THREADS=4
const int REPS = 16;

#pragma omp parallel for schedule(static)
for (int i = 0; i < REPS; i++) { // thr 0: 0-3, thr 1: 4-7
    ...                          // thr 2: 8-11, thr 4: 12-15
}

#pragma omp parallel for schedule(static,2)
for (int i = 0; i < REPS; i++) { // thr 0: 0,1,8,9   thr 1: 2,3,10,11
    ...                          // thr 2: 4,5,12,13 thr 3: 6,7,14,15
}

#pragma omp parallel for schedule(dynamic,2)
for (int i = 0; i < REPS; i++) { // varies, all start with 2 iters
    ...                          // request more as completed
}

#pragma omp parallel for schedule(guided)
for (int i = 0; i < REPS; i++) {
    ...                          // varies, start with large chunks
}                                // request smaller chunks

#pragma omp parallel for schedule(runtime)
for (int i = 0; i < REPS; i++) {
    ...                          // controlled via environment var
}                                // ex: OMP_SCHEDULE=static
```



# Code for Loop Scheduling

- ▶ `omp_loop_scheduling.c` demonstrates loops of each kind with printing
- ▶ `omp_guided_schedule.c` longer loop to demonstrate iteration scheduling during Guided execution

## Exercise: Spell Checking

- ▶ Consider a spell checking problem
- ▶ Look up each word in a document in a dictionary to determine correct spelling
- ▶ If document word is not in the dictionary, report a misspelling

```
// fragment from spellcheck_omp.c
for (int i=0; i < document->word_count; i++) {
    int result =
        linear_search(dictionary, document->words[i]);
    if(result == -1){
        misspelled++;
    }
}
```

## Questions

1. Parallelize the “outer” loop over words or the “inner” loop that is `linear_search()`
2. Which type of loop schedule seems to make the most sense? Static? Dynamic? Guided?

## Answers: Spell Checking

1. Parallelize the “outer” loop over words or the “inner” loop that is `linear_search()`

*For a large number of words, outer “word” loop makes more sense than inner loop : induces less thread startup overhead. For a small number of words, may be more worthwhile to parallelize inner loop.*

2. Which type of loop schedule seems to make the most sense? Static? Dynamic? Guided?

*Dynamic or Guided makes more sense. Especially with `linear_search()`, expect that some checks will take longer than others which means a Static schedule may lead to some threads with much more work than others.*

# Example Runs on Spellcheck w/ Word Loop Parallelized

```
>> time OMP_SCHEDULE=static spellcheck_omp ...
threads = 8
misspelled: 0
Thread 0 work: 110803941
Thread 1 work: 332426710
Thread 2 work: 554049479
Thread 3 work: 775672248
Thread 4 work: 997295017
Thread 5 work: 1218917786
Thread 6 work: 1440540555
Thread 7 work: 1662044229
Total work: 7091749965
```

```
real    0m12.110s
user    0m53.495s
sys     0m0.008s
```

```
>> time OMP_SCHEDULE=guided spellcheck_omp ...
threads = 8
misspelled: 0
Thread 0 work: 901203843
Thread 1 work: 892041145
Thread 2 work: 897067217
Thread 3 work: 895931158
Thread 4 work: 850295834
Thread 5 work: 892967175
Thread 6 work: 896993276
Thread 7 work: 865250317
Total work: 7091749965
```

```
real    0m8.853s
user    1m9.492s
sys     0m0.031s
```

```
>> time OMP_SCHEDULE=dynamic spellcheck_omp ...
threads = 8
misspelled: 0
Thread 0 work: 851351653
Thread 1 work: 887921206
Thread 2 work: 908569538
Thread 3 work: 893075776
Thread 4 work: 882219930
Thread 5 work: 873179476
Thread 6 work: 904986970
Thread 7 work: 890445416
Total work: 7091749965
```

```
real    0m7.877s
user    1m0.578s
sys     0m0.011s
```

```
>> time OMP_SCHEDULE=static,1 spellcheck_omp ...
threads = 8
misspelled: 0
Thread 0 work: 886431528
Thread 1 work: 886446415
Thread 2 work: 886461302
Thread 3 work: 886476189
Thread 4 work: 886491076
Thread 5 work: 886505963
Thread 6 work: 886520850
Thread 7 work: 886416642
Total work: 7091749965
```

```
real    0m7.665s
user    1m0.295s
sys     0m0.011s
```

# Notes on Spellcheck

- ▶ Pure static scheduling does not balance the work well
- ▶ Dynamic / Guided gives reasonable performance improvement over pure Static scheduling
- ▶ Specific instance of
  - >> `spellcheck_omp english-words.txt english-words.txt`  
allows for block-cyclic distribution for 0-overhead fair distribution of work
- ▶ Most problems where work distribution is unknown benefit from dynamic or guided scheduling

## Sections: Non-loopy Parallelism

- ▶ Independent code can be “sectioned” with threads taking different sections.
- ▶ Good to parallelize distinct independent execution paths
- ▶ See `omp_sections.c`

```
#pragma omp sections
{
    #pragma omp section
    {
        printf("Thread %d computing d[]\n",
               omp_get_thread_num());
        for (i=0; i < N; i++)
            d[i] = a[i] * b[i];
    }

    #pragma omp section
    printf("Thread %d chillin' out\n",
           omp_get_thread_num());
}
```

# Locks in OpenMP

- ▶ Implicit parallelism/synchronization is awesome but...
- ▶ Occasionally need more fine-grained control
- ▶ Lock facilities provided to enable mutual exclusion
- ▶ Each of these have analogues in PThreads we will discuss later

```
void omp_init_lock(omp_lock_t *lock);    // create
void omp_destroy_lock(omp_lock_t *lock); // destroy
void omp_set_lock(omp_lock_t *lock);     // wait to obtain
void omp_unset_lock(omp_lock_t *lock);   // release
int  omp_test_lock(omp_lock_t *lock);    // check, don't wait
```