

# CUDA Programming and GPU Architecture

Chris Kauffman

*Last Updated:  
Tue Nov 16 10:18:01 PM CST 2021*

# Logistics

# GPUs will Feel Different

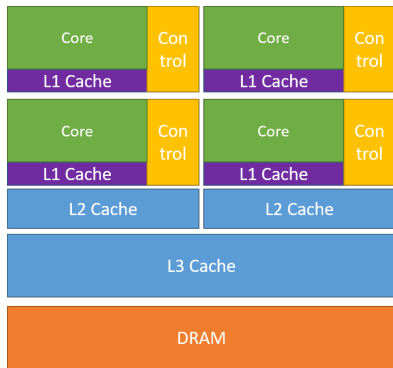
## Distributed / Threaded Programming

- ▶ Most effective strategies looked for ways to assign lots of work to limited number of procs/threads
- ▶ Poo-pooed the idea of “Assume length  $N$  array and  $N$  processors”


## GPU Programming

- ▶ Threads are essentially cost-free, close to theoretical models so...
- ▶ Assume length  $N$  array and  $N$  processors
- ▶ Will require some mental adjustment

# GPU vs CPU



CPU



GPU

Source: NVidia Docs "CUDA C++ Programming Guide"

# GPUs are a Co-Processor or Accelerator

- ▶ CPU is still in charge, has access to main memory
- ▶ GPU is a partner chip, has a distinct set of memory
- ▶ Sections of code will feel like Distributed architecture
  - ▶ CPU / GPU memory transfers
  - ▶ Barriers / synchronization as CPU waits for GPU to finish
- ▶ GPU itself is like a multicore system on steroids

# Why do GPUs Look like this?

# CUDA : NVidia's General Purpose GPU Technology

# CUDA Terminology

**Thread** A set of operations; can be as small as a single addition

**Kernel** A function which expresses what a thread should do

**Block** A group of executing threads which can share some local memory

**Execution Context** Run a Kernel function with a specified number of Blocks, Threads per Block, and amount of shared memory

**Host** The CPU - sets Execution Context, launches Kernels on GPU

**Device** The GPU which runs Kernels



# Hello CUDA

Examine `hello.cu`: `.cu` extension favored for CUDA programs; it is C++ w/ a few extras

## Kernel

```
__global__ void helloFromGPU() { // __global__ => called from CPU/GPU
                                // runs on GPU

    printf("Block %02d Thread %02d: Hello World\n",
           blockIdx.x,           // ever-present structs which gives
           threadIdx.x);         // each GPU thread indexing info
}
```

## Execution Context

```
int main (int argc, char *argv[]){
    printf("CPU: Running 1 block w/ 16 threads\n");
    helloFromGPU<<<1,16>>>>(); // executes in 1 block, 16 threads per block
    cudaDeviceSynchronize();    // ensures GPU completes operations
}
```