

# Parallel Languages and Platforms

Chris Kauffman

*Last Updated:  
Mon Dec 13 07:39:17 AM CST 2021*

# Logistics

## Today

- ▶ Finish Applications
- ▶ Parallel Languages and Platforms

## Assignments

- ▶ Questions on A2?
- ▶ A3 Canceled: all students who **Submit A2** will get +10% for A3\*

## Schedule

Mon 12/13	Applications Parallel Languages
Wed 12/15	Review A2 Due Course Evals Due
Fri 12/17	A2 Late Deadline
Mon 12/20	Final Exam 1:30-3:30pm Lecture Location

\* Originally was “take Final Exam” but “submit A2” now

# Course Feedback

## Official Student Rating of Teaching (SRTs)

- ▶ Official UMN Evals are done online this semester
- ▶ Available here: <https://srt.umn.edu/blue>
- ▶ **EVALUATE YOUR LECTURE SECTION: 001**
- ▶ **Due** Wed 12/15/2021 by 9:45am
- ▶ Response Rate  $\geq 80\%$   $\rightarrow$  One Final Exam Question Revealed during lecture on 12/15

# Menagerie of Parallel Languages and Platforms

## Distributed Memory Only

Erlang, Map+Reduce / Hadoop, Job Schedulers

## Shared Memory Only

Cilk, Clojure

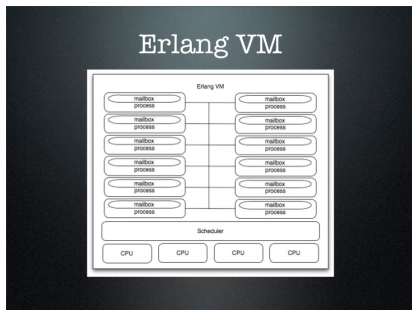
## Distributed + Shared

Unified Parallel C, Chapel (later)

## Device Concurrency / GPUs

CUDA / OpenCL (later)

# Erlang



Source

- ▶ Developed for distributed computation, telephony systems
- ▶ Virtual machine which mirrors many OS functions
- ▶ Process spawn to create *lightweight* procs
- ▶ send/receive clauses to share information among processes
- ▶ Facilities to contact a remote Erlang VM and talk to its processes

# Erlang Sample straight from Wikipedia

```
% Create a process on this machine and invoke the function
%   web:start_server(Port,MaxConnections)
```

```
ServerProcess = spawn(web, start_server, [Port, MaxConnections]),
```

```
% Create a remote process and invoke the function
%   web:start_server(Port, MaxConnections)
% on machine RemoteNode
```

```
RemoteProcess = spawn(RemoteNode, web, start_server, [Port, MaxConnections]),
```

```
% Send a message to ServerProcess (asynchronously). The message
% consists of a tuple with the atom "pause" and the number "10".
```

```
ServerProcess ! {pause, 10},
```

```
% Receive messages sent to this process
```

```
receive
```

```
    a_message -> do_something;
```

```
    {data, DataContent} -> handle(DataContent);
```

```
    {hello, Text} -> io:format("Got hello message: ~s", [Text]);
```

```
    {goodbye, Text} -> io:format("Got goodbye message: ~s", [Text])
```

```
end.
```

# Erlang's Nature and Target

- ▶ Syntax and semantics are somewhat odd/archaic but can be “gotten used to”
- ▶ Targeted at client server architectures, computation distributed across many nodes
- ▶ Well known for robustness of the VM, fault-tolerance features to keep application going if participating nodes go down
- ▶ Not targeted at high-performance computation / scientific problems, more towards business, IT, web services

# MapReduce (or more properly Map, Shuffle, Reduce)

- ▶ A style of programming, inspired by functional programming  

```
(def doub-sum (reduce + 0 (map double '(1 2 3 4 5))))
```
- ▶ Targeted at big data: large distributed stores of data
  - ▶ Map: Transform / filter data in some way
  - ▶ Shuffle: Move data with same properties to same node
  - ▶ Reduce: Combine results on individual nodes

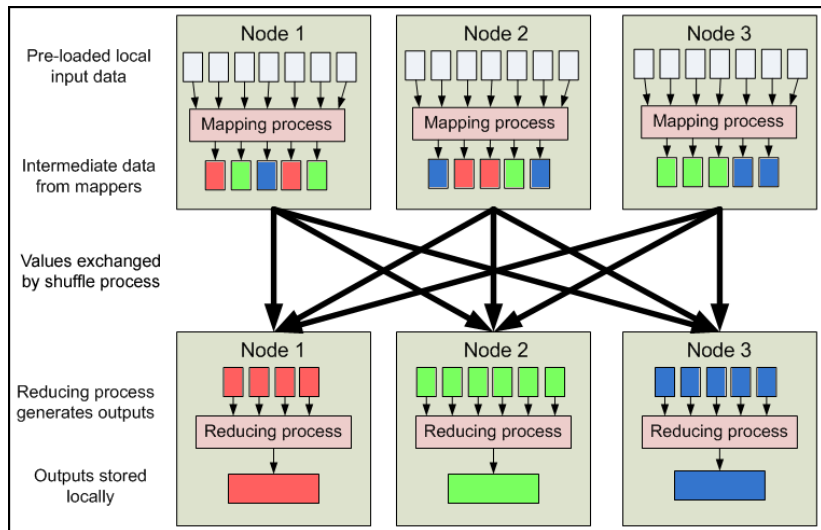
## Fault-tolerance

by @jrecursive





# Basic Architecture of MapReduce



Source: Yahoo Developer Tutorial on MapReduce

# Shameless Wikipedia Example: Document Word Counts

## Pseudocode

```
function map(String name,
               String document):
    // name: document name
    // document: document contents
    for each word w in document:
        emit (w, 1)

function reduce(String word,
                Iterator partialCounts):
    // word: a word to count
    // partialCounts: list of
                    partial counts

    sum = 0
    for each pc in partialCounts:
        sum += pc
    emit (word, sum)
```

- ▶ Goal: produce frequency of each word in a document
- ▶ Nodes are each fed the document
- ▶ During `reduce()` *emit* pairs like ("apple",1) and ("Dell",1)
- ▶ System automatically sends pairs with key apple to the same nodes (redistribute)
- ▶ Nodes run `reduce()` to count apple occurrences, may redistribute further

# Variety of Languages for MapReduce Framework

```
public static class MapClass extends MapReduceBase
    implements Mapper<LongWritable, Text, Text, IntWritable>
{
    private final static
        IntWritable one = new IntWritable(1);
    private Text word = new Text();
    public void
map(LongWritable key, Text value,
    OutputCollector<Text, IntWritable> output,
    Reporter reporter) throws IOException
    {
        String line = value.toString();
        StringTokenizer itr = new StringTokenizer(line);
        while (itr.hasMoreTokens()) {
            word.set(itr.nextToken());
            output.collect(word, one);
        }
    }
}

public static class Reduce extends MapReduceBase
    implements Reducer<Text, IntWritable, Text, IntWritable>
{
    public void
reduce(Text key, Iterator<IntWritable> values,
    OutputCollector<Text, IntWritable> output,
    Reporter reporter) throws IOException
    {
        int sum = 0;
        while (values.hasNext()) {
            sum += values.next().get();
        }
        output.collect(key, new IntWritable(sum));
    }
}
```

Problem: Word  
Frequencies in a  
Document

← Java (Source)

Pig Latin ↓ (Source)

```
A = load './input.txt';
B = foreach A generate
    flatten(TOKENIZE((chararray)$0))
    as word;
C = group B by word;
D = foreach C generate COUNT(B), group;
store D into './wordcount';
```

# MapReduce Framework Notes

- ▶ Primary contribution of implementations is distributing load across many machines efficiently
- ▶ When machines both store some data and participate in MapReduce, gain locality for speed
- ▶ Alternative to large database processing, may open up opportunities for parallelism to avoid read/write locks in traditional DBs
- ▶ Most frequently referenced implementation is Apache Hadoop



But other implementations exist, some proprietary

- ▶ All implementation implement a MapReduce server/scheduler to which jobs are submitted

```
> javac MRWordCount.java
> java MRWordCount &
> bin/hadoop job -list
1 jobs currently running
```

JobId	State	StartTime	UserName
job_0001	1	1218506470390	kauffman

# Job Schedulers

- ▶ Has long been the need for many parallel jobs to be run on individual systems/clusters
- ▶ Job schedulers offer frameworks for this: submit many programs to run, scheduler assigns resources
- ▶ `slurm`: scheduler on `medusa` cluster which you used to schedule jobs running
- ▶ Generic form of easy concurrency for variety of different programs, resources etc.

## Cilk and CilkPlus (supported in gcc/g++)

- ▶ [Tutorial Here](#)
- ▶ Extensions to C/C++ which enable easy spawning of threads (*strands*) to run functions concurrently
- ▶ Primary Additions are keywords to easily spawn functions

```
// Run func concurrently - separate thread
```

```
int x = cilk_spawn func(n);
```

```
...
```

```
// Wait for all running functions to finish
```

```
cilk_sync;
```

```
// Compile with cilk features enabled
```

```
> gcc -fcilkplus cilk-fib.c
```

- ▶ Examine: cilk-fib.c
- ▶ Contrast with PThread startup

## cilk\_for loop, Reducers

- ▶ Similar to OpenMP, Cilk provides loop parallelization and reduction
- ▶ See `cilk-for-picalc.cpp`


```
#include <cilk/cilk.h>
#include <cilk/reducer_opadd.h>

// C++ class for reductions
cilk::reducer_opadd<int> total_hits;

// Appropriate for parallel contexts
total_hits++;

// Retrive final value
int th = total_hits.get_value();

// Automatic loop parallelization
cilk_for (int i = 0; i < npoints; i++) {
    ...;
}
```



## Array-based Operations

- ▶ Cilk provides some convenient syntax for array operations
- ▶ **Vectorized** operations automatically created
- ▶ Compile can sometimes do this but special notation helps hint

```
// standard element-wise vector multiply
void axpy(int n, double alpha, const double *x, double *y)
{
    for (int i = 0; i < n; i++) {
        y[i] += alpha * x[i];
    }
}

// Cilk Plus abbreviated syntax
void axpy_cilk(int n, double alpha, const double *x, double *y)
{
    y[0:n] += alpha * x[0:n];
}
```

- ▶ Matlab/Octave are well known for this style
- ▶ See `cilk-array-syntax.c` for uses



# Clojure and Software Transactional Memory

- ▶ A lisp which runs on the Java Virtual Machine



- ▶ Design Goal: allow for shared memory parallelism to be exploited
- ▶ Realization:
  - ▶ Each data element has well-defined local/shared semantics
  - ▶ Data is immutable by default
  - ▶ Provides atom data types for atomic alterations
  - ▶ Other alterations to shared resources occur in dosync blocks
  - ▶ Software Transactional Memory (STM) system: try changing a shared area, if it changes, try again with the new current value
- ▶ Runs as an executable JAR


```
> java -jar clojure-1.8.0.jar
Clojure 1.8.0
user=> (def x (atom 0))
#'user/x
user=> x
#object[clojure.lang.Atom 0x6c372fe6 {:status :ready, :val 0}]
user=> @x
0
```

## Picalc in Clojure

```
;; Serial version with atomic updates
(defn calc-pi-atoms [iterations]
  (let [hits (atom 0)]
    (dotimes [i iterations]
      (let [x (rand) y (rand)]
        (if (<= (+ (* x x) (* y y)) 1)
          (swap! hits inc))))
    (double (* (/ @hits iterations) 4))))

;; Parallel version with atomic updates
(defn calc-pi-atoms [iterations nthreads]
  (let [hits (atom 0)]
    (doall (pmap ; parallel map, force evaluation
              (fn [x]
                (dotimes [i (/ iterations nthreads)]
                  (let [x (rand) y (rand)]
                    (if (<= (+ (* x x) (* y y)) 1)
                      (swap! hits inc))))
              (range nthreads))) ; map onto number of threads
    (double (* (/ @hits iterations) 4))))
```

# Common HPC Parallel Platform



- ▶ Similar config to medusa cluster we used
- ▶ Cluster of machines, each with multiple cores
- ▶ Options to program:
  - ▶ Serial execution on each core/machine
  - ▶ Parallel shared memory execution on each machine
  - ▶ Parallel distributed memory execution on each core
  - ▶ **Mixed**: Distributed/Shared parallel execution



**Mixing** MPI  
and OpenMP

## Unified Parallel C (requires special compiler)

- ▶ Extensions to the C language
- ▶ Aimed at BOTH shared memory and distributed memory
- ▶ Automatic THREADS and id MYTHREAD variables
- ▶ Thread is more generalized: might be same machine (shared) or different machine (distributed)
- ▶ Shared memory blocks  
`shared int all_hits[THREADS];`  
Access like `int x = all_hits[1];` will work locally (shared) or via MPI-style message passing if remote (distributed)
- ▶ Compiler/runtime automatically sets up sharing
- ▶ Standard locks
- ▶ Automatic loop parallelization via `upc_forall(..)` with some *affinity* control: which thread executes which iteration
- ▶ Control over layout of shared blocks of memory: which thread gets what section
- ▶ Examine `upc-shared-picalc.c`